



CUDA-Q and Quantum Accelerated Supercomputing

Monica VanDieren, Sr Technical Marketing Engineer | August 2024

Agenda

- What is ~~Quantum~~ Accelerated Supercomputing

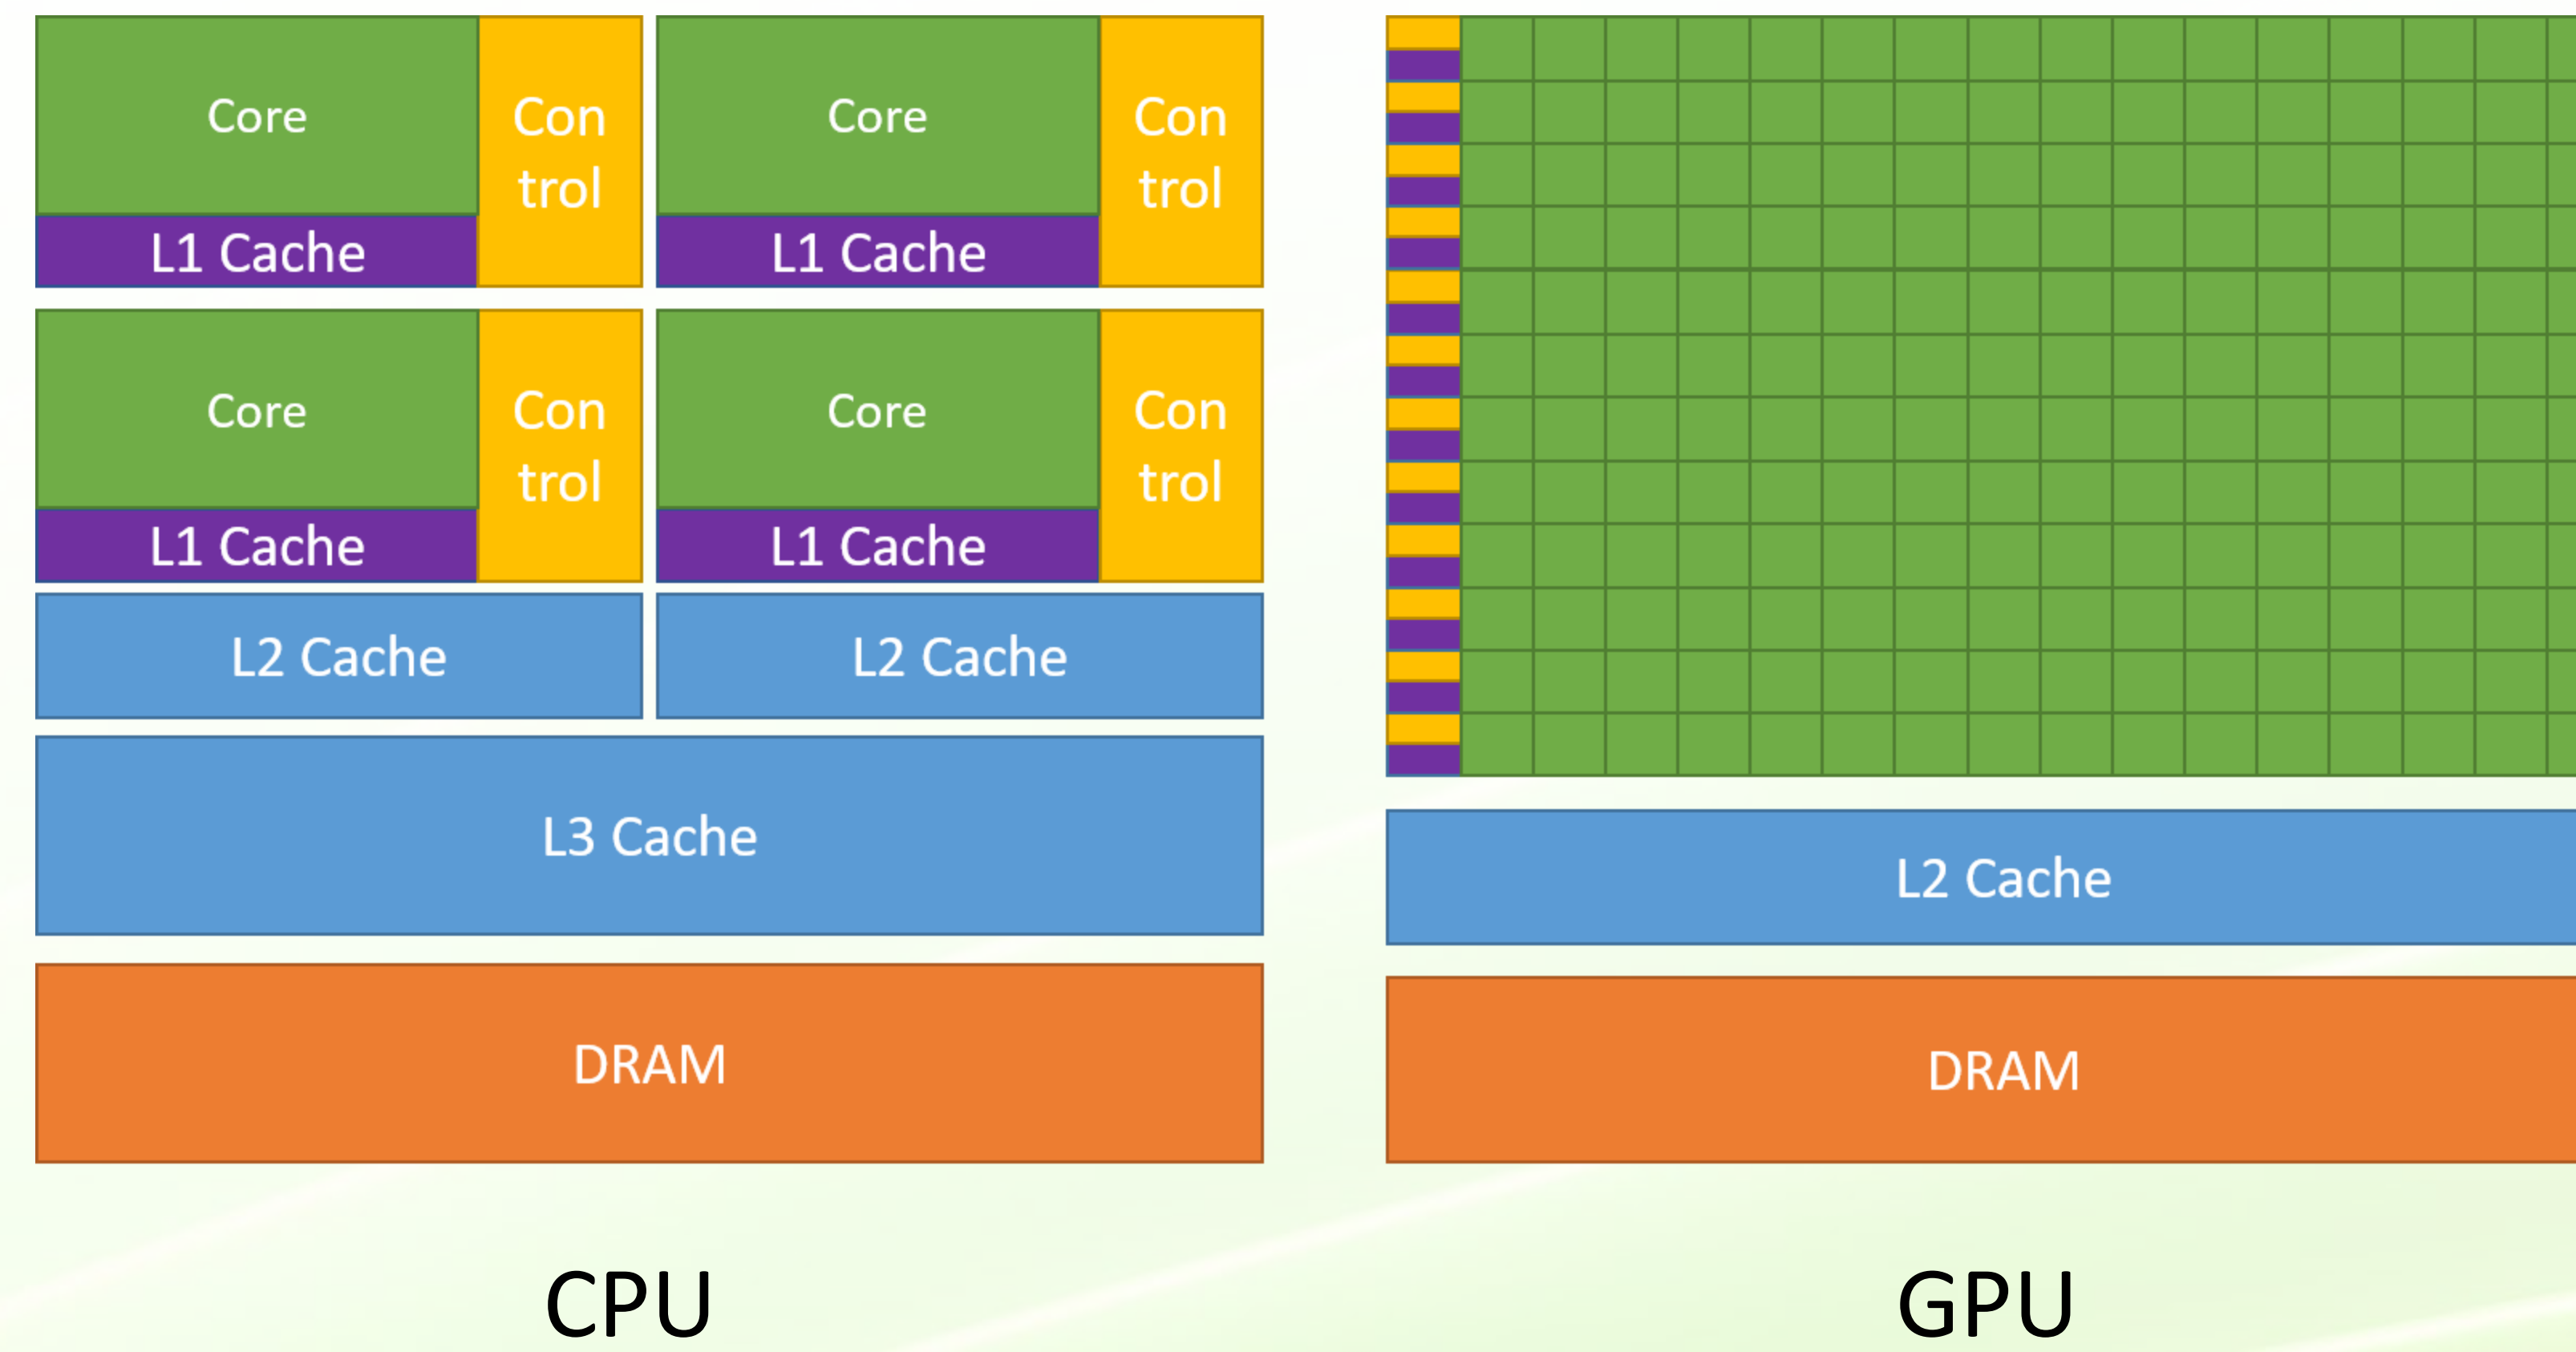
- Useful Quantum Simulation

- How-to Guide to CUDA-Q

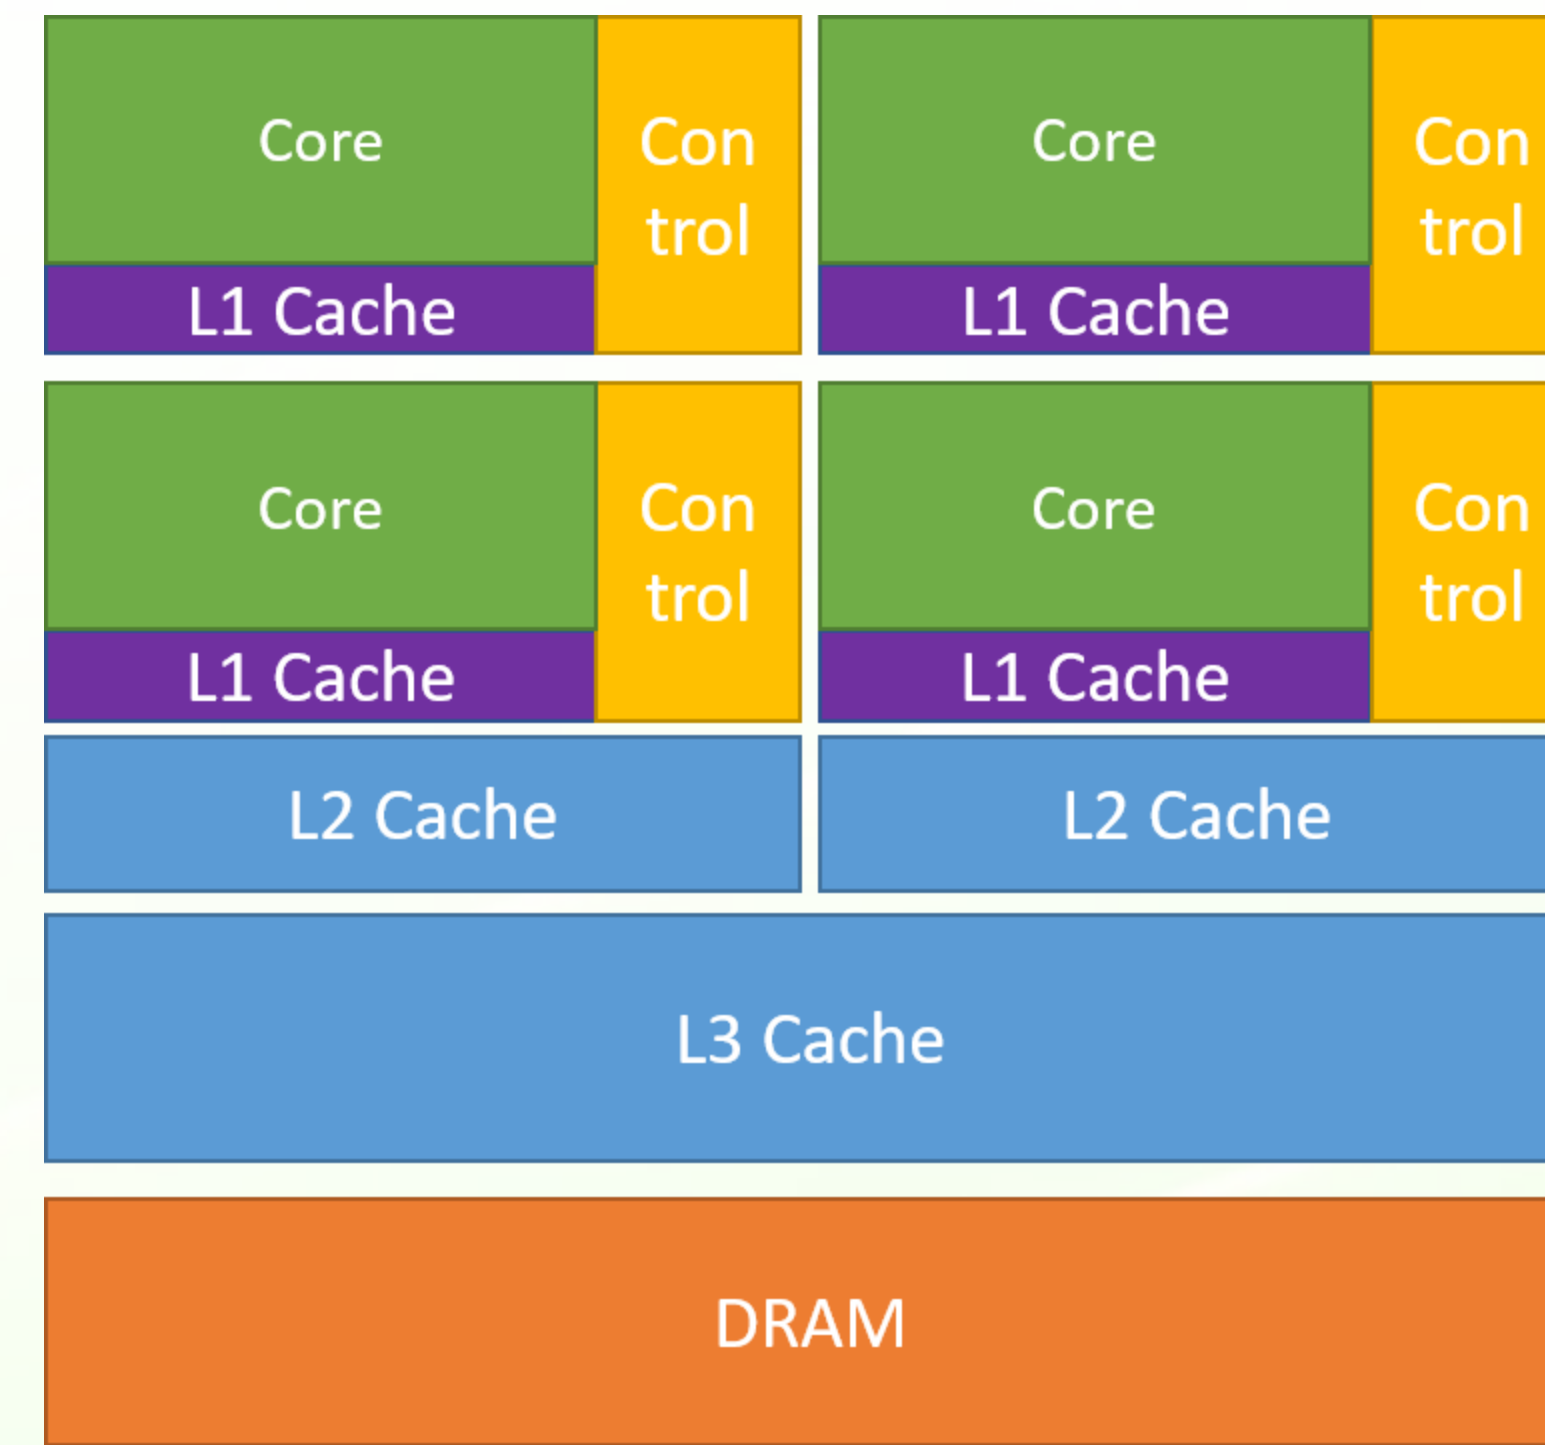
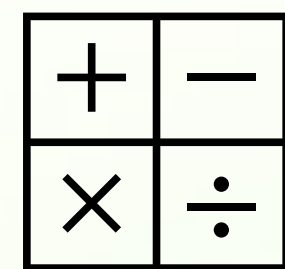
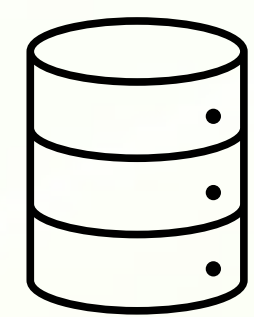
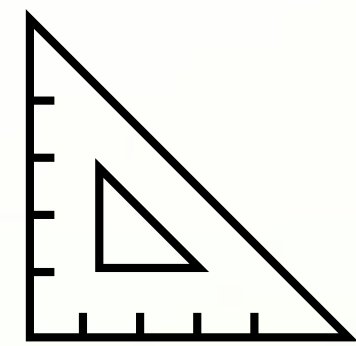
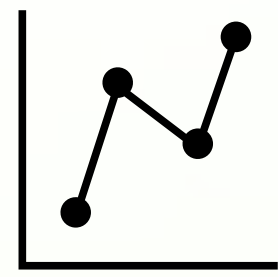
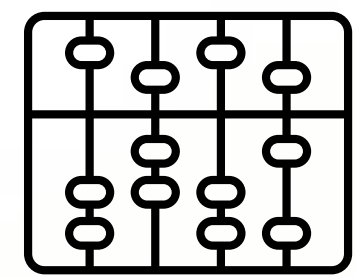
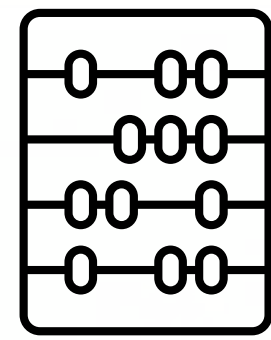
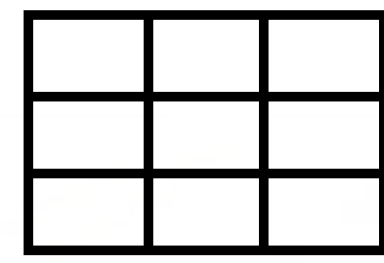
- Distributed Quantum Computing

- Conclusion

Accelerated Supercomputing



Accelerated Supercomputing

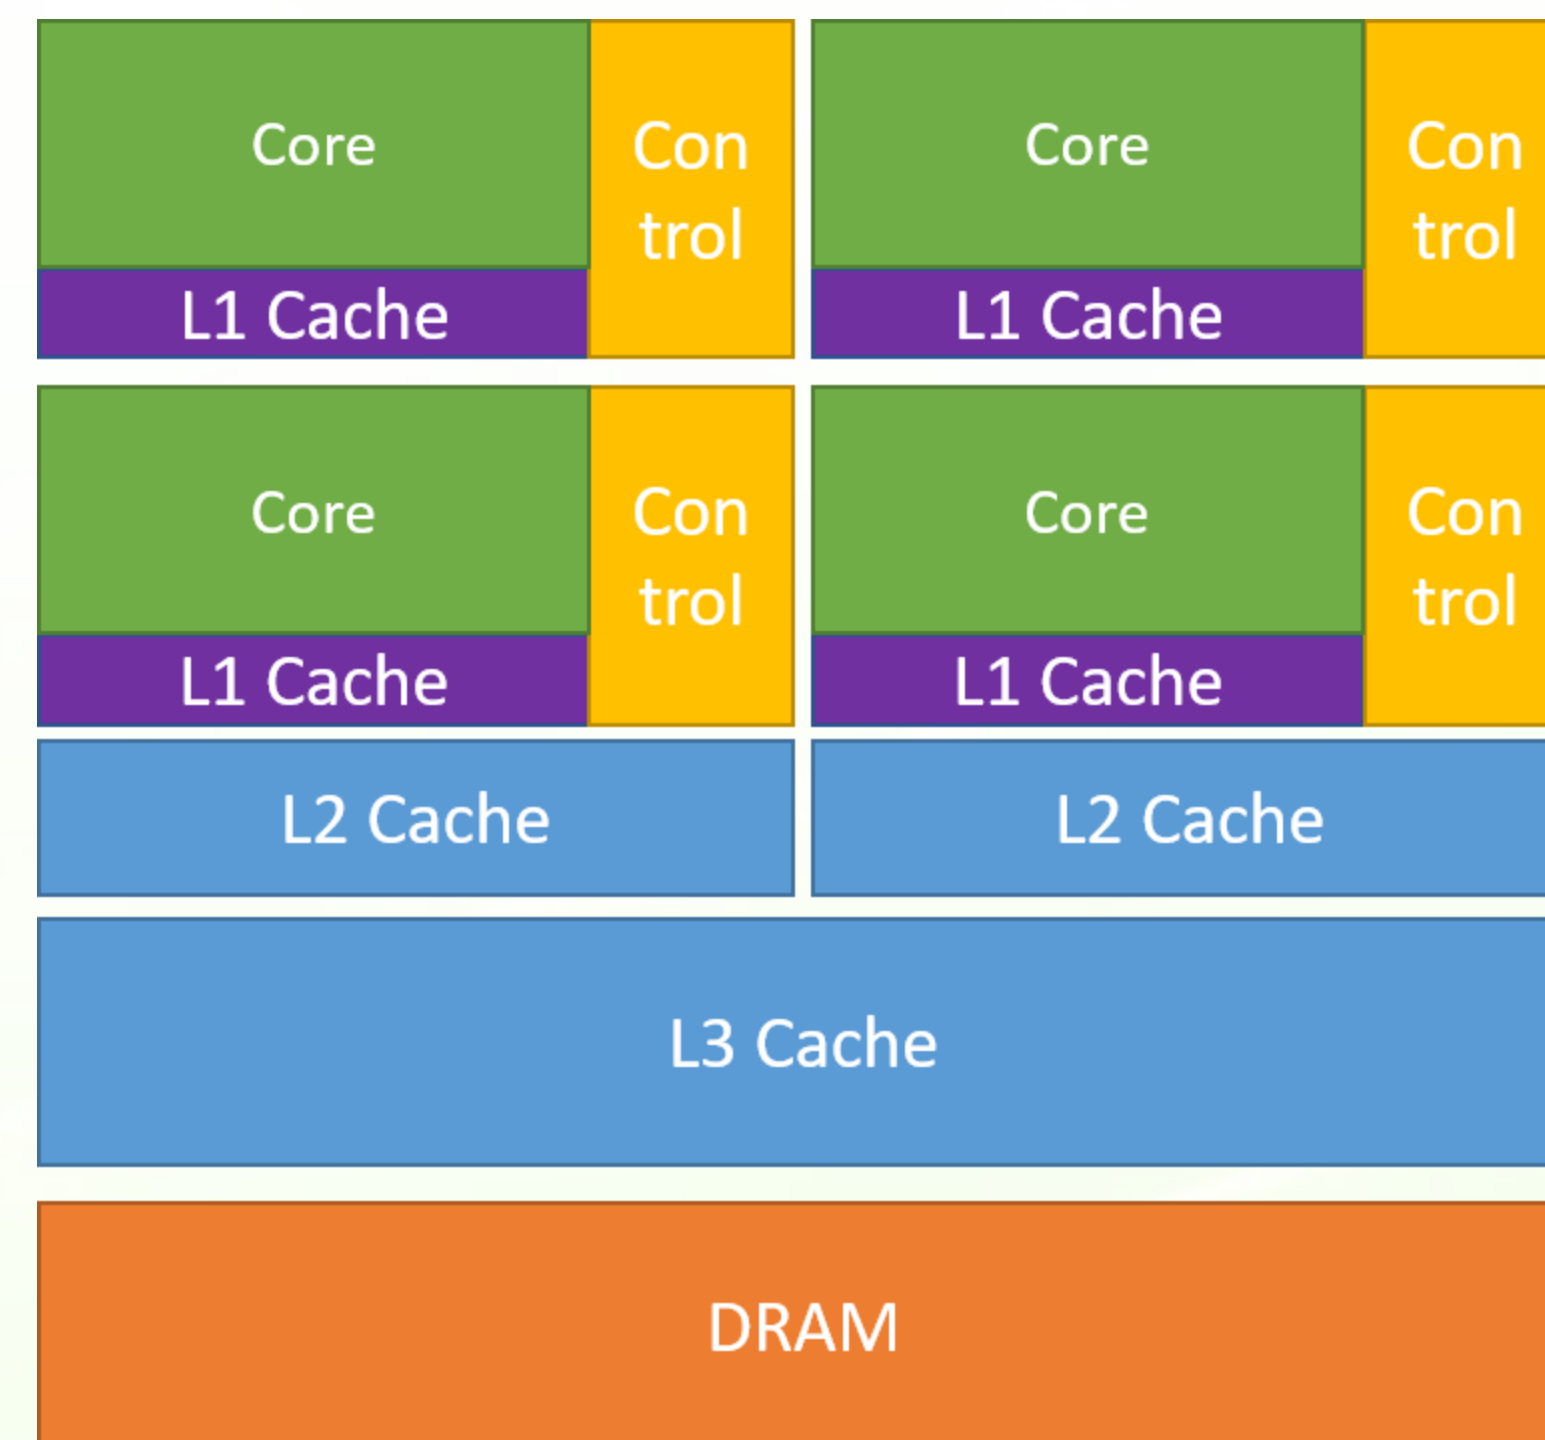
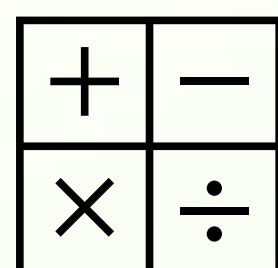
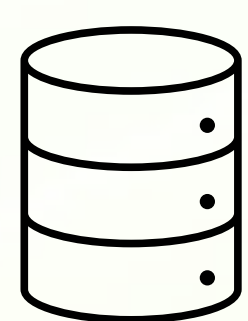
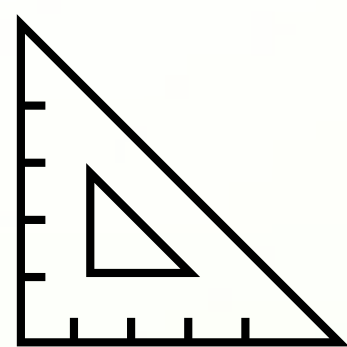
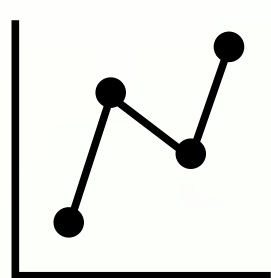
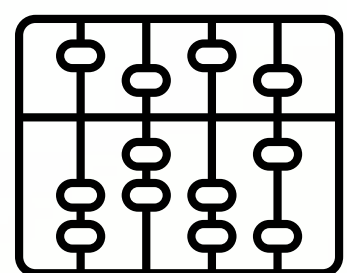
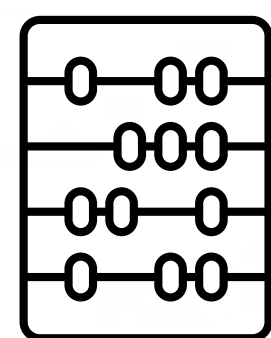
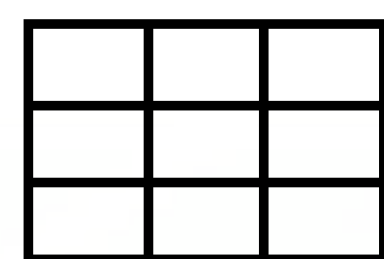


CPU



GPU

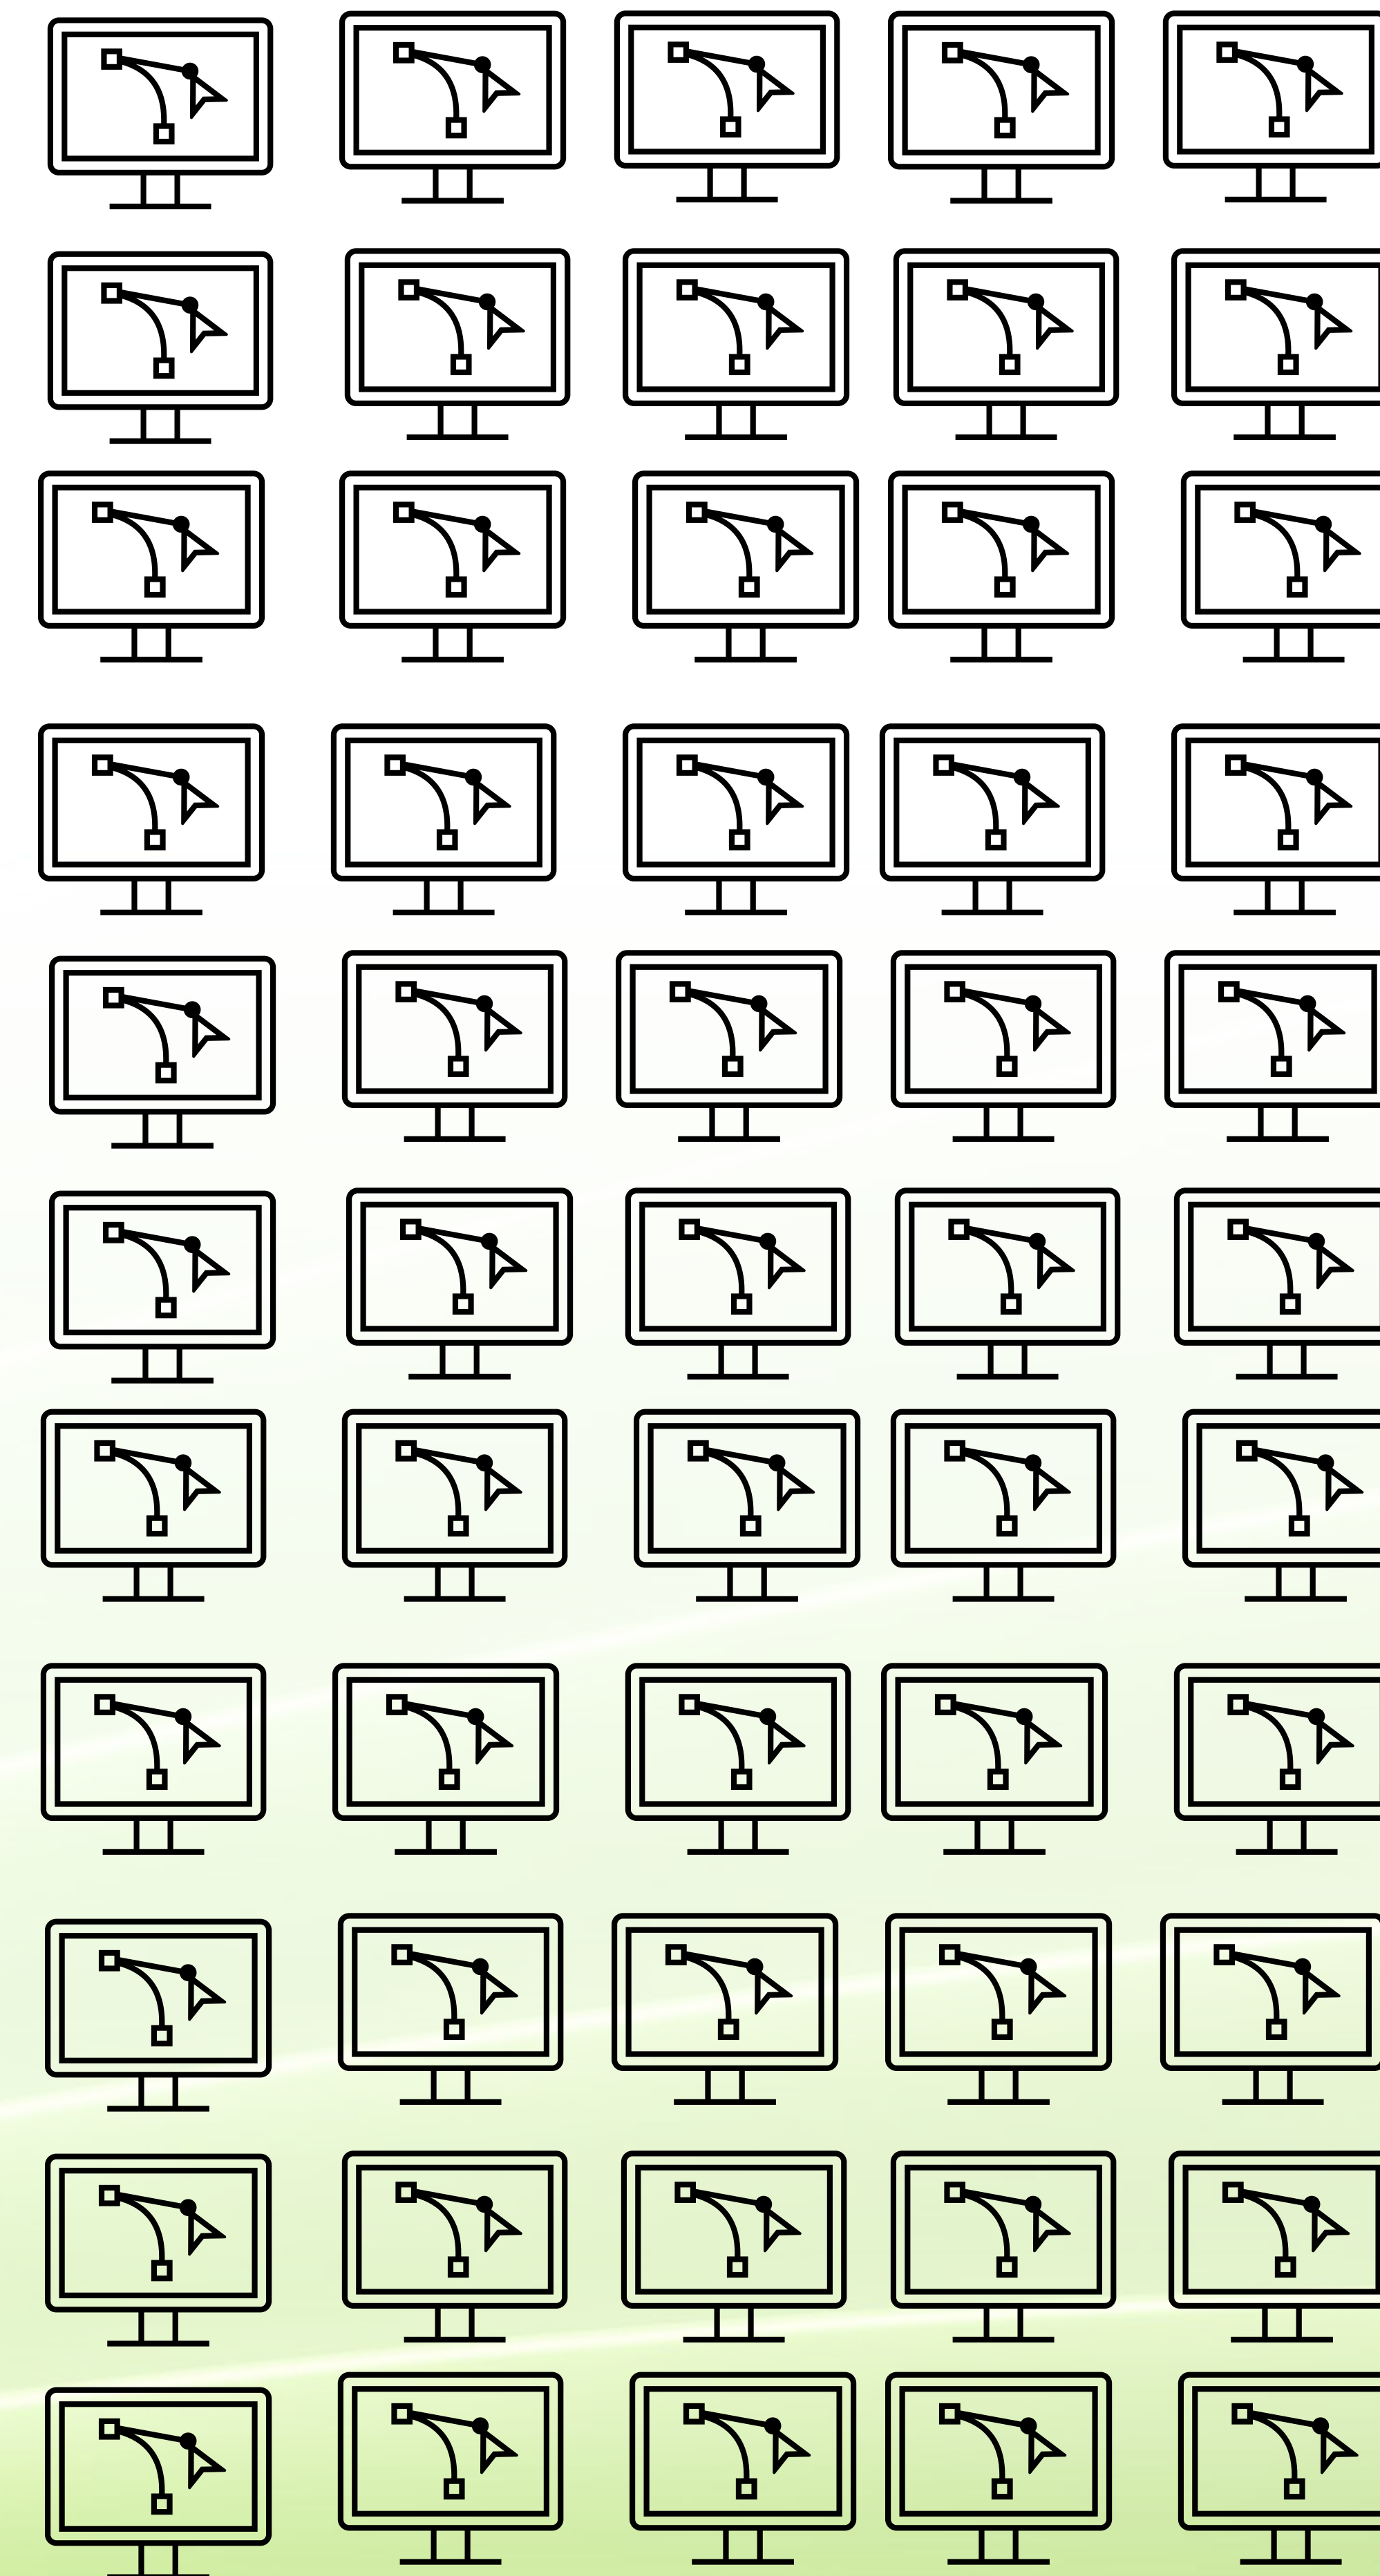
Accelerated Supercomputing



CPU

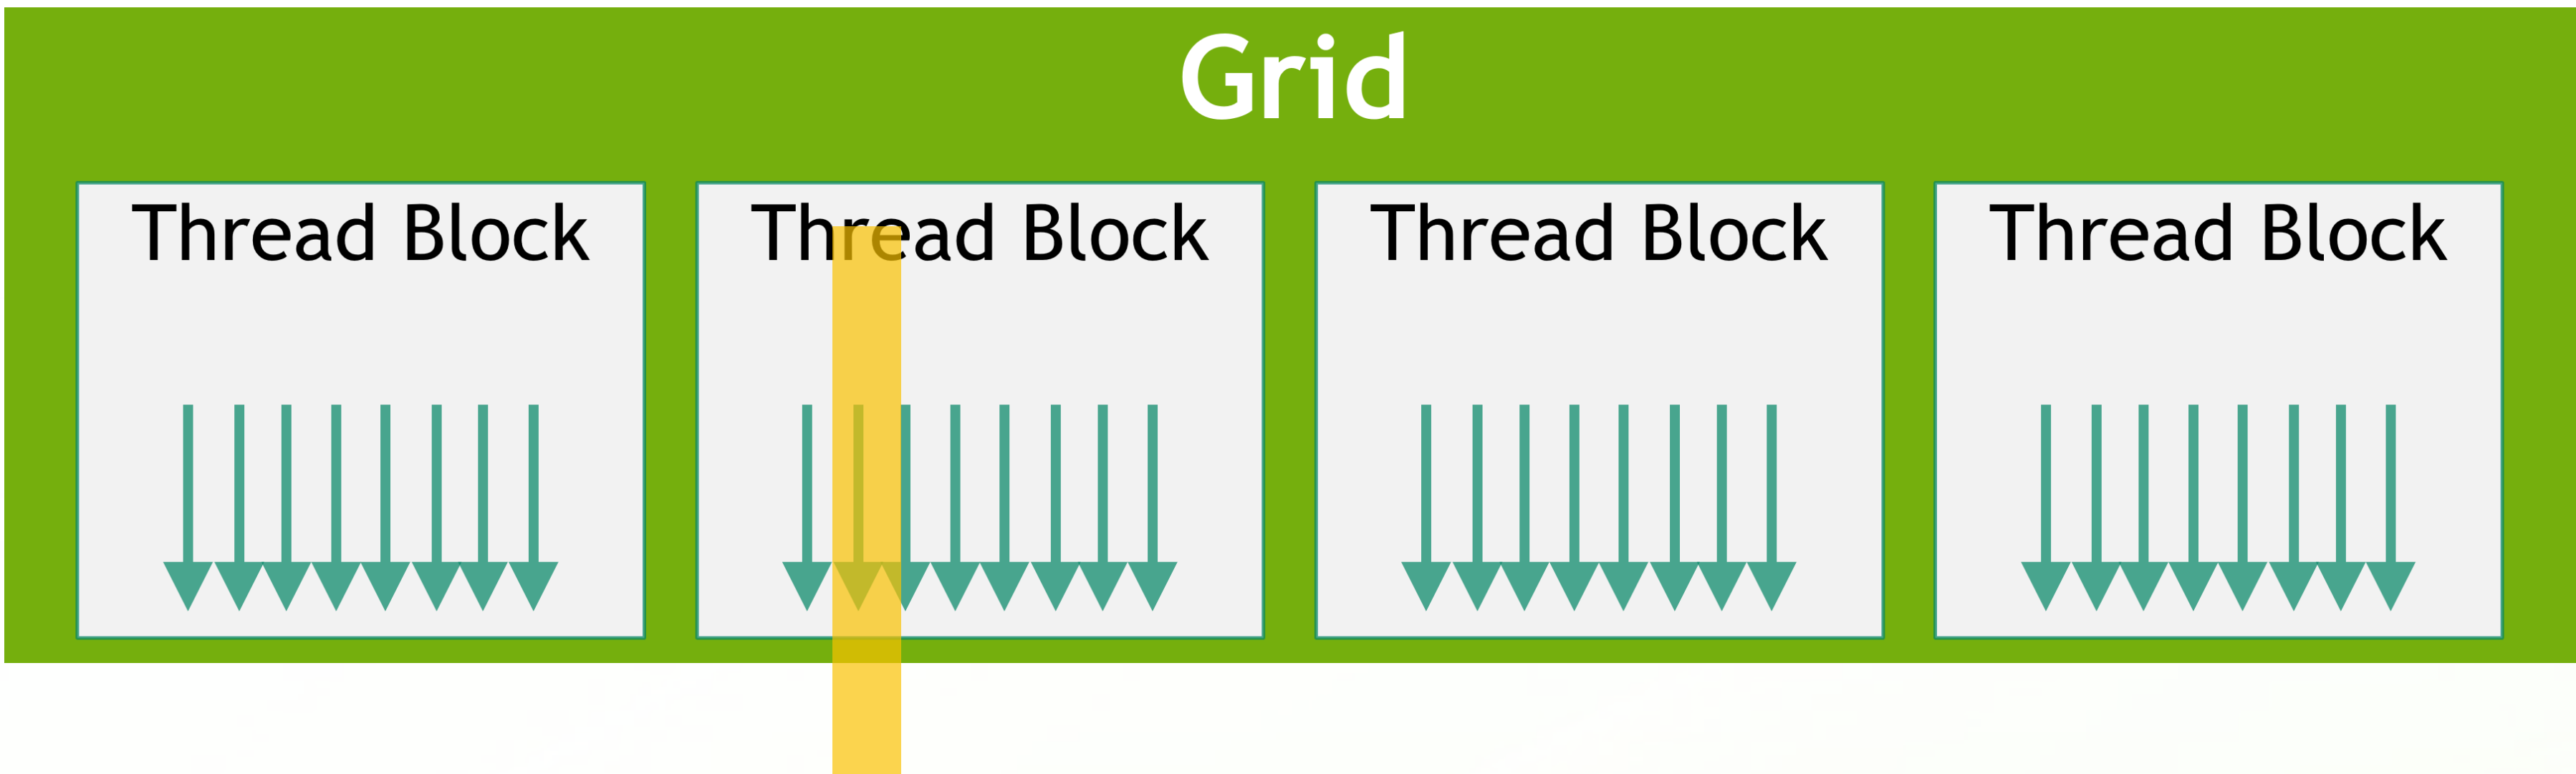


GPU



Matrix Multiplication in Parallel on a GPU

$A \times B = C$



$$\begin{pmatrix} 1 & -1 & 1 & -1 & 1 & -1 & 1 & -1 \\ -1 & 1 & -1 & 1 & -1 & 1 & -1 & 1 \\ 1 & -1 & 1 & -1 & 1 & -1 & 1 & -1 \\ -1 & 1 & -1 & 1 & -1 & 1 & -1 & 1 \end{pmatrix}$$

$$\begin{pmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \\ 4 & 3 & 2 & 1 \\ 8 & 7 & 6 & 5 \\ 12 & 11 & 10 & 9 \\ 1 & 1 & 1 & 1 \\ 0 & 1 & 0 & 1 \end{pmatrix}$$

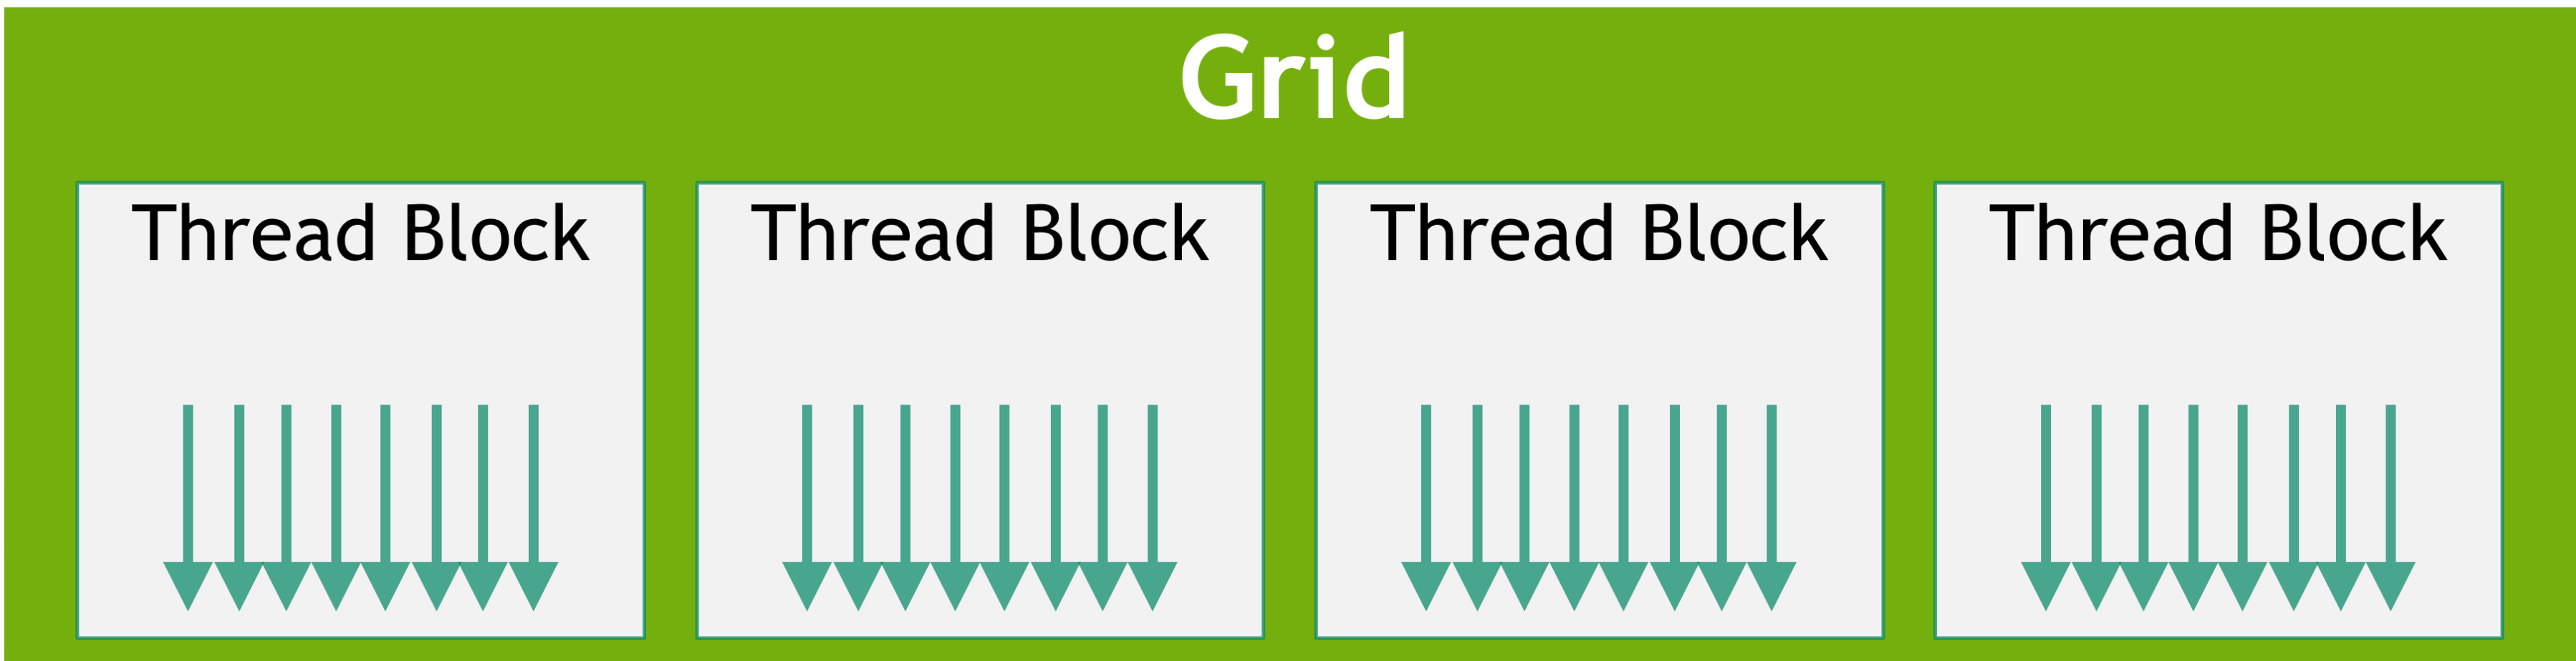
$$\begin{pmatrix} \square & \square & \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square & \square & \square \end{pmatrix}$$

Kernel = instruction for each thread to follow

Kernel for matrix multiplication:
compute the dot product of
an assigned row in A with an assigned column of B

Matrix Multiplication in Parallel on a GPU

$$A \times B = C$$

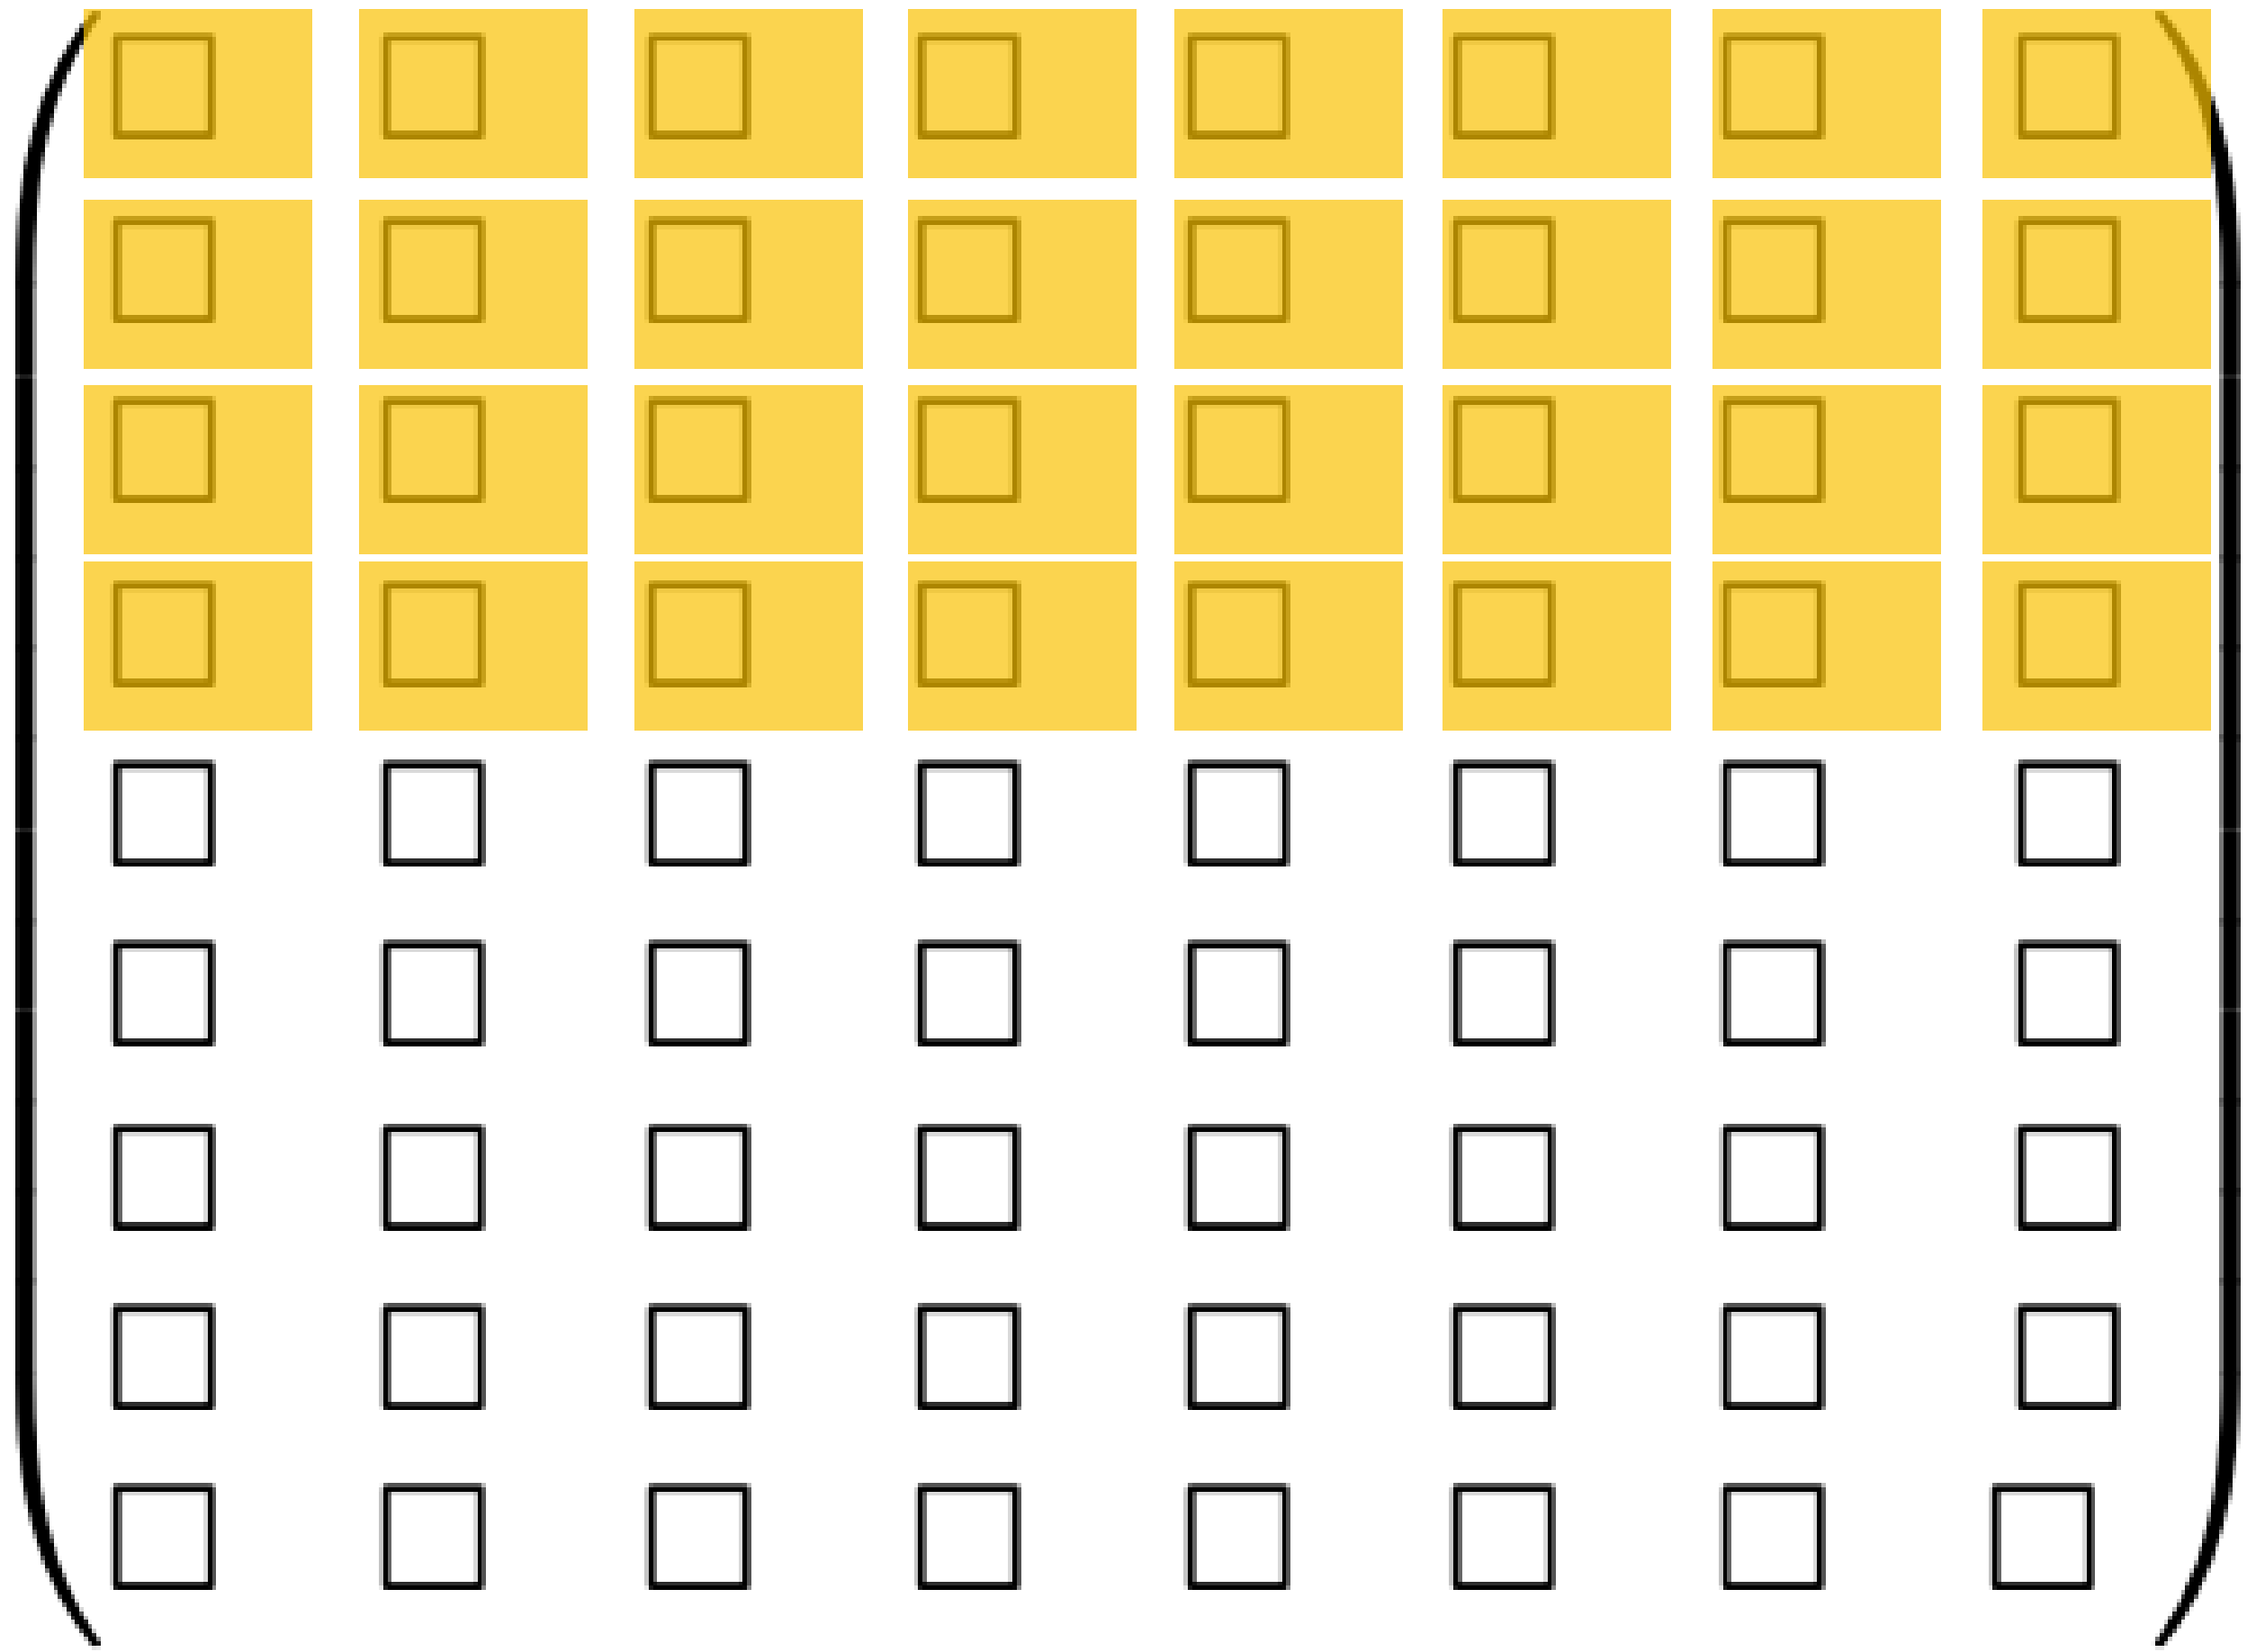


$$\begin{pmatrix} 1 & -1 & 1 & -1 & 1 & -1 & 1 & -1 \\ -1 & 1 & -1 & 1 & -1 & 1 & -1 & 1 \\ 1 & -1 & 1 & -1 & 1 & -1 & 1 & -1 \\ -1 & 1 & -1 & 1 & -1 & 1 & -1 & 1 \end{pmatrix}$$

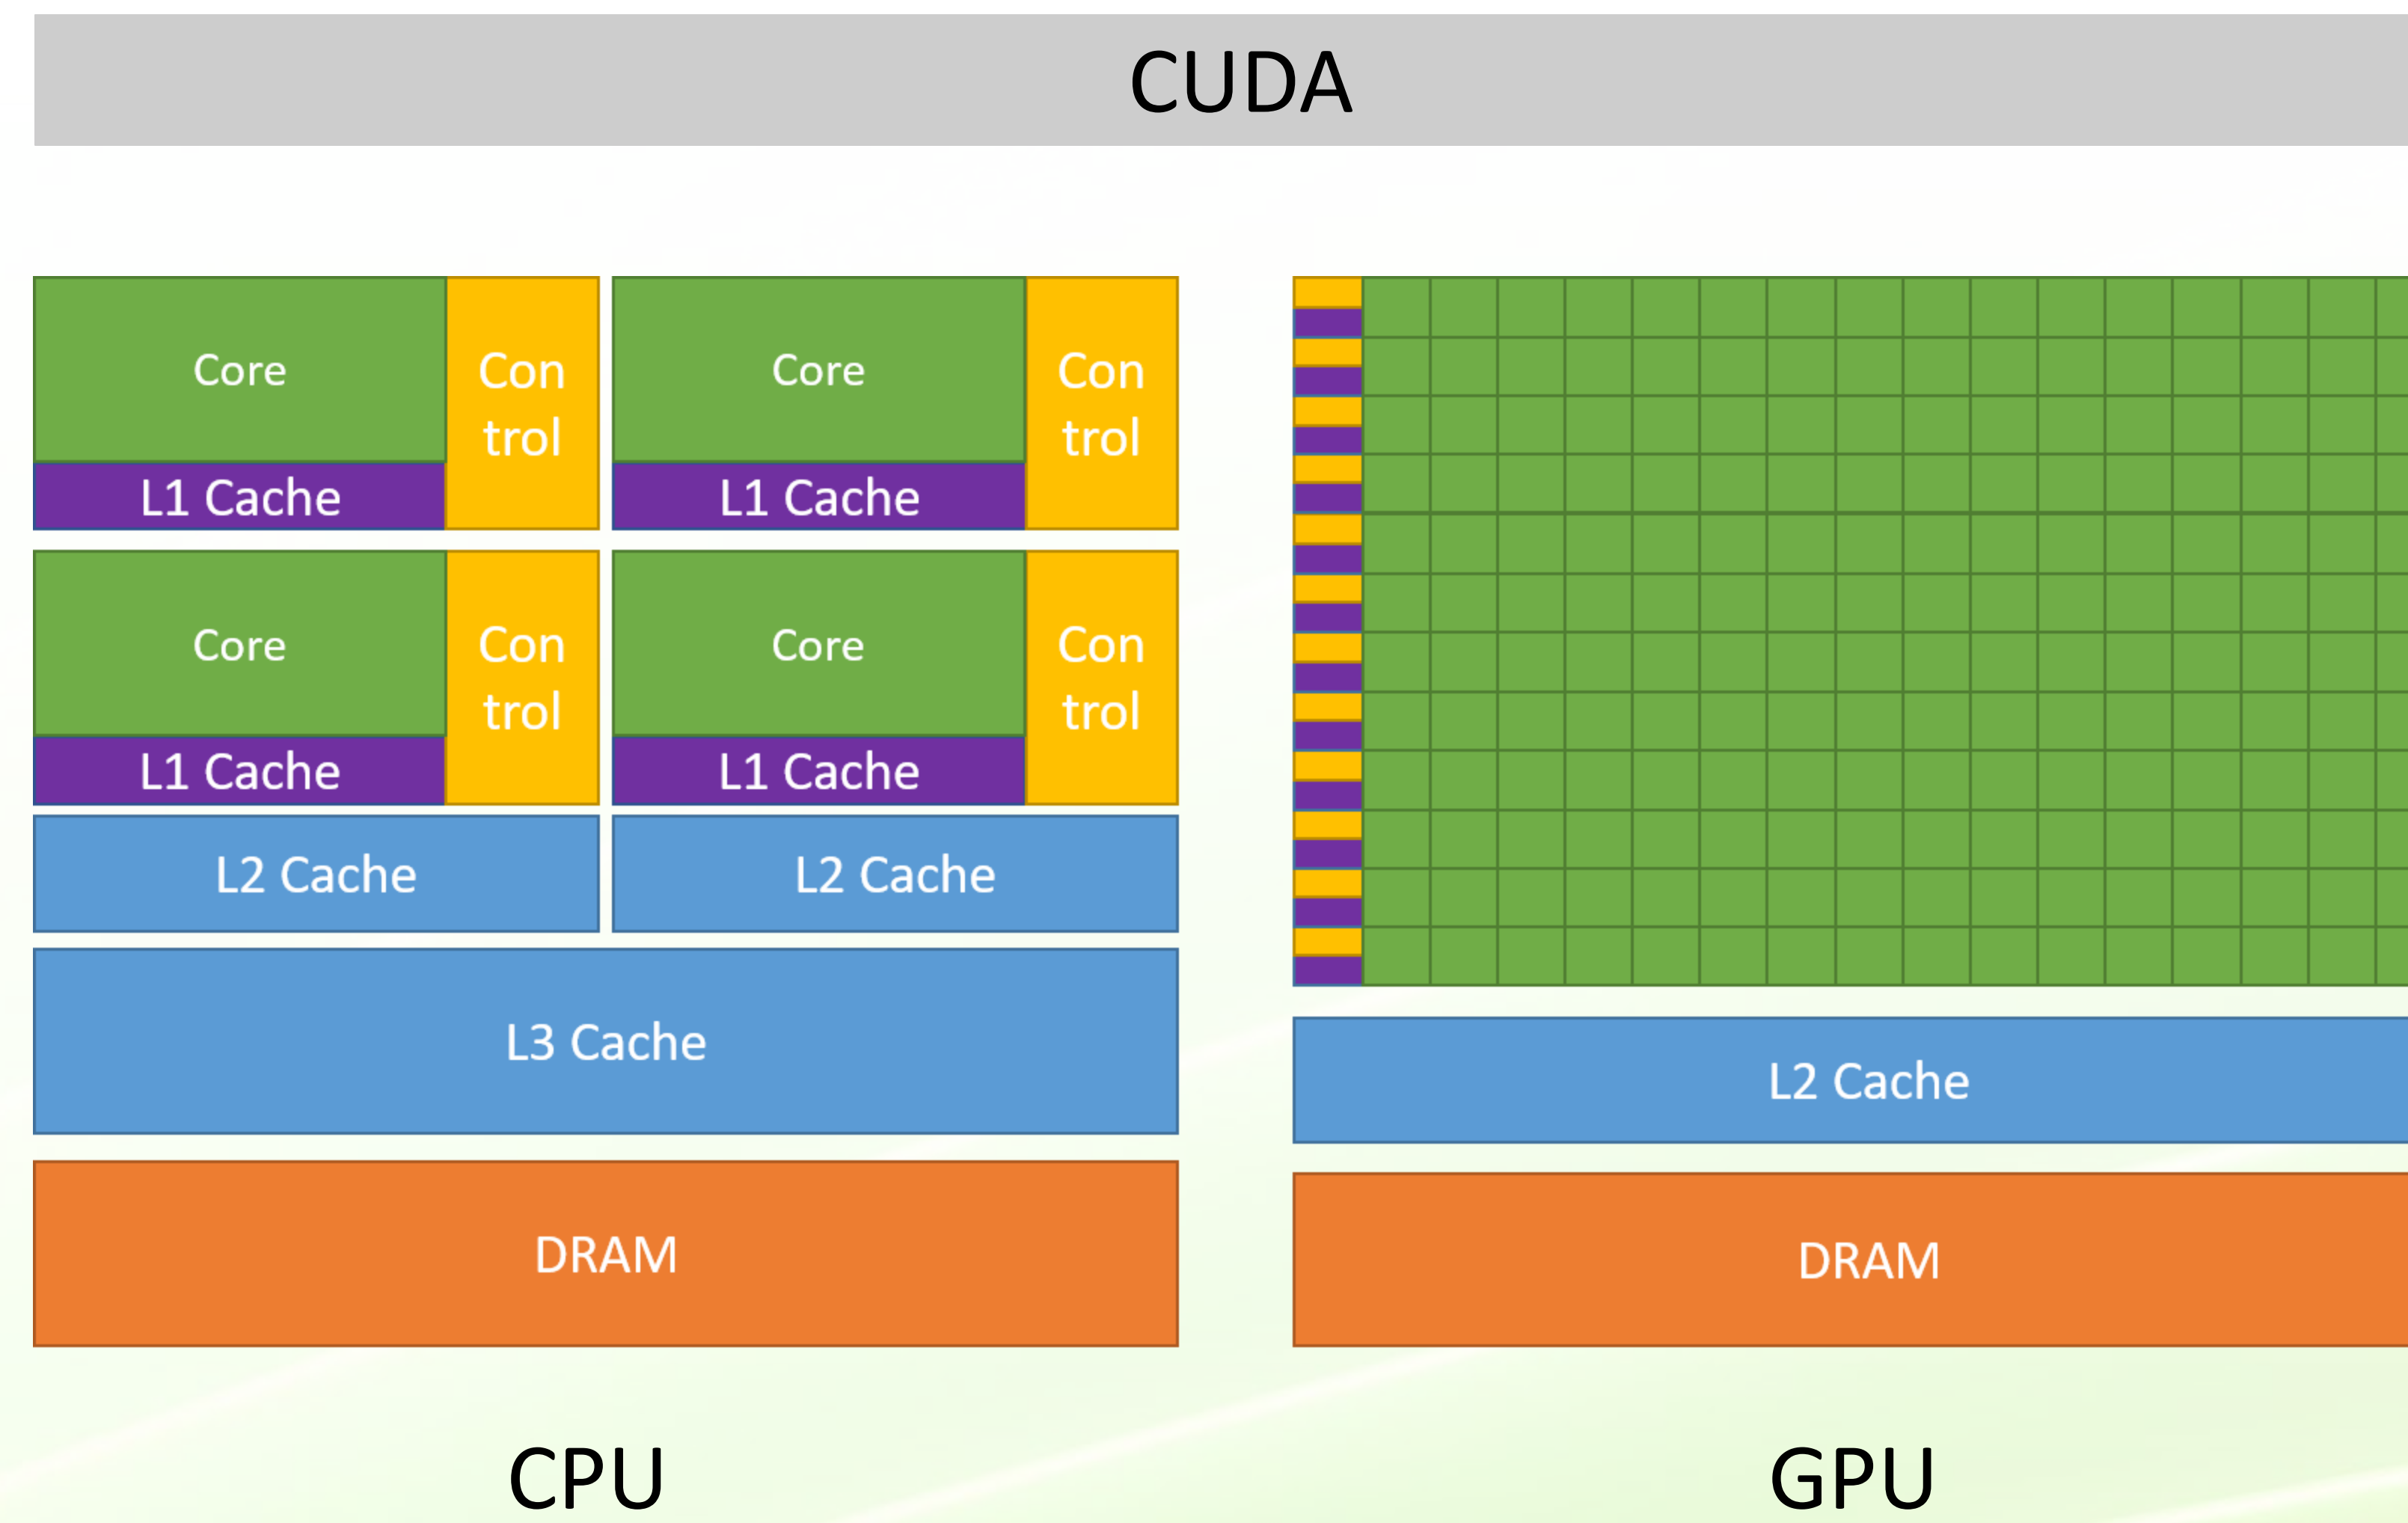
Kernel = instruction for each thread to follow

*Kernel for matrix multiplication:
compute the dot product of
an assigned row in A with an assigned column of B*

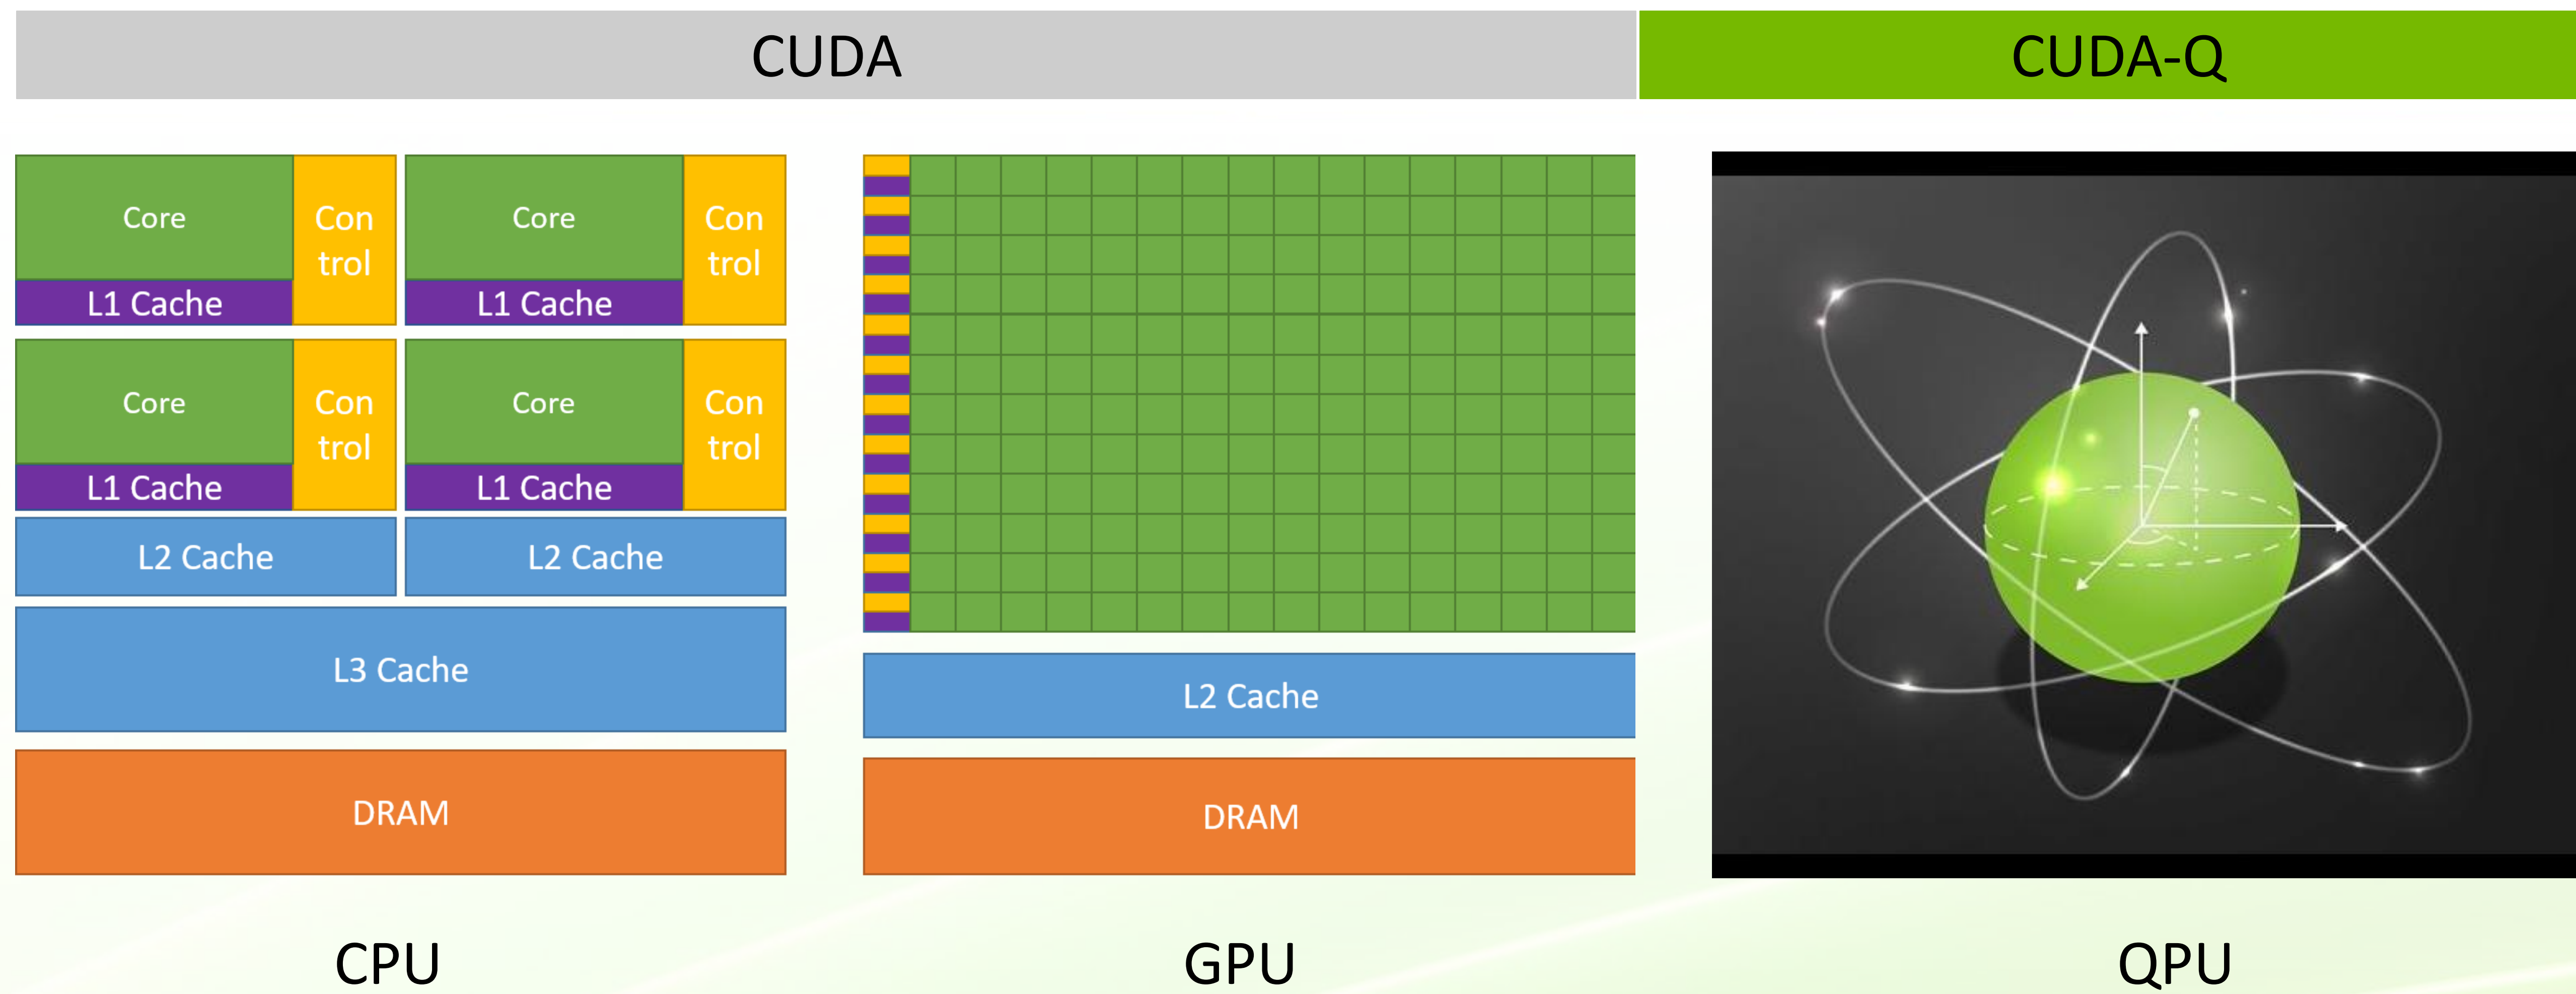
$$\begin{pmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \\ 4 & 3 & 2 & 1 \\ 8 & 7 & 6 & 5 \\ 12 & 11 & 10 & 9 \\ 1 & 1 & 1 & 1 \\ 0 & 1 & 0 & 1 \end{pmatrix}$$



Accelerated Supercomputing



Quantum Accelerated Supercomputing

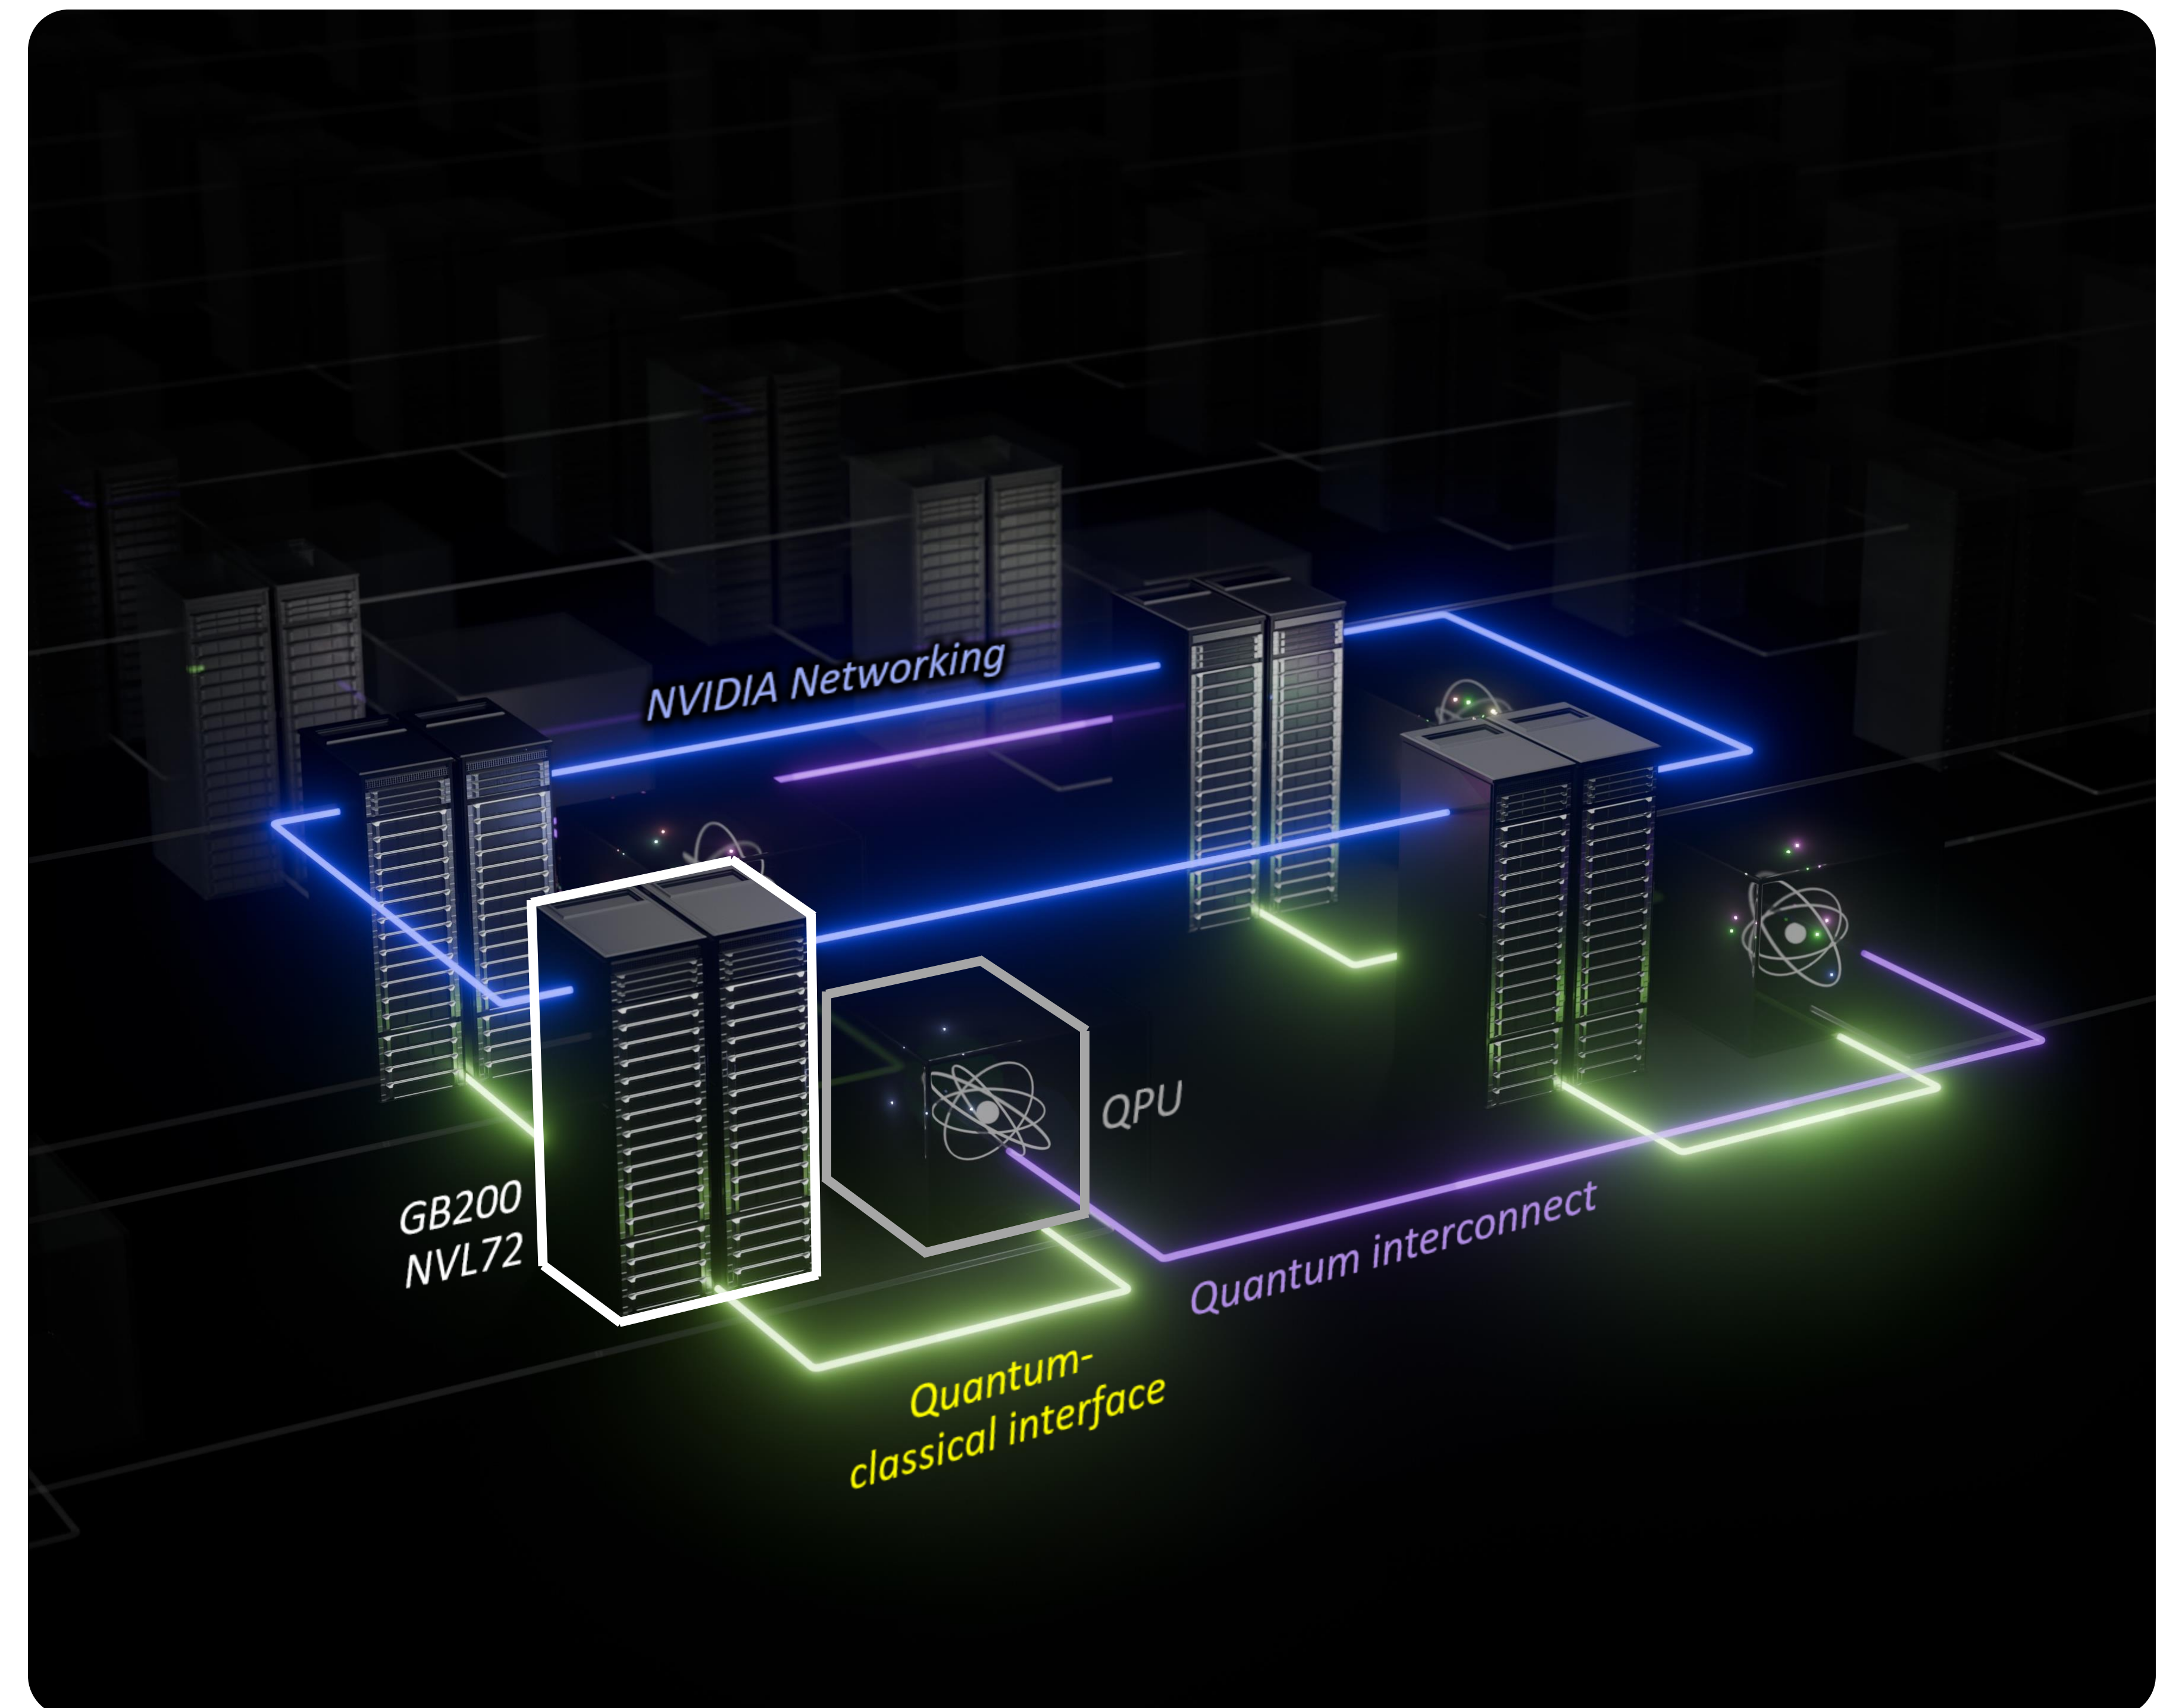


Tomorrow's Accelerated Quantum Supercomputers are GPU Supercomputers

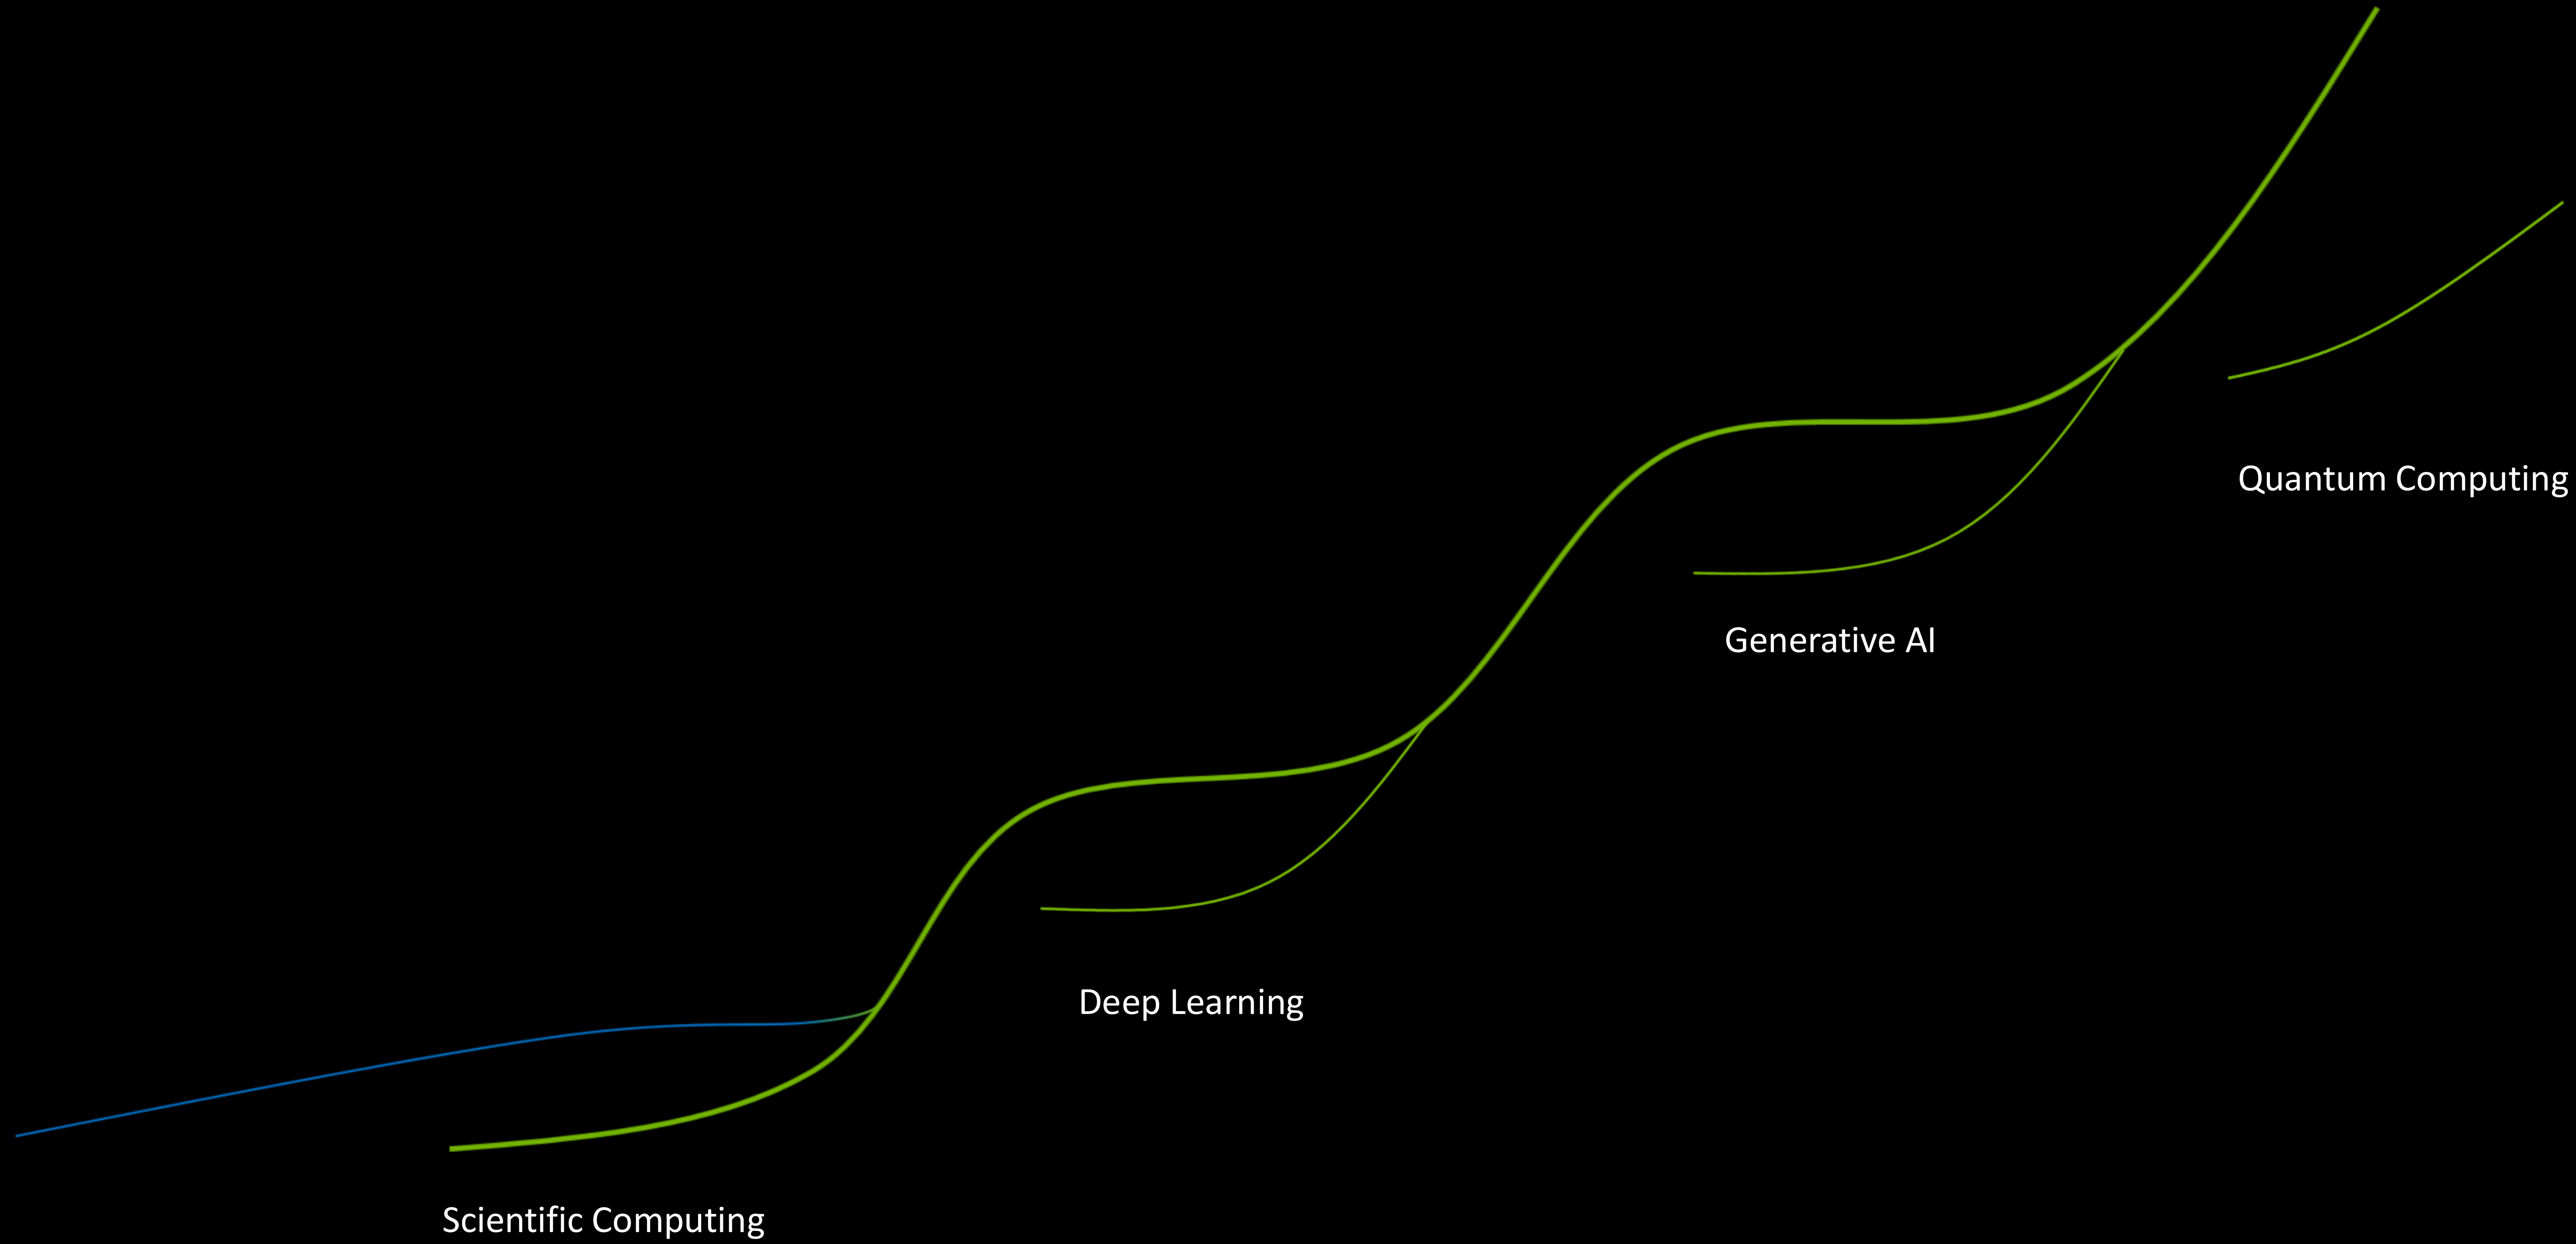
Accelerated Quantum Supercomputer

A hybrid quantum-classical device that uses GPU-supercomputing to turn qubit technology into a computer able to run useful applications

- Useful quantum computers are mostly an **AI supercomputer**
- NV supercomputers are QPU-agnostic
- **Hybrid applications** use CPUs, GPUs and QPUS
- AI supercomputing needed to **control and operate** QPU hardware

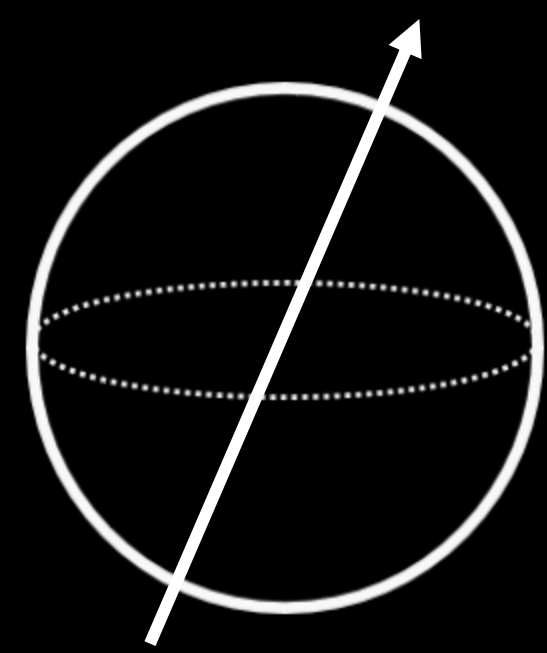


Accelerated Computing



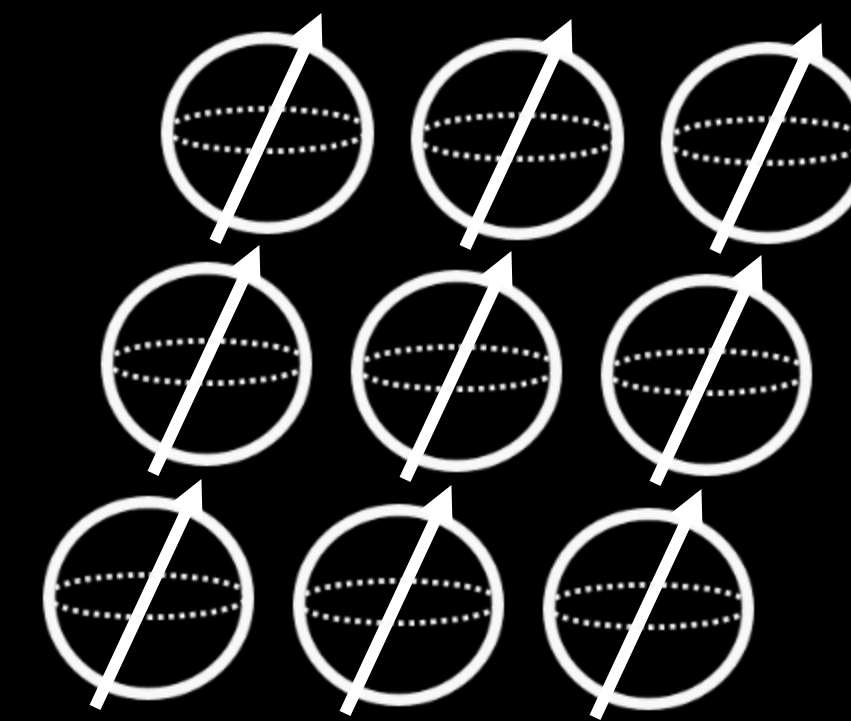
Quantum Challenges

What's Standing Between Today and Useful Quantum Computing?



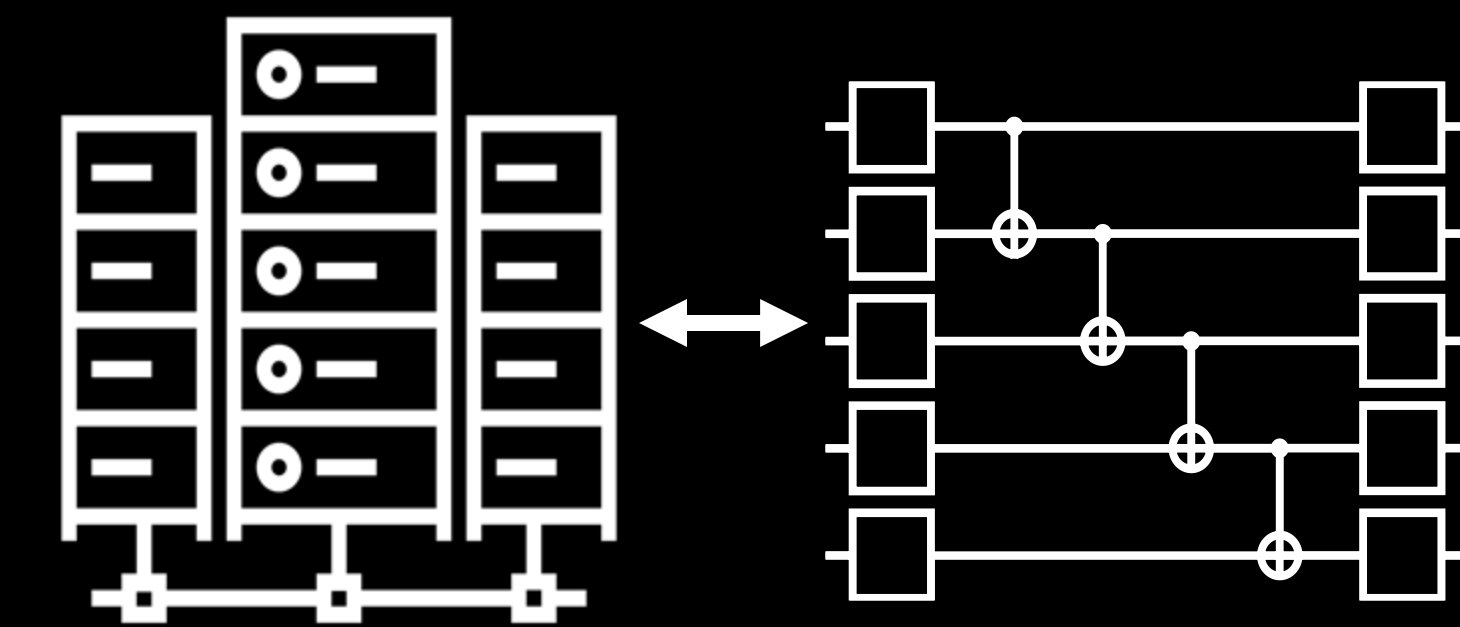
Qubit Fidelity

99.99% 2-Qubit Gate Fidelity



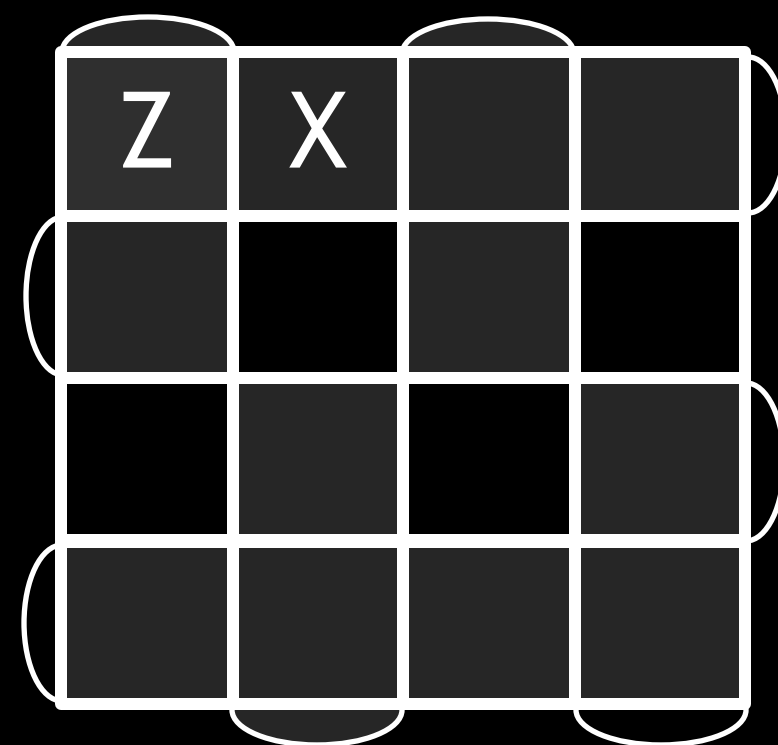
Qubit Scale

100k-1M+ Qubits for FTQC



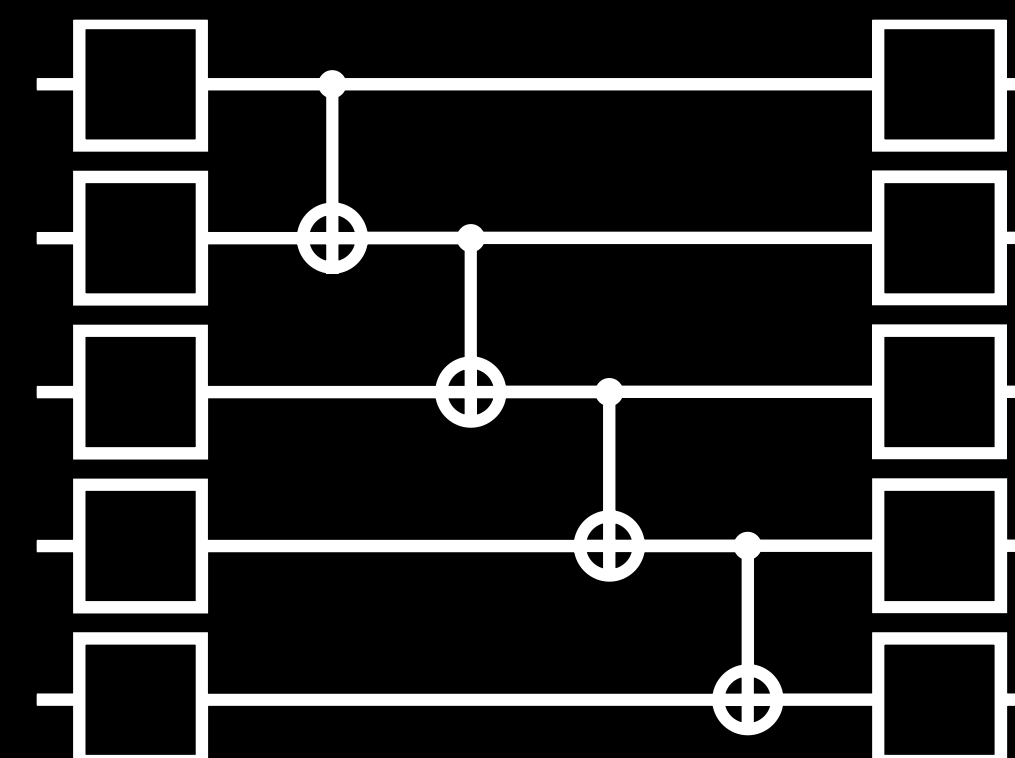
HPC Integration

Sub-Microsecond HPC-QC Latency



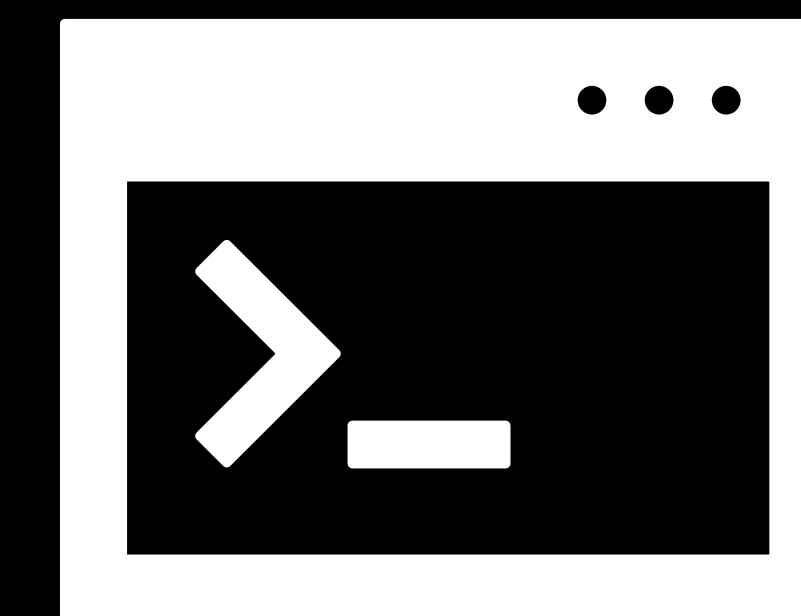
Error Correction

Methods that Scale to Large Quantum Systems



Algorithms

Algorithms with Exponential Speed-up

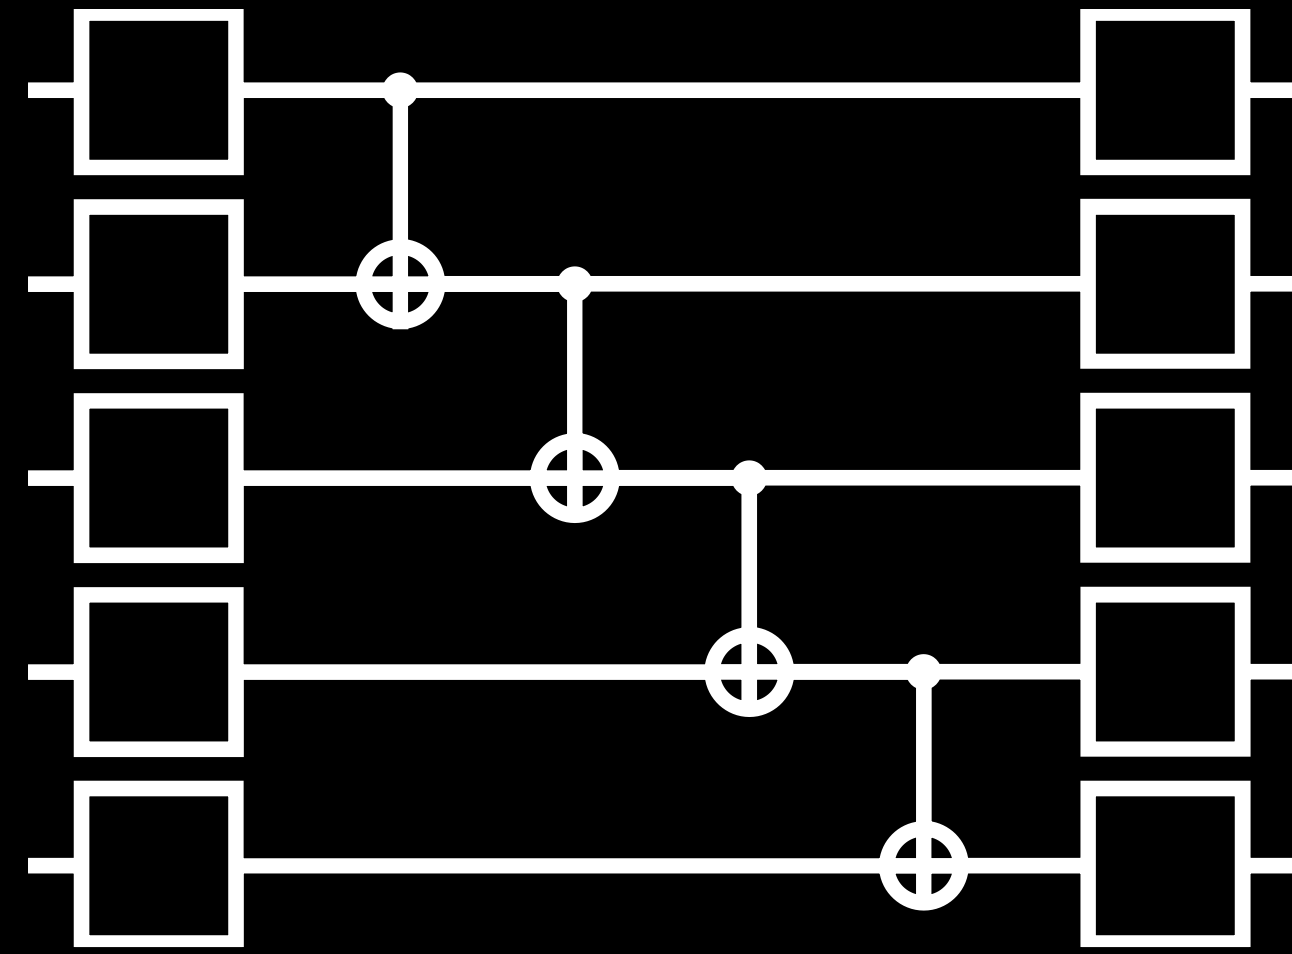


Developer Tools

Integrate with Scientific Computing
Familiar to non-Quantum Physicists

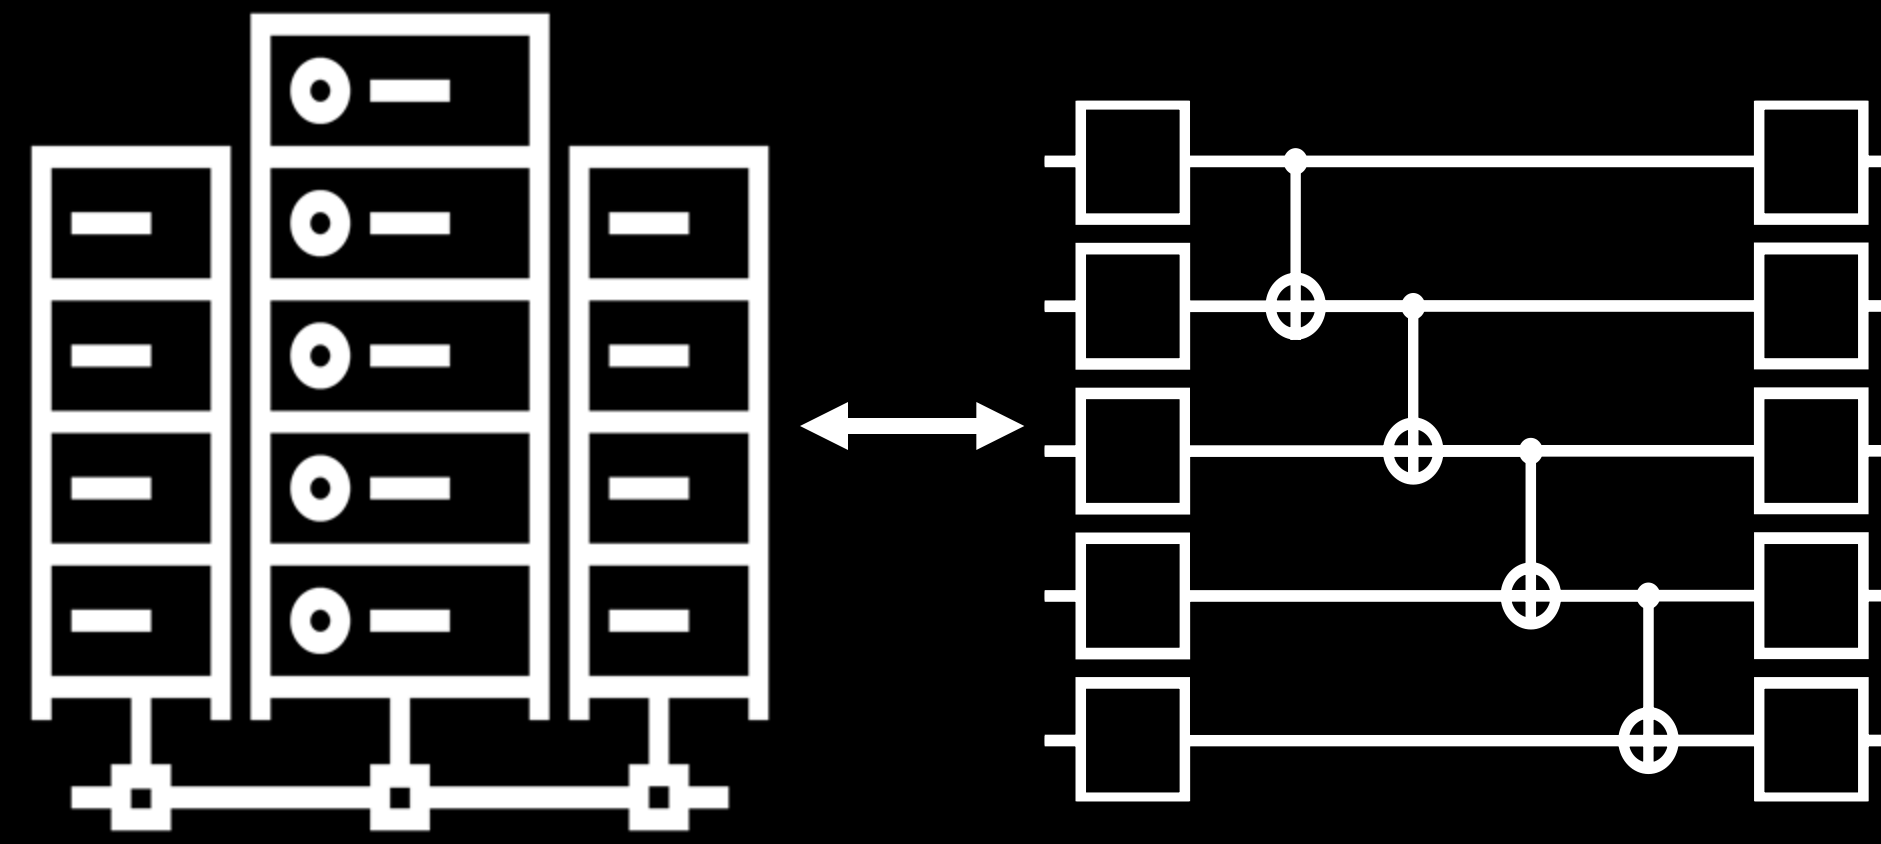
NVIDIA Quantum

Powering the Global Quantum Computing Community



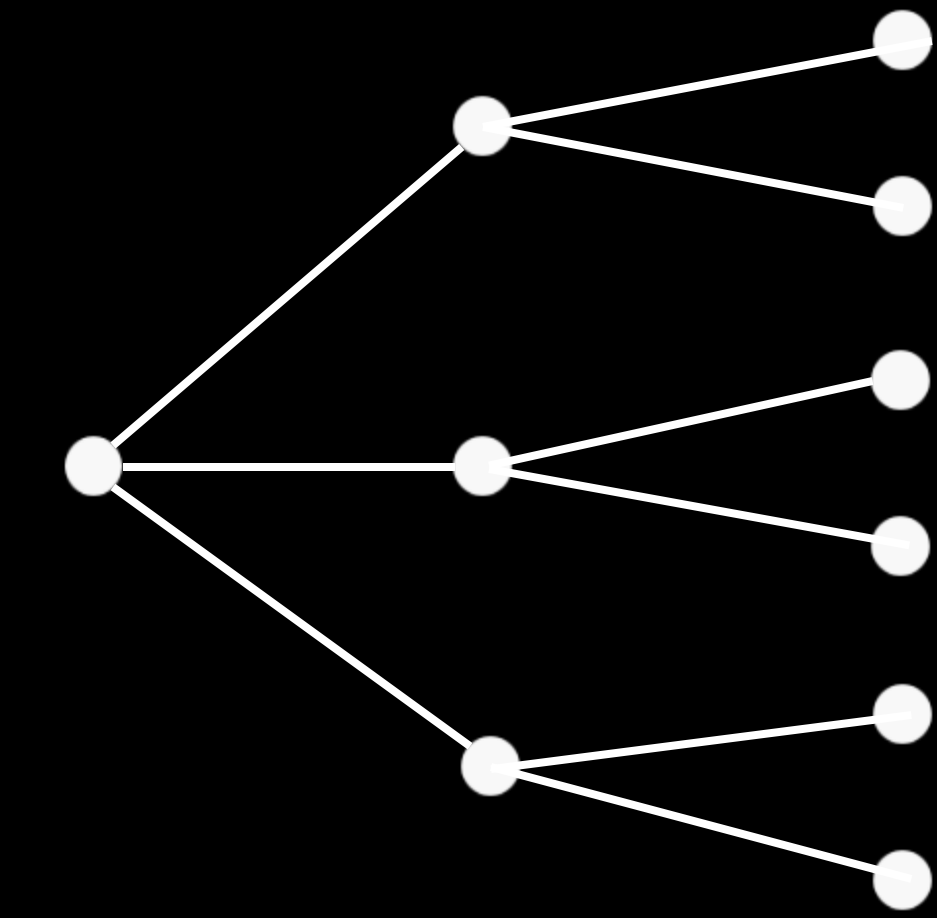
Simulation

Algorithm Design, Resource Estimation, QPU Design



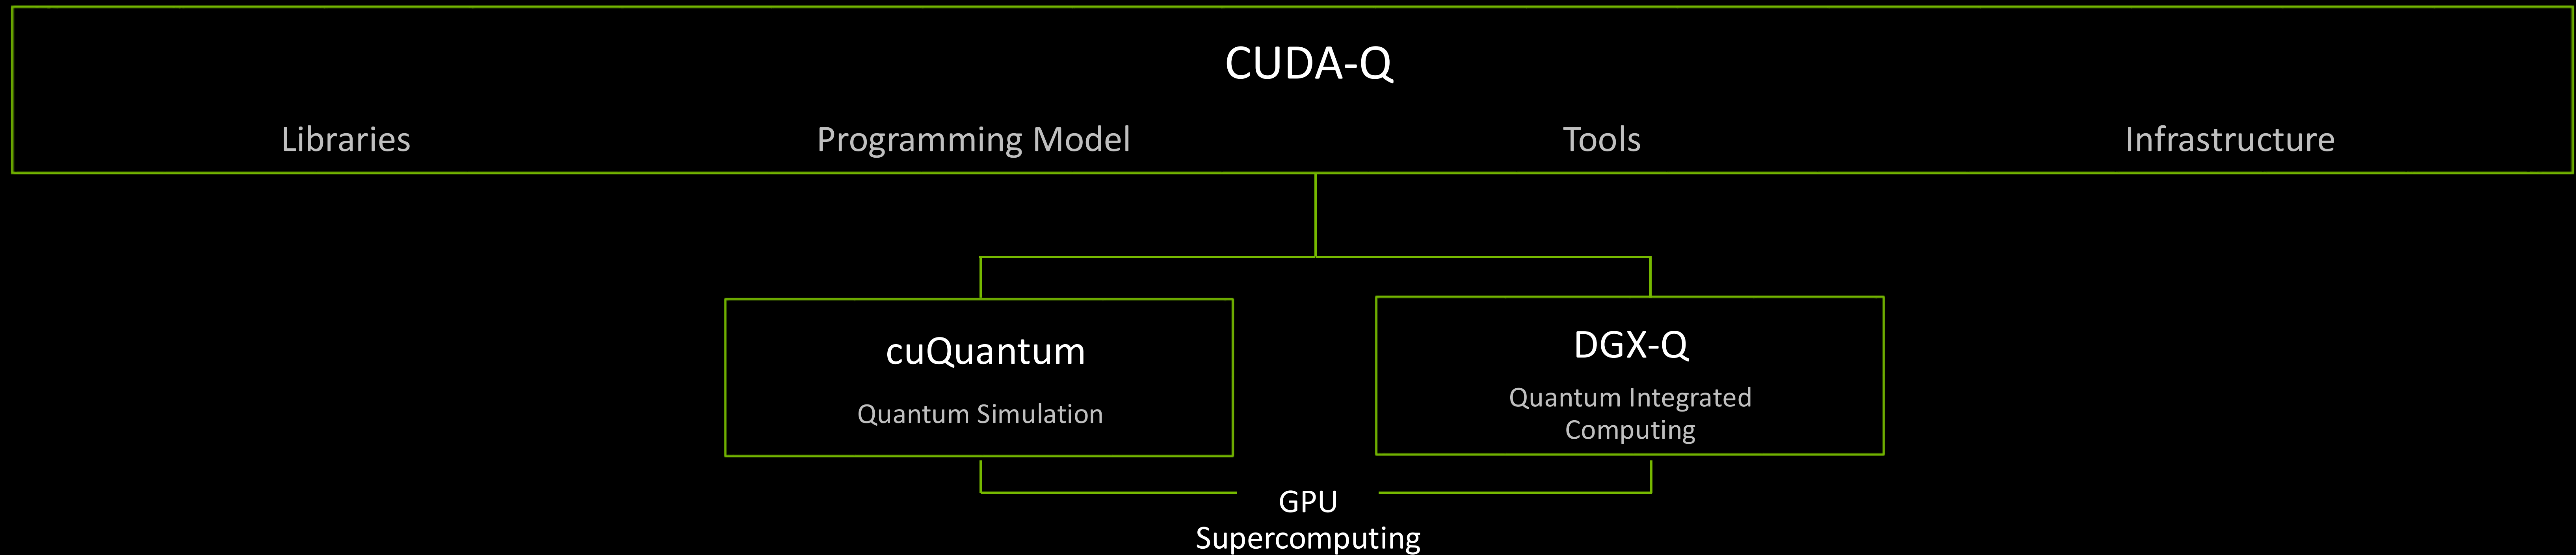
HPC Quantum Integration

Integrated Applications, QEC, Sub-Microsecond Latency



AI for Quantum

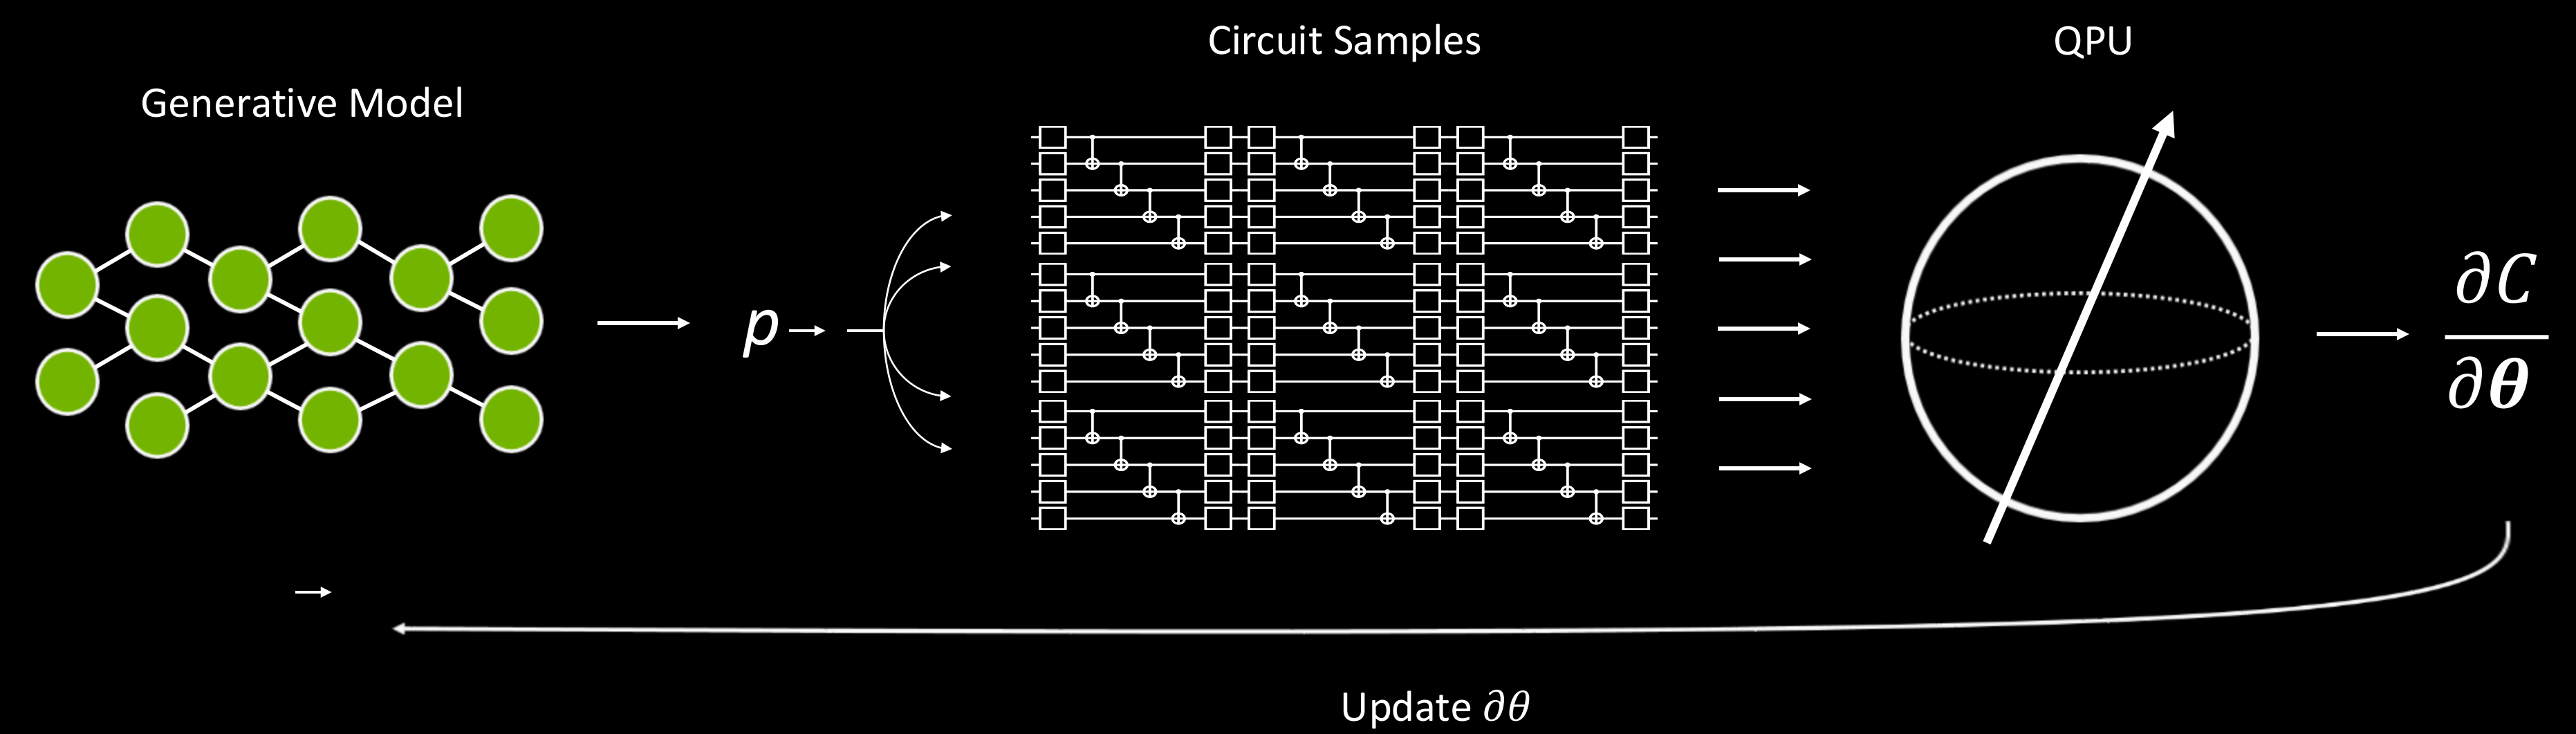
QEC, Calibration, Algorithms



Generative AI + Quantum Algorithms

University of Toronto, St Jude's, and NVIDIA partner to invent GPT-QE

- Generative Pre-Trained Transformer-based (GPT) method for computing the ground state energies
- First GPT-generated quantum circuit
- Run via CUDA-Q on NERSC Perlmutter





Agenda

- What is Quantum Accelerated Supercomputing

- Useful Quantum Simulation

- How-to Guide to CUDA-Q

- Distributed Quantum Computing

- Conclusion

The Case for Quantum Computing Simulation

Quantum research is limited by access

Availability

~500k

Quantum Developers

~50

Publicly available QPUs

Uptime

10-20%

Typical uptime for deployed QPU

Scale

0

Fault-Tolerant Qubits Available

$O(100)$ - $O(1000)$

Needed for useful applications

Accuracy

19

Circuit Depth, best case

3-4

Circuit Depth, typical

Iteration Time

40Q sim

1 hour

GPU Cluster

7.5 years

CPU

Integration

4 million

CUDA Developers worldwide

Quantum Simulation on a GPU

The screenshot shows a code editor with two Python files side-by-side. The left file, `QML_cpu_demo.py`, sets the target to `'qpp-cpu'`. The right file, `QML_nvqc_demo.py`, sets the target to `'nvqc'`. Both files perform a quantum simulation with a kernel, qubits, and parameters, and print the expectation value and runtime. The editor interface includes a terminal at the bottom with the prompt `cudaq@221def261680:~$`. The status bar at the bottom indicates the current position is `Ln 53, Col 25`.

```
File Edit Selection View Go Run ... < > cudaq [Container nvcr.io/nvidia/nightly/cuda-quantum:0.7.0 (vibrant_ramanujan)]
```

```
QML_cpu_demo.py x QML_nvqc_demo.py x
```

```
QML_cpu_demo.py QML_nvqc_demo.py
```

```
44
45
46 for q1, q2 in zip(qubits_list[0::2], qubits_list[1::2]):
47     kernel.cz(qubits[q1], qubits[q2])
48
49 kernel.x(qubits)
50 kernel.y(qubits)
51 kernel.h(qubits)
52
53 cudaq.set_target('qpp-cpu')
54
55 start = timeit.default_timer()
56
57 exp_vals = cudaq.observe(kernel, h, parameters)
58
59 end = timeit.default_timer()
60
61 print("Expectation Value: ", exp_vals[0].expectation())
62 print("Runtime:", end - start)
63
```

```
44
45
46 for q1, q2 in zip(qubits_list[0::2], qubits_list[1::2]):
47     kernel.cz(qubits[q1], qubits[q2])
48
49 kernel.x(qubits)
50 kernel.y(qubits)
51 kernel.h(qubits)
52
53 cudaq.set_target('nvqc')
54
55 start = timeit.default_timer()
56
57 exp_vals = cudaq.observe(kernel, h, parameters)
58
59 end = timeit.default_timer()
60
61 print("Expectation Value: ", exp_vals[0].expectation())
62 print("Runtime:", end - start)
63
```

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
```

```
o cudaq@221def261680:~$
```

```
o cudaq@221def261680:~$
```

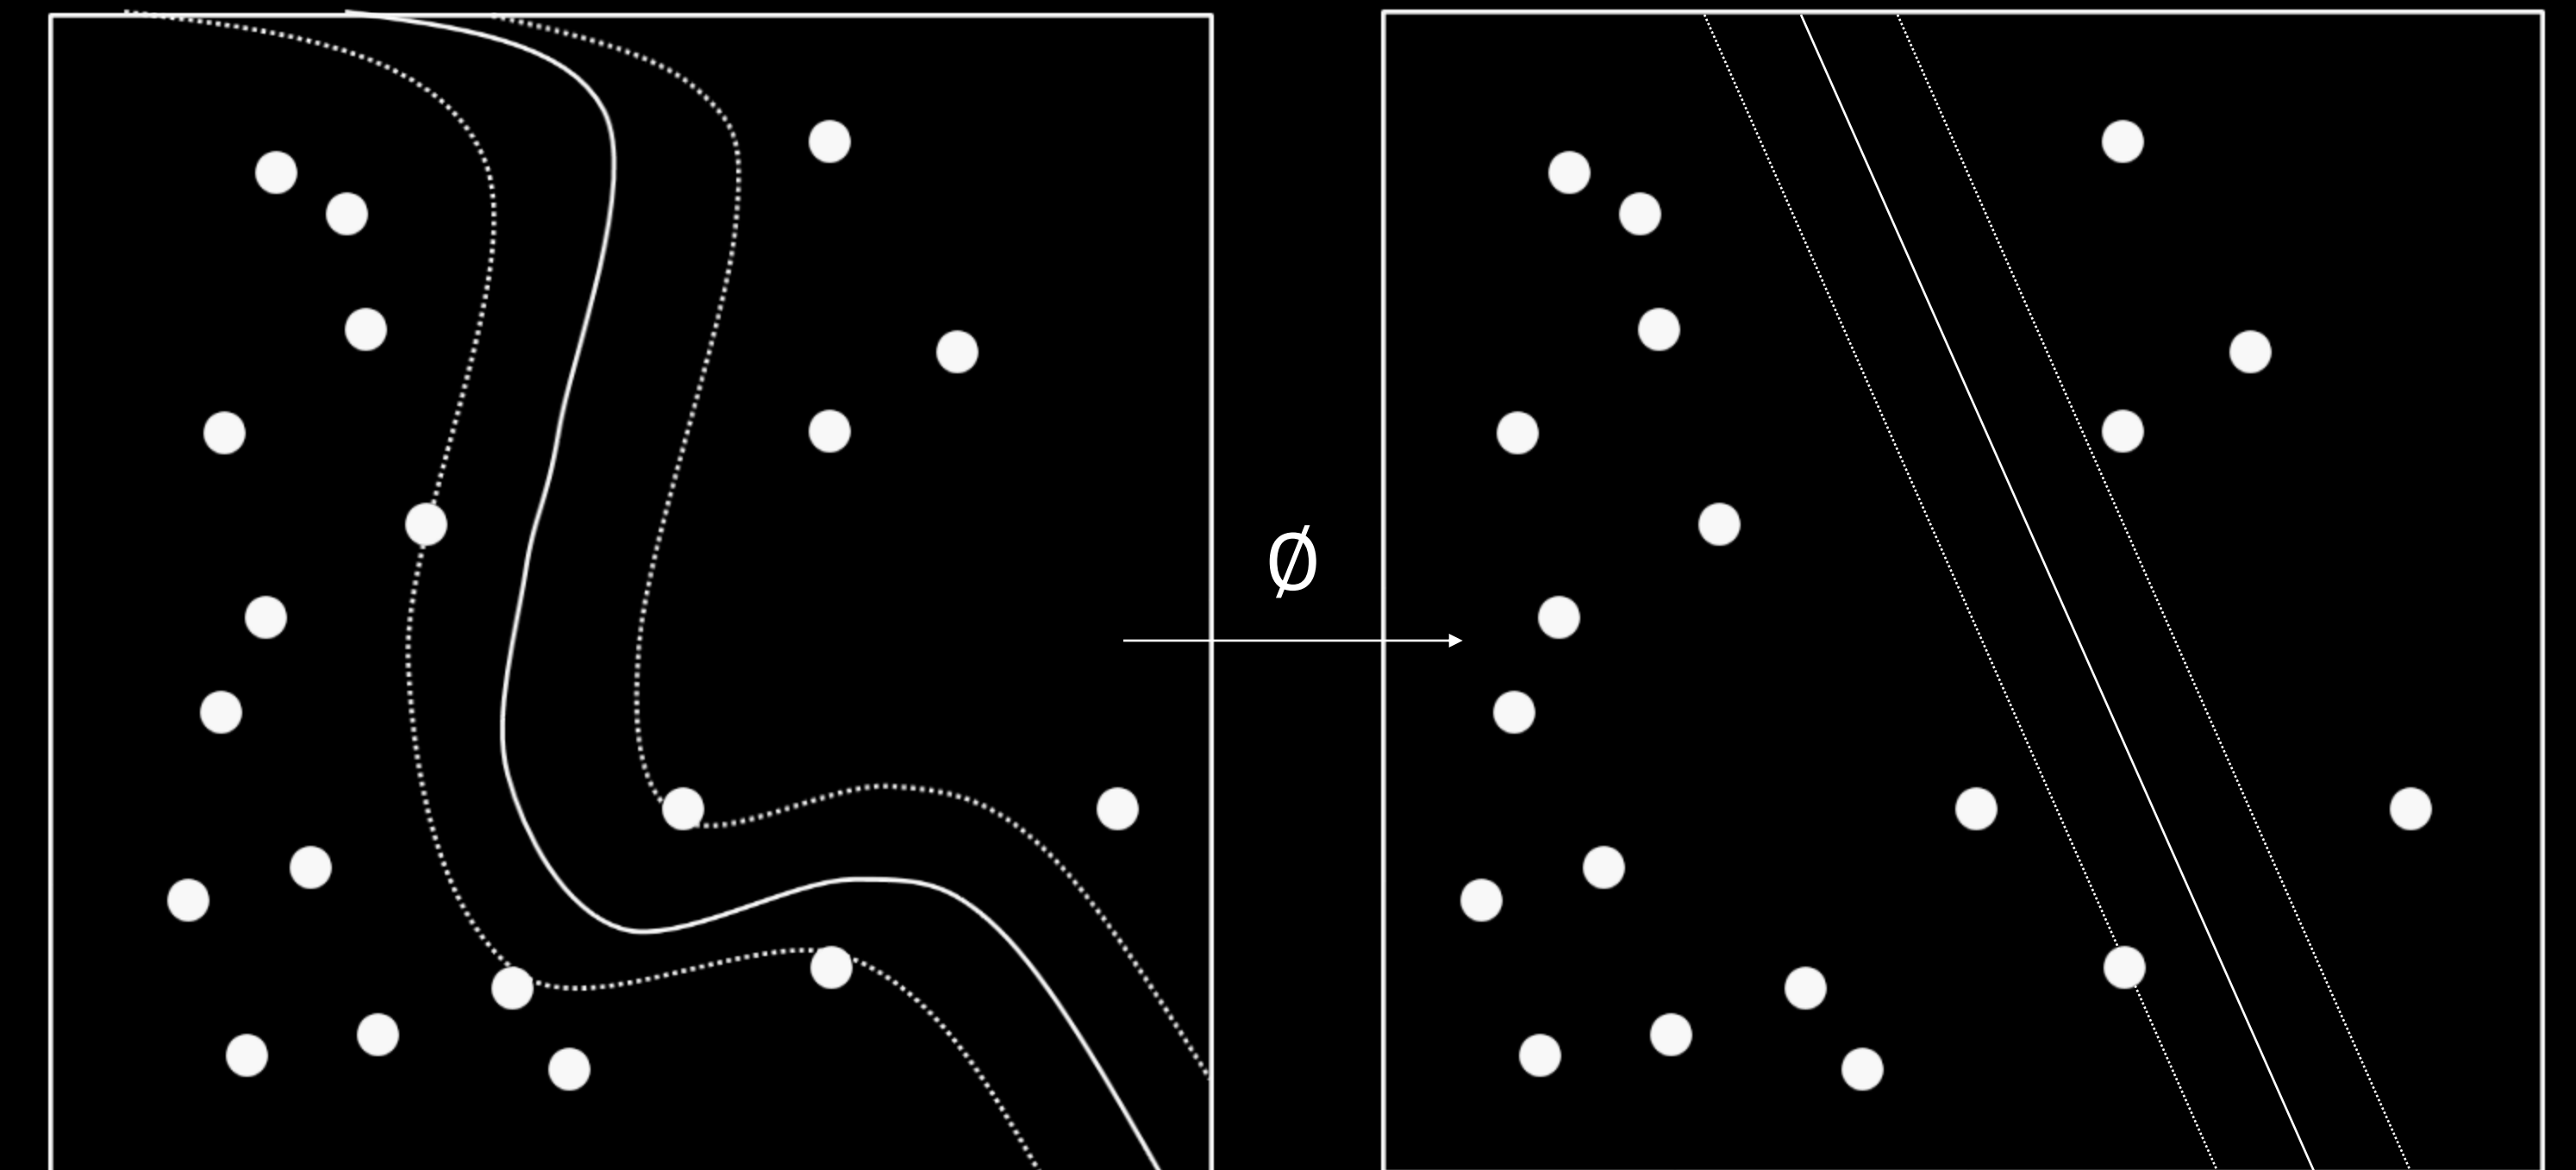
```
bash - cudaq
```

```
Container nvcr.io/nvidia/nightly/cuda-qu... 0 0 0 Ln 53, Col 25 Spaces: 4 UTF-8 LF Python
```

Fraud Detection

HSBC Leverages CUDA-Q to Develop Improved Fraud Detection

- Fraudulent transactions: loss of \$1.9BN per year for UK alone
- Quantum-inspired methods may improve fraud detection
- Reduced false positives by 4%, improved true positives by 2%
- Run as 165 qubit classification problem with CUDA-Q





Agenda

- What is Quantum Accelerated Supercomputing

- Useful Quantum Simulation

- How-to Guide to CUDA-Q

- Distributed Quantum Computing

- Conclusion

Quantum States

Single-qubit states

The zero state

The one state

Quantum states

Ket notation

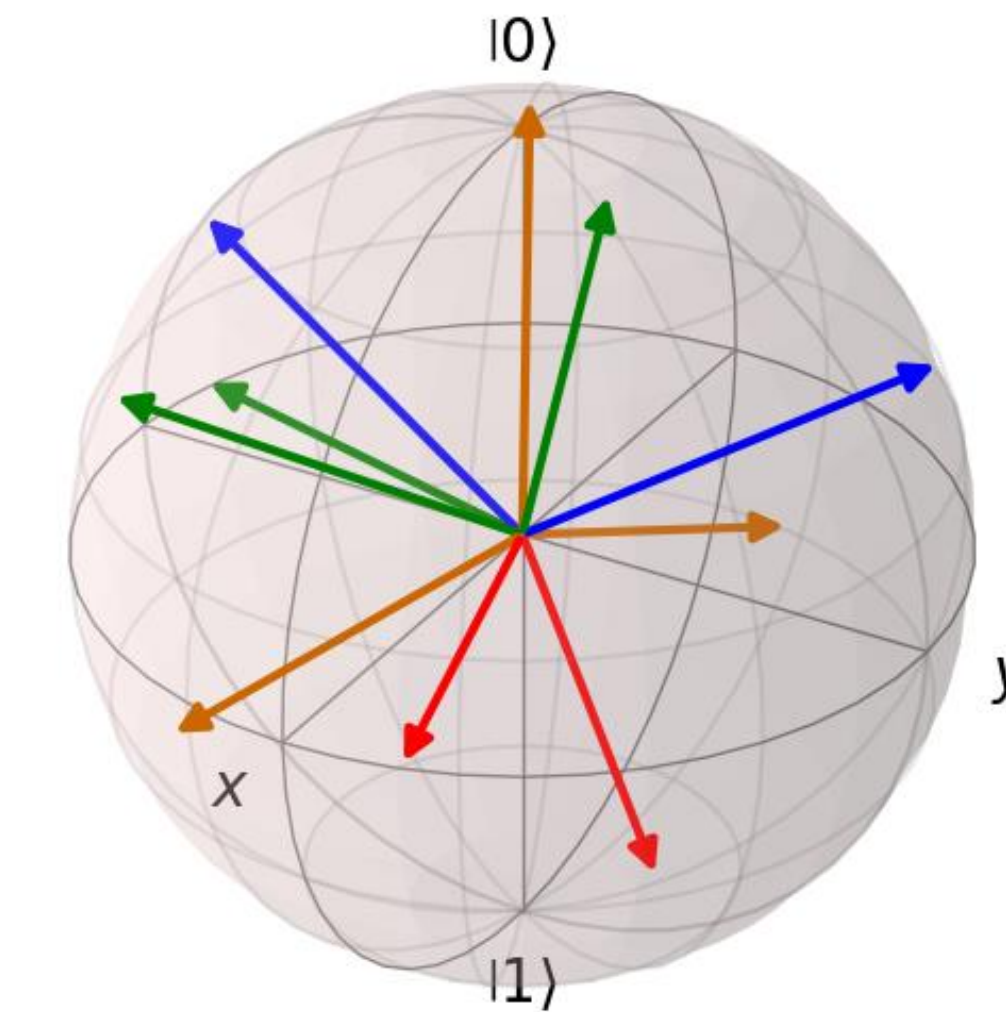
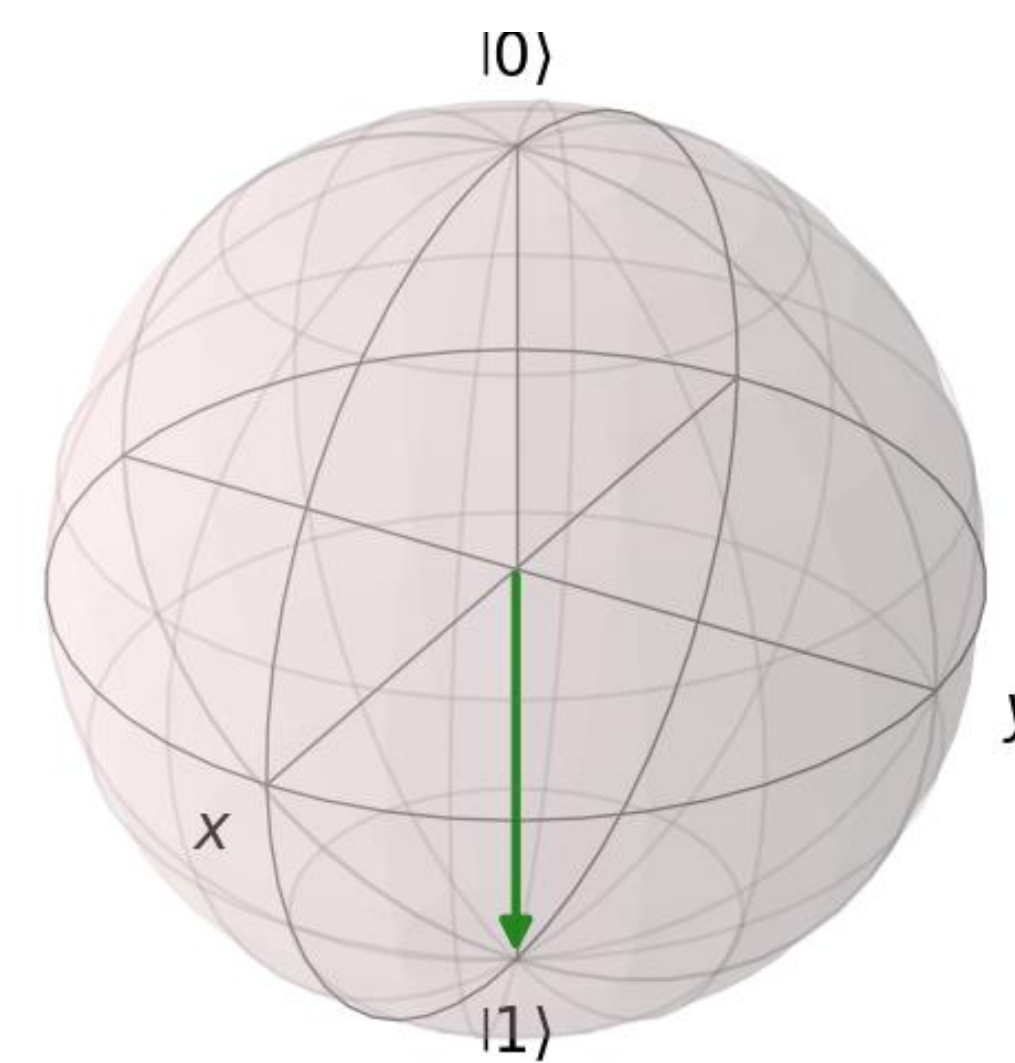
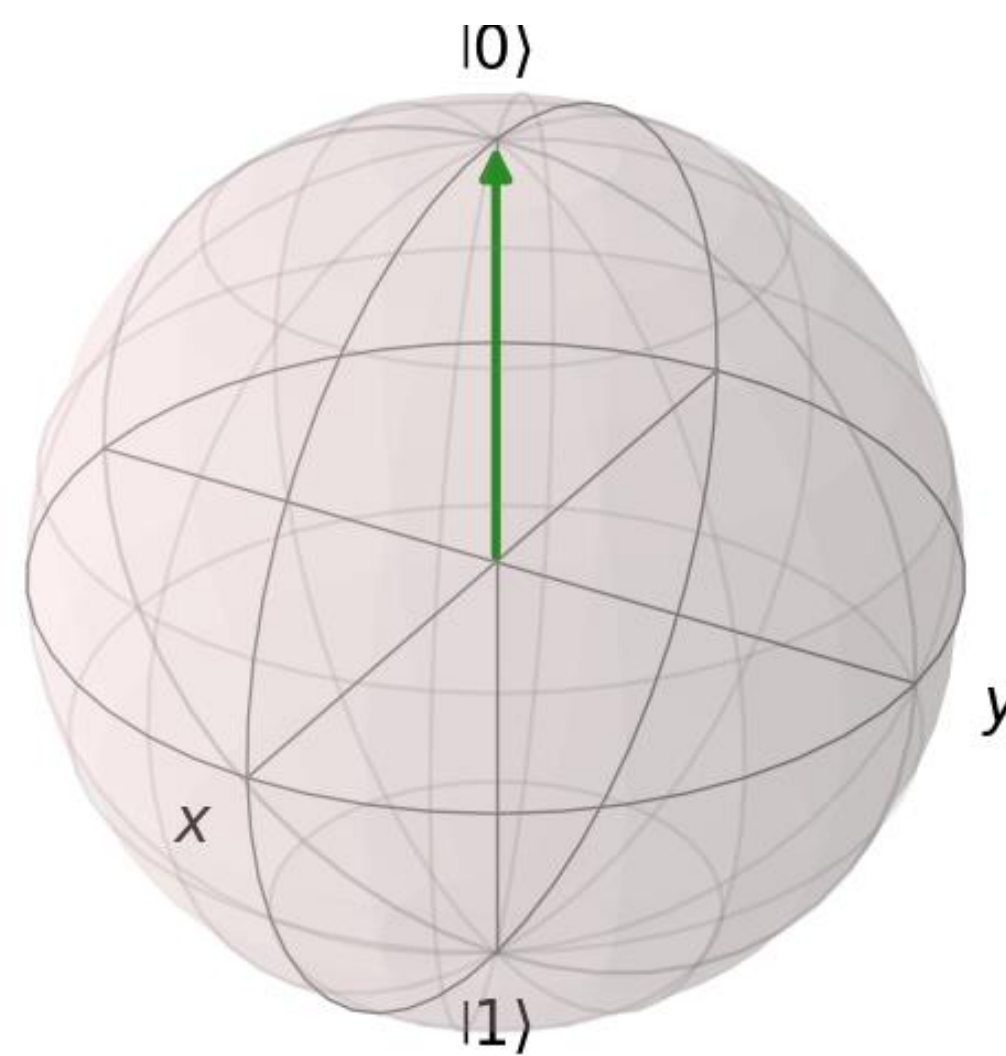
$$|0\rangle = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$$

$$|1\rangle = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$

$$|\psi\rangle = \alpha |0\rangle + \beta |1\rangle,$$

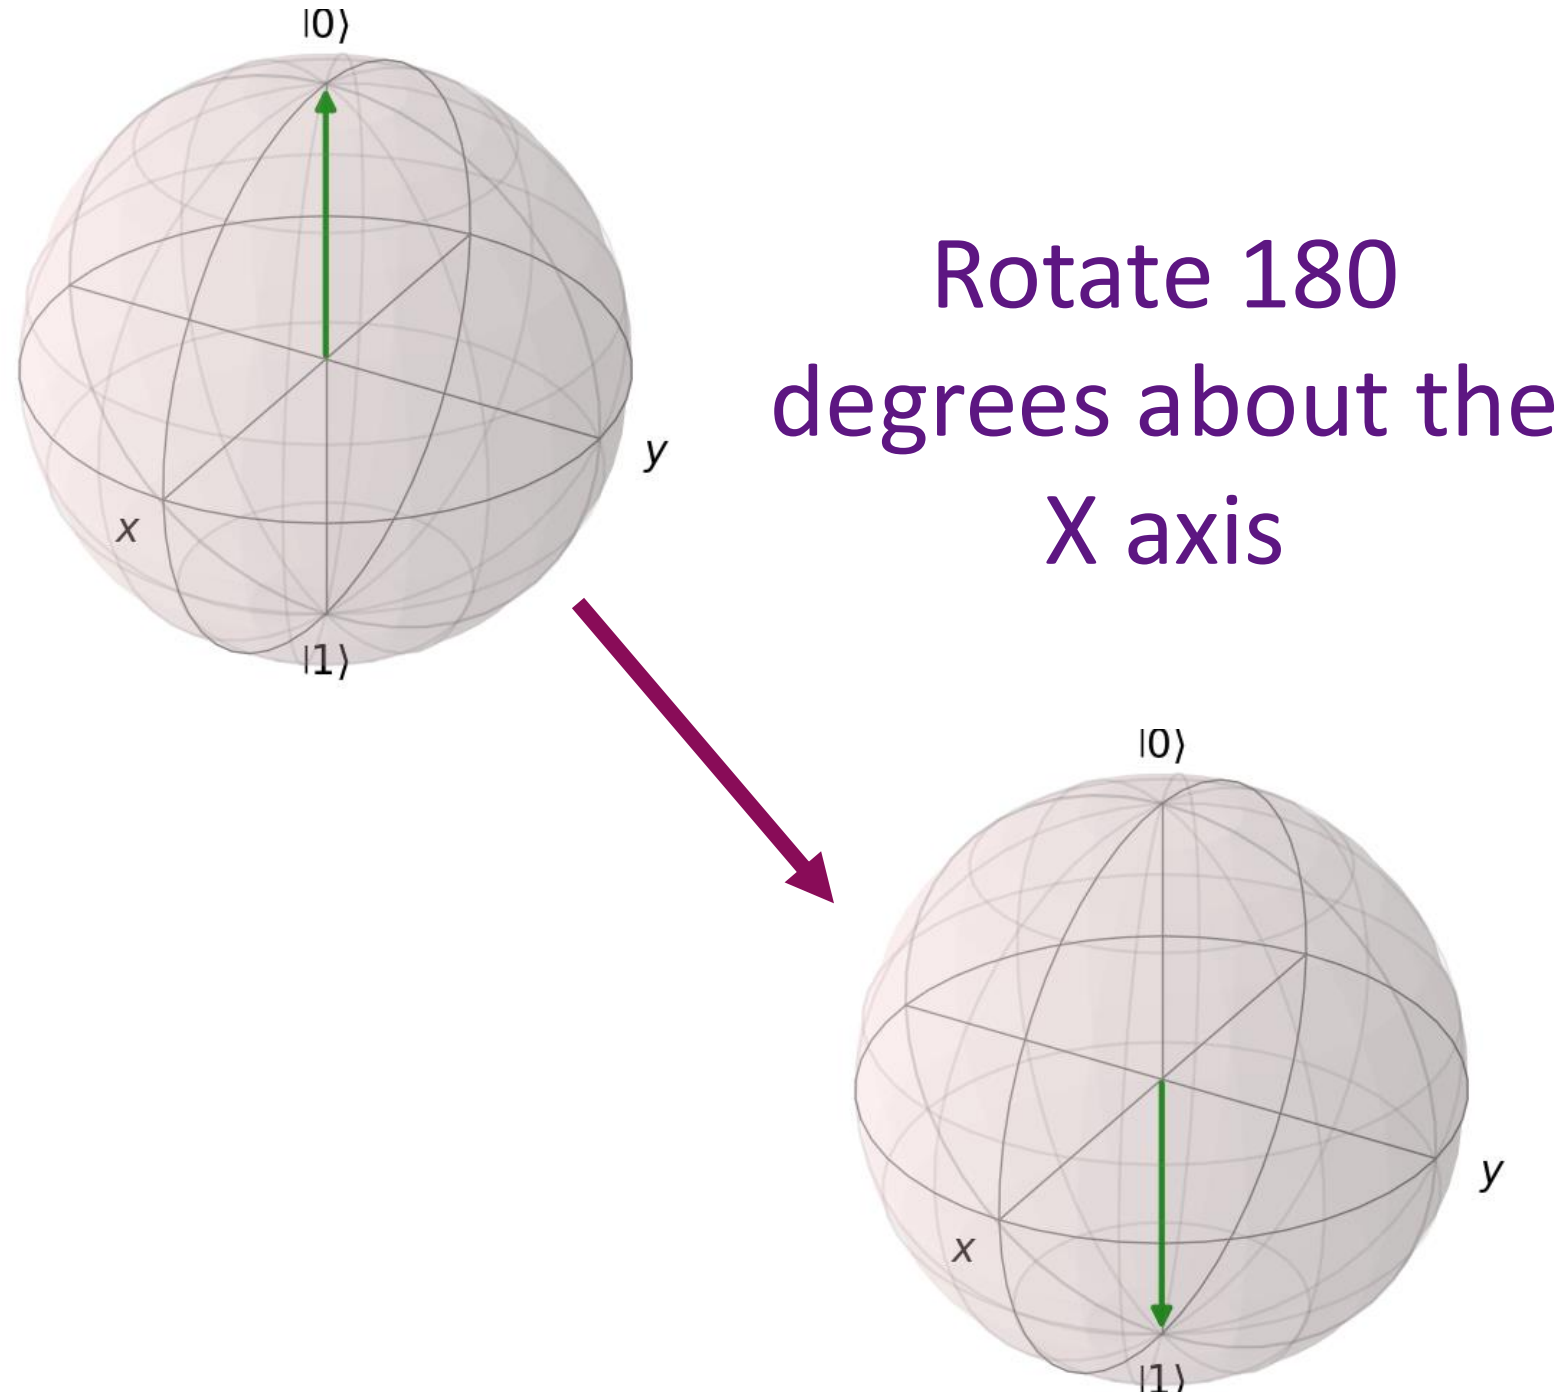
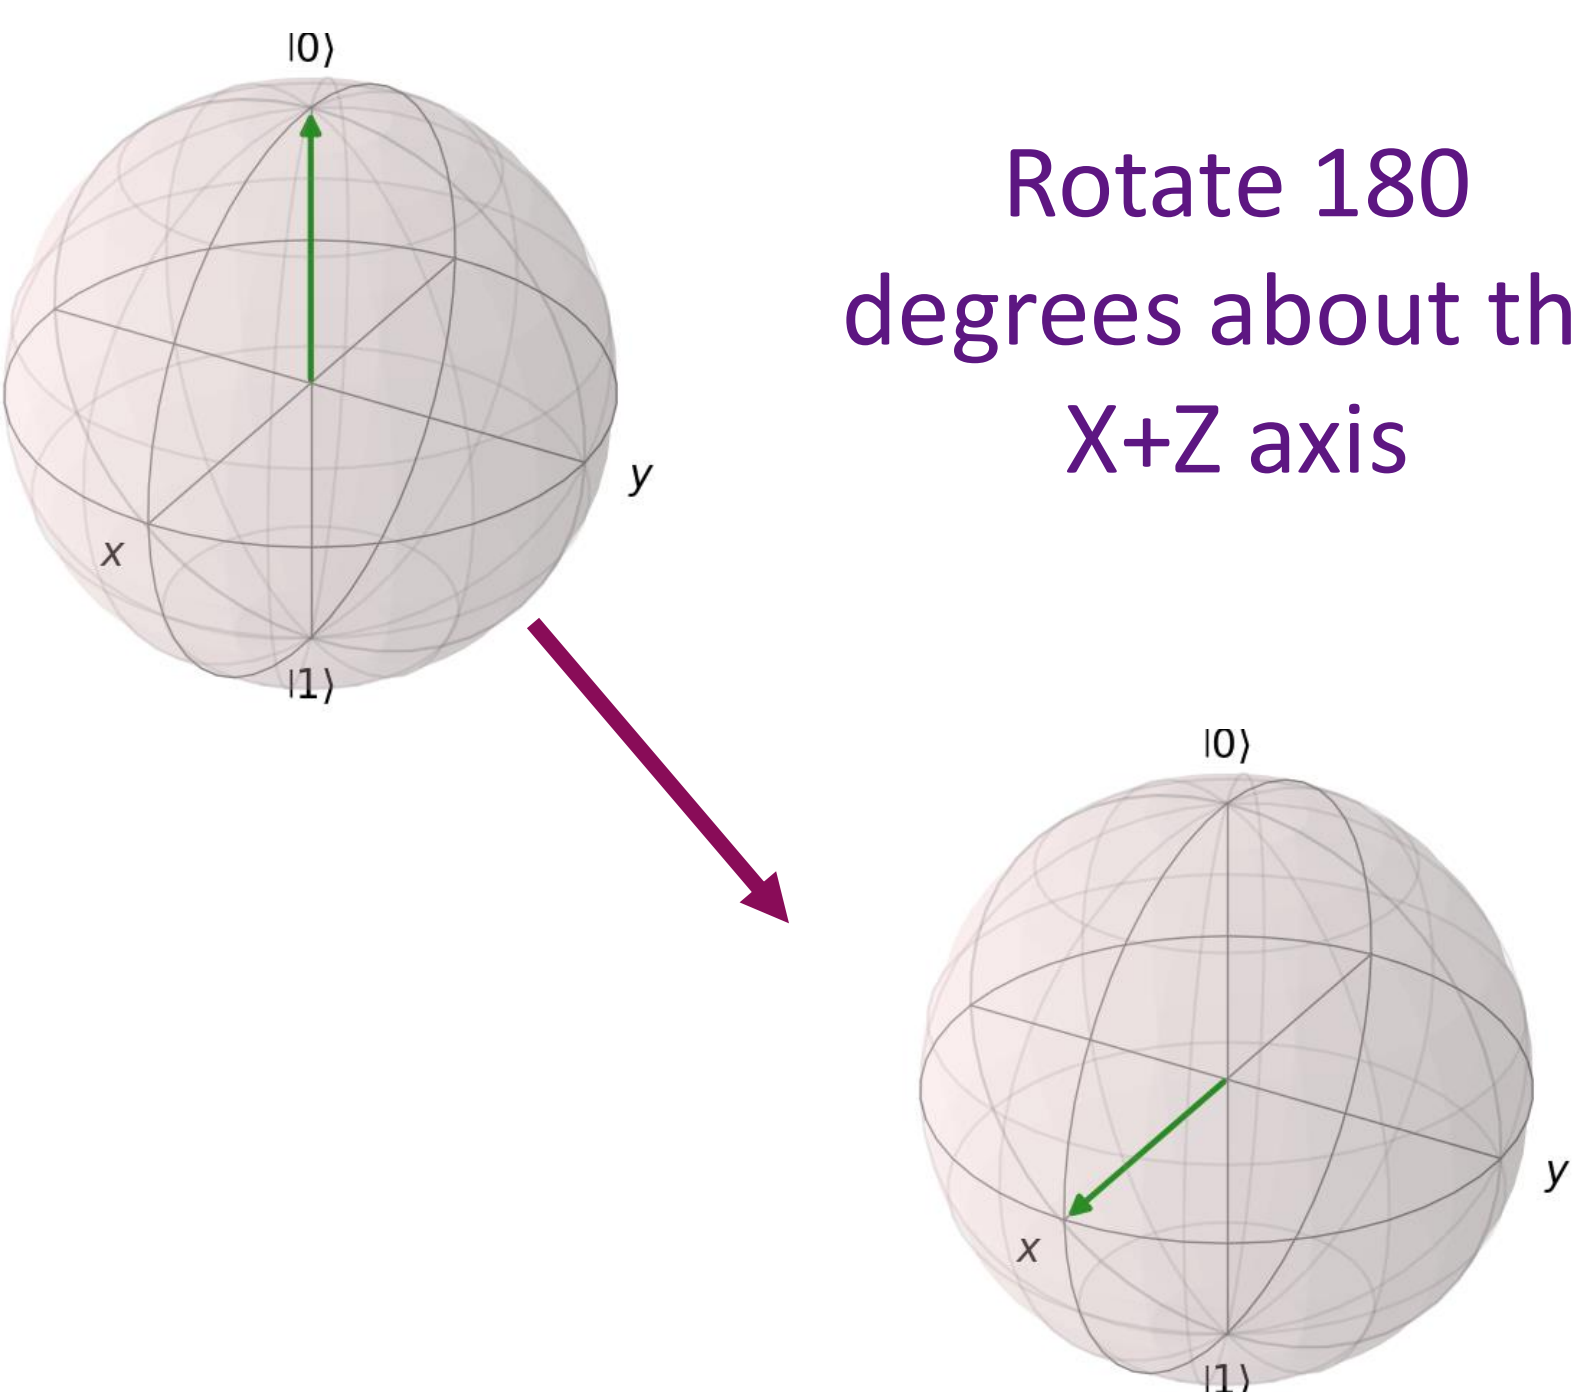
where α and β are complex numbers satisfying the equation $|\alpha|^2 + |\beta|^2 = 1$. The coefficients α and β are referred to as *probability amplitudes*, or *amplitudes* for short.

Bloch Sphere

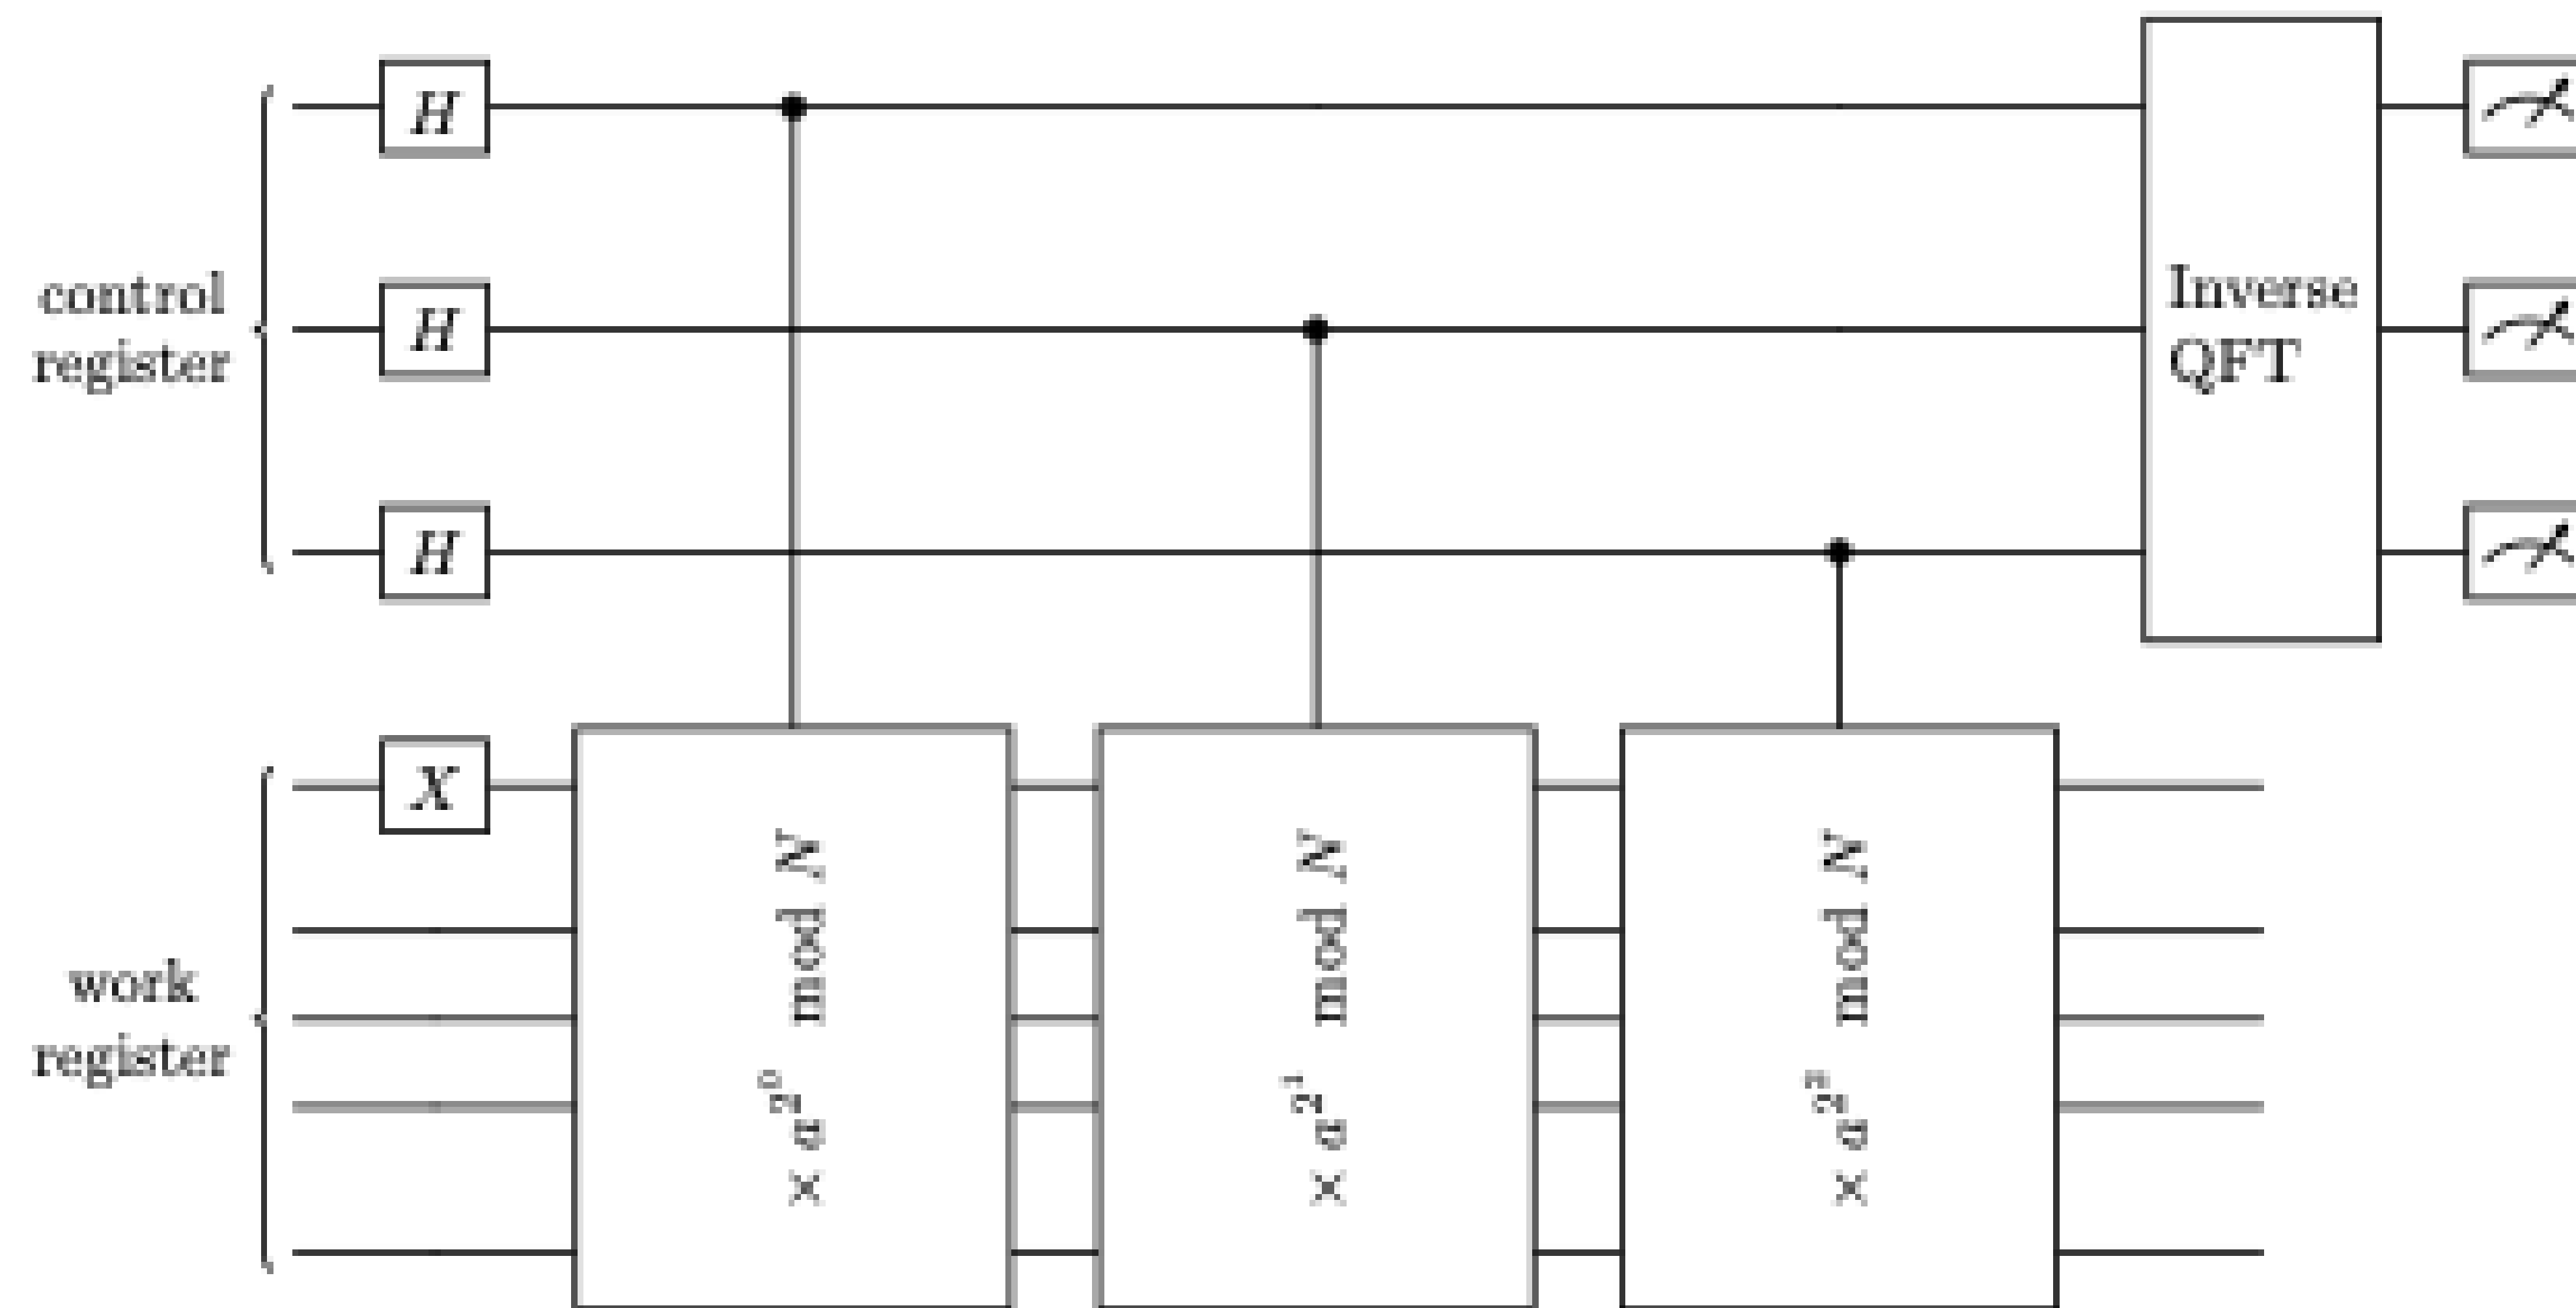


Quantum Gates or Operations

Some examples

	The Bit Flip Gate	The Hadamard Gate	CNOT Gate
Ket notation	$X = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$ $X 0\rangle = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix} \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \begin{pmatrix} 0 \\ 1 \end{pmatrix} = 1\rangle$	$H = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$ $H 0\rangle = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix} \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ 1 \end{pmatrix} = +\rangle$	$CNOT = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \end{pmatrix}$
Bloch Sphere Representation	 <p>Rotate 180 degrees about the X axis</p>	 <p>Rotate 180 degrees about the X+Z axis</p>	$CNOT\left(\frac{1}{\sqrt{2}} 00\rangle + \frac{1}{\sqrt{2}} 10\rangle\right)$ $= \frac{1}{\sqrt{2}}CNOT 00\rangle + \frac{1}{\sqrt{2}}CNOT 10\rangle$ $= \frac{1}{\sqrt{2}} 00\rangle + \frac{1}{\sqrt{2}} 11\rangle$

Quantum Kernels or Circuits



Template for a Quantum Program

- Initialize/allocate the qubits
- Manipulate the quantum with gates
- Extract information from the quantum state by taking measurement(s)

Building your First CUDA-Q Kernel

```
import cudaq

qubit_count = 2

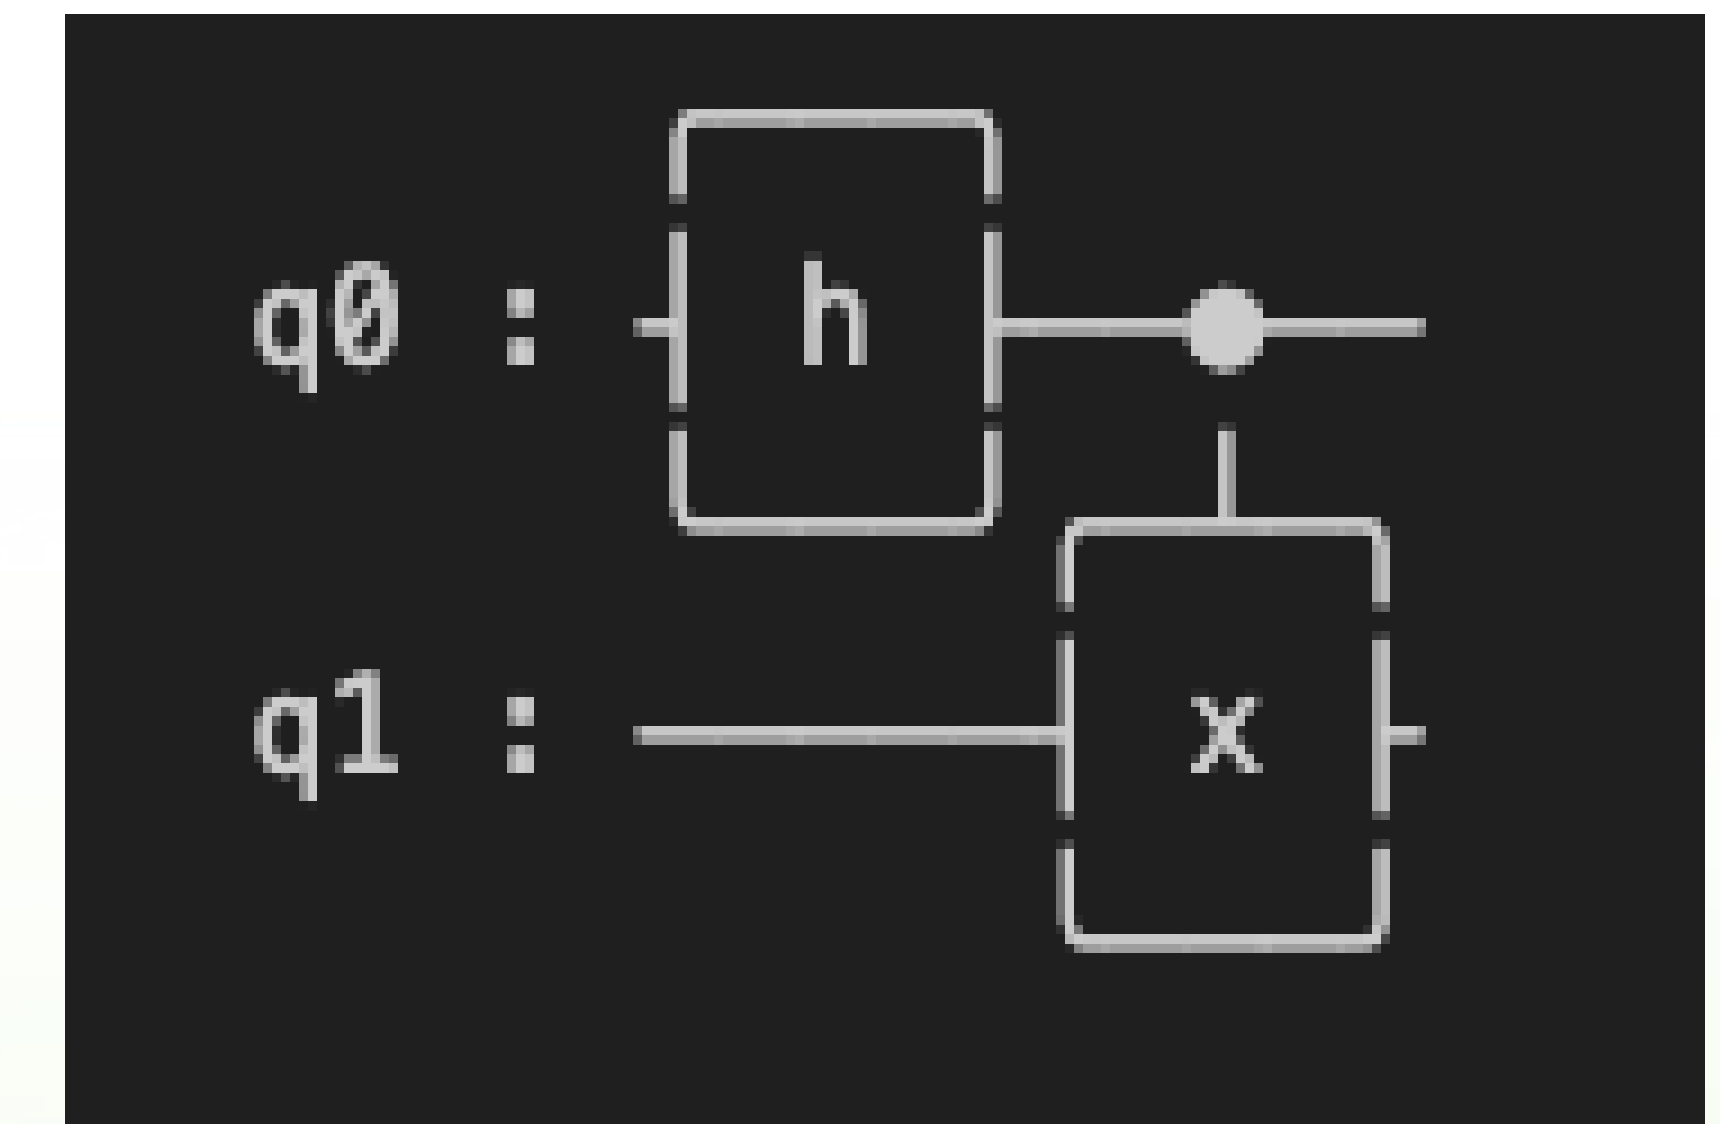
# Define the kernel
@cudaq.kernel
def my_kernel(qubit_count: int):
    # Allocate our `qubit_count` to the kernel.
    qubits = cudaq.qvector(qubit_count)

    # Apply a Hadamard gate to the qubit indexed by 0.
    h(qubits[0])

    # Apply a Controlled-X gate between qubit 0 (acting
    # as the control) and each of the remaining qubits.
    for i in range(qubit_count - 1):
        x.ctrl(qubits[i], qubits[i + 1])

    # Measure the qubits
    # If we don't specify measurements, all qubits are measured in the Z-basis by default.
    mz(qubits)

print(cudaq.draw(my_kernel, qubit_count))
```



- Initialize/allocate the qubits
- Manipulate the quantum with gates
- Extract information from the quantum state by taking measurement(s)

Sampling your First CUDA-Q Kernel

```
# First set the backend for kernel execution
cudaq.set_target('qpp-cpu') # selects a CPU backend

if cudaq.num_available_gpus() > 0:
    cudaq.set_target('nvidia') # selects a GPU backend

# cudaq.set_target('nvqc') # selects the NVIDIA Quantum Cloud
# cudaq.set_target('ionq') # select an available QPU backend

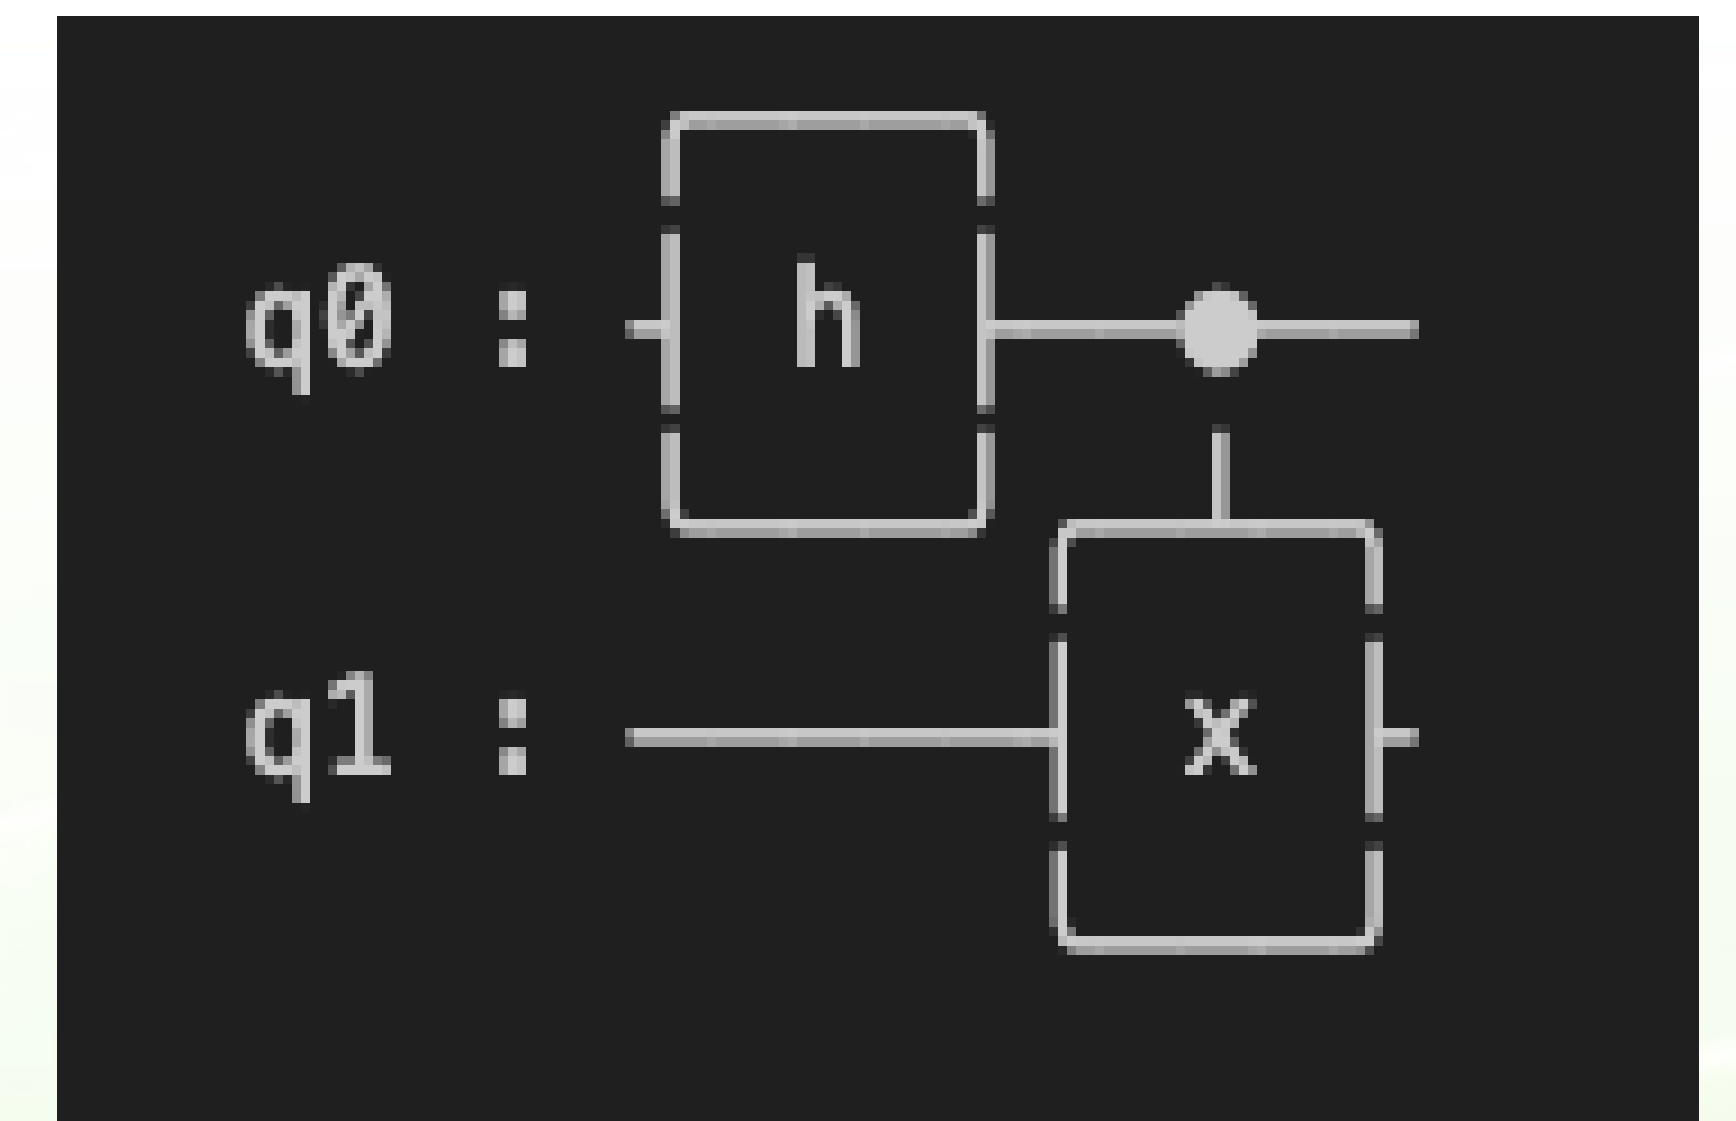
qubit_count = 2

results = cudaq.sample(my_kernel, qubit_count, shots_count = 10000)

print(results) # Example: {00:5005, 11: 4995}

print(results.most_probable()) # prints: `00`

print(results.probability(results.most_probable())) # prints: `0.5005`
```



Computing Expectation Values with CUDA-Q

```
import cudaq
from cudaq import spin

# First set the backend for kernel execution
cudaq.set_target('qpp-cpu') # selects a CPU backend

if cudaq.num_available_gpus() > 0:
    cudaq.set_target('nvidia') # selects a GPU backend

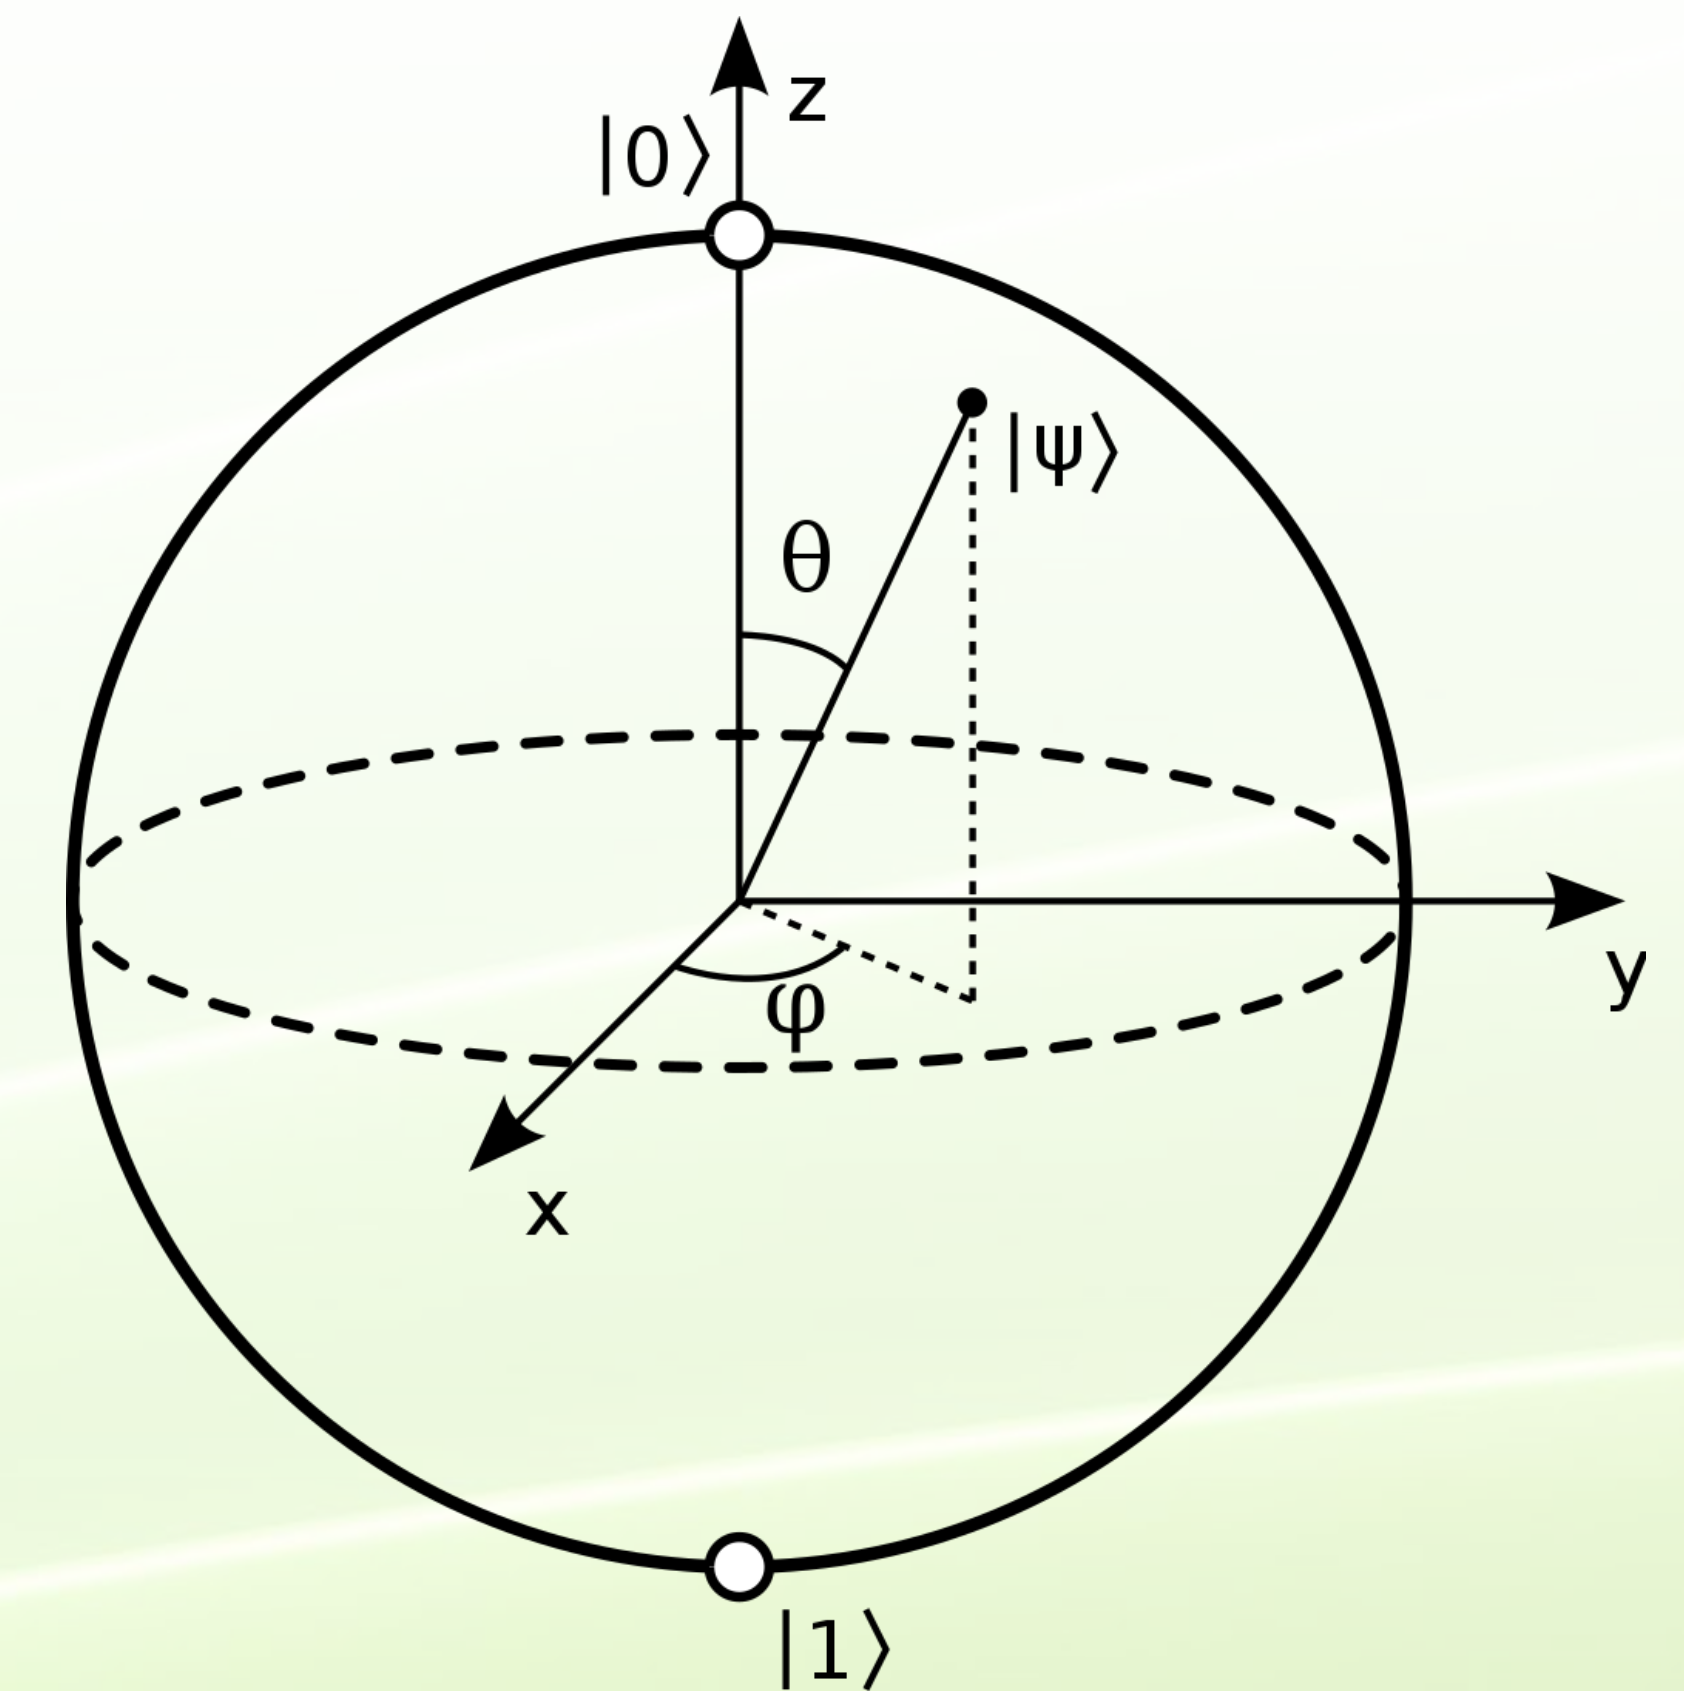
# Define your Hamiltonian Operator
operator = spin.z(0)
print(operator) # prints: [1+0j] Z

# Define your kernel to generate the plus state
@cudaq.kernel
def plus_state():
    qubit = cudaq.qubit()
    h(qubit)

result = cudaq.observe(plus_state, operator, shots_count = 10000)

print(result.expectation()) # prints the approximate expectation value computed from 10000 shots
```

$$\langle + | Z | + \rangle$$

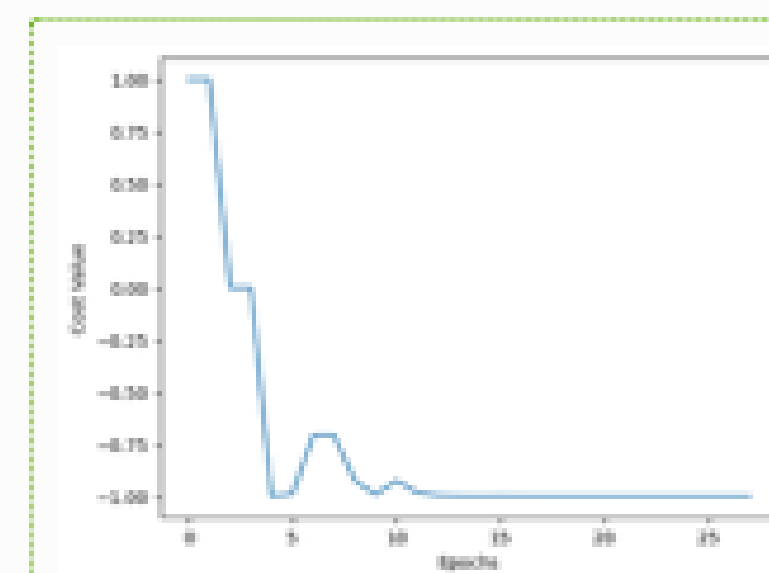


More Sample Code on our Website

<https://nvidia.github.io/cuda-quantum>

CUDA-Q Tutorials

Tutorials that give an in depth view of CUDA-Q and its applications in Python.



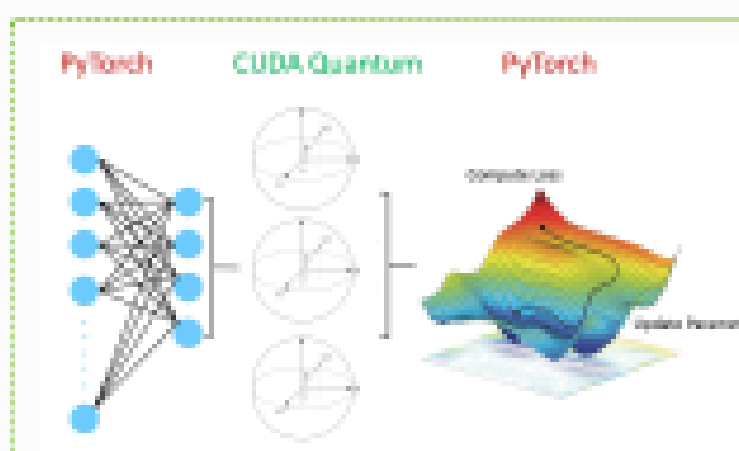
Cost Minimization



Deutsch's Algorithm



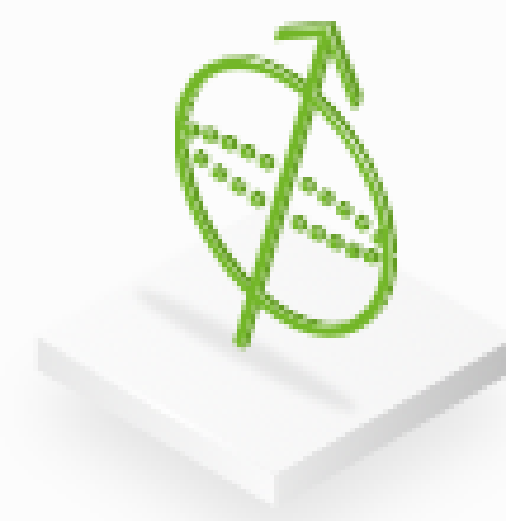
Executing Quantum Circuits



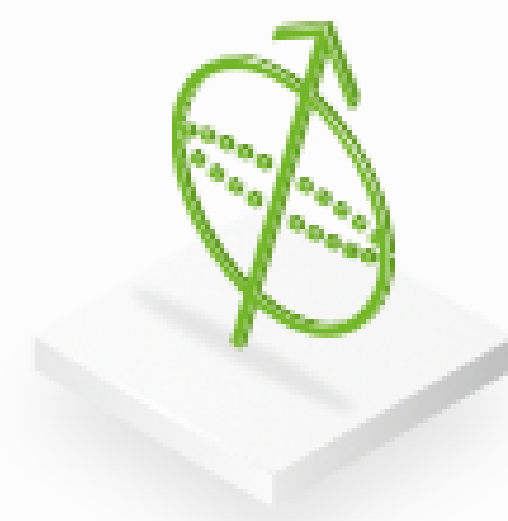
Hybrid Quantum Neural Networks



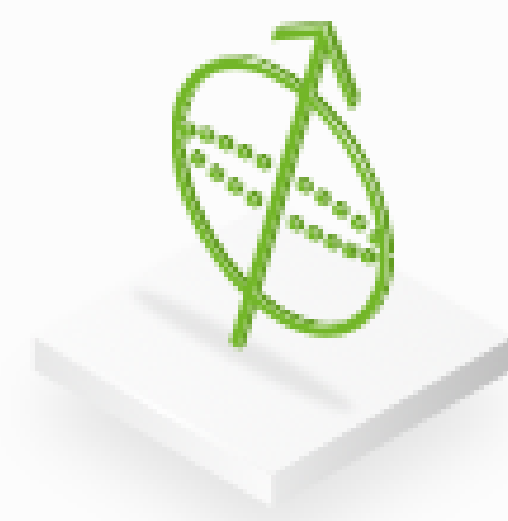
Multi-GPU Workflows



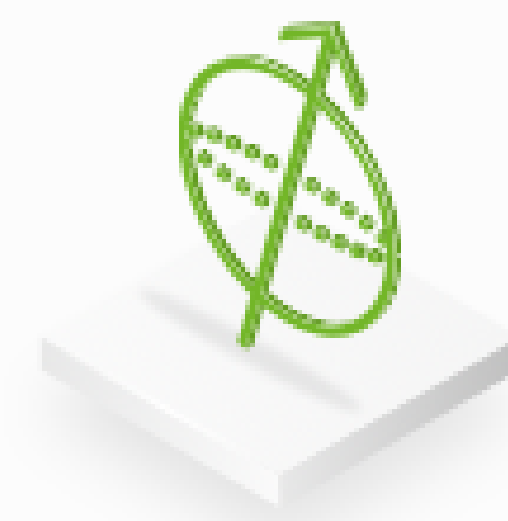
Multiple Qubits



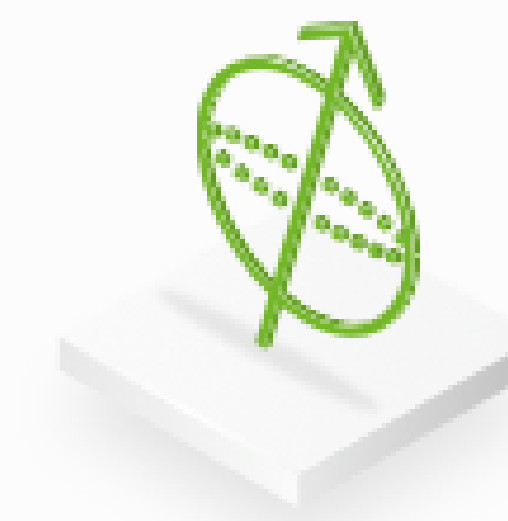
Noisy Simulation



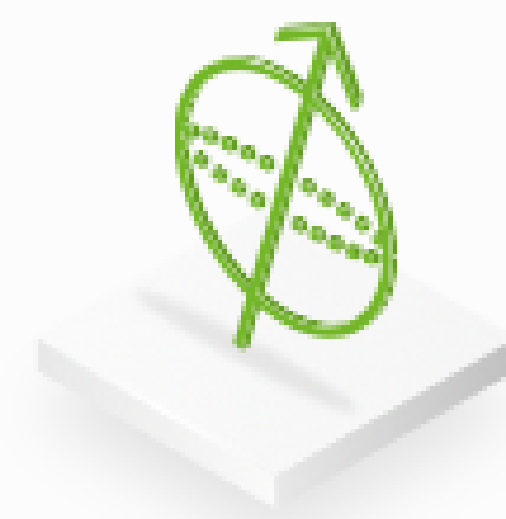
Quantum Approximate Optimization Algorithm for Max Cut



Quantum Bits



Quantum Fourier Transform



Variational Quantum Eigensolver



Agenda

- What is Quantum Accelerated Supercomputing

- Useful Quantum Simulation

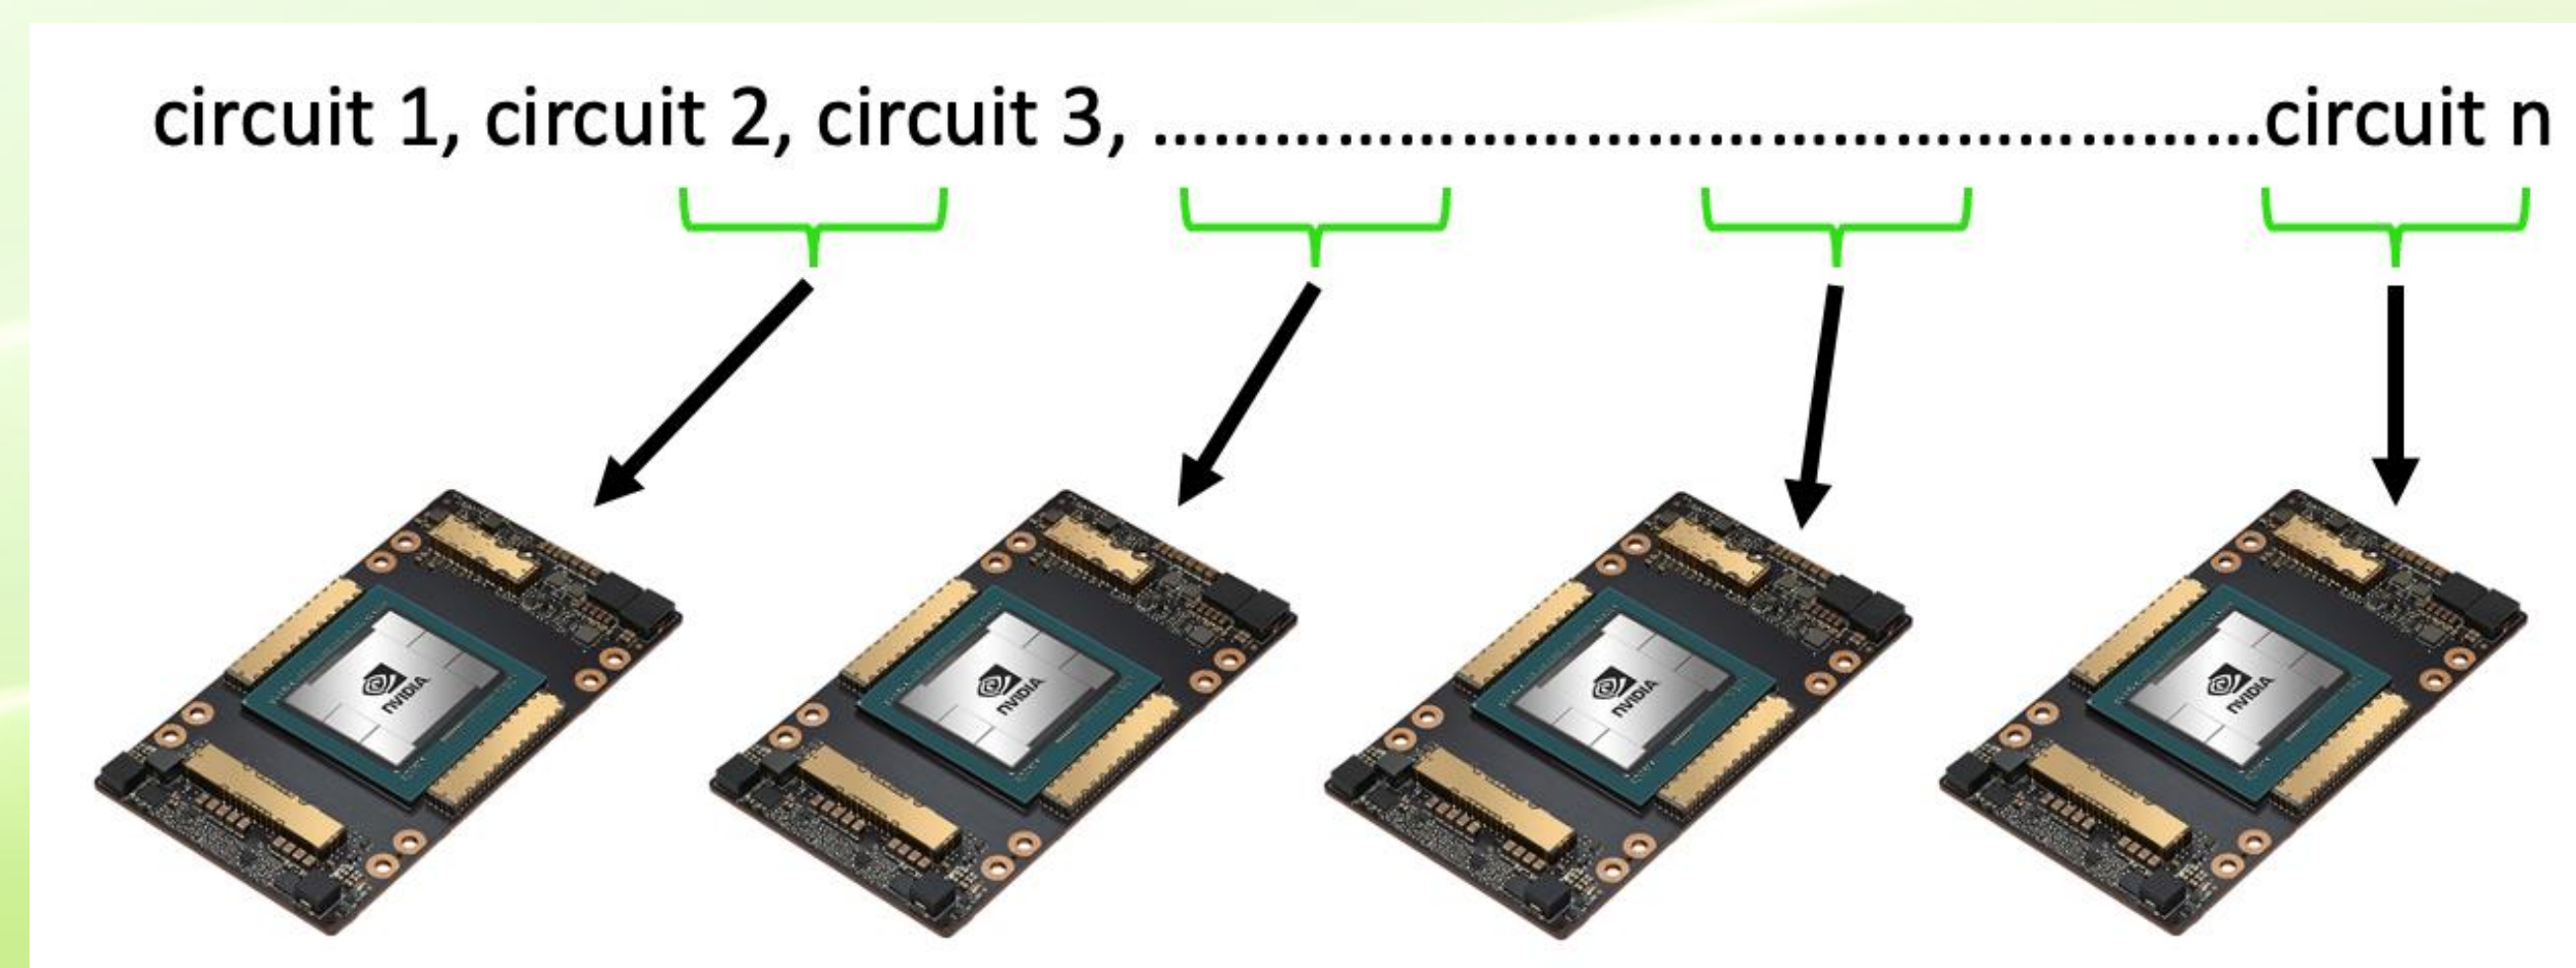
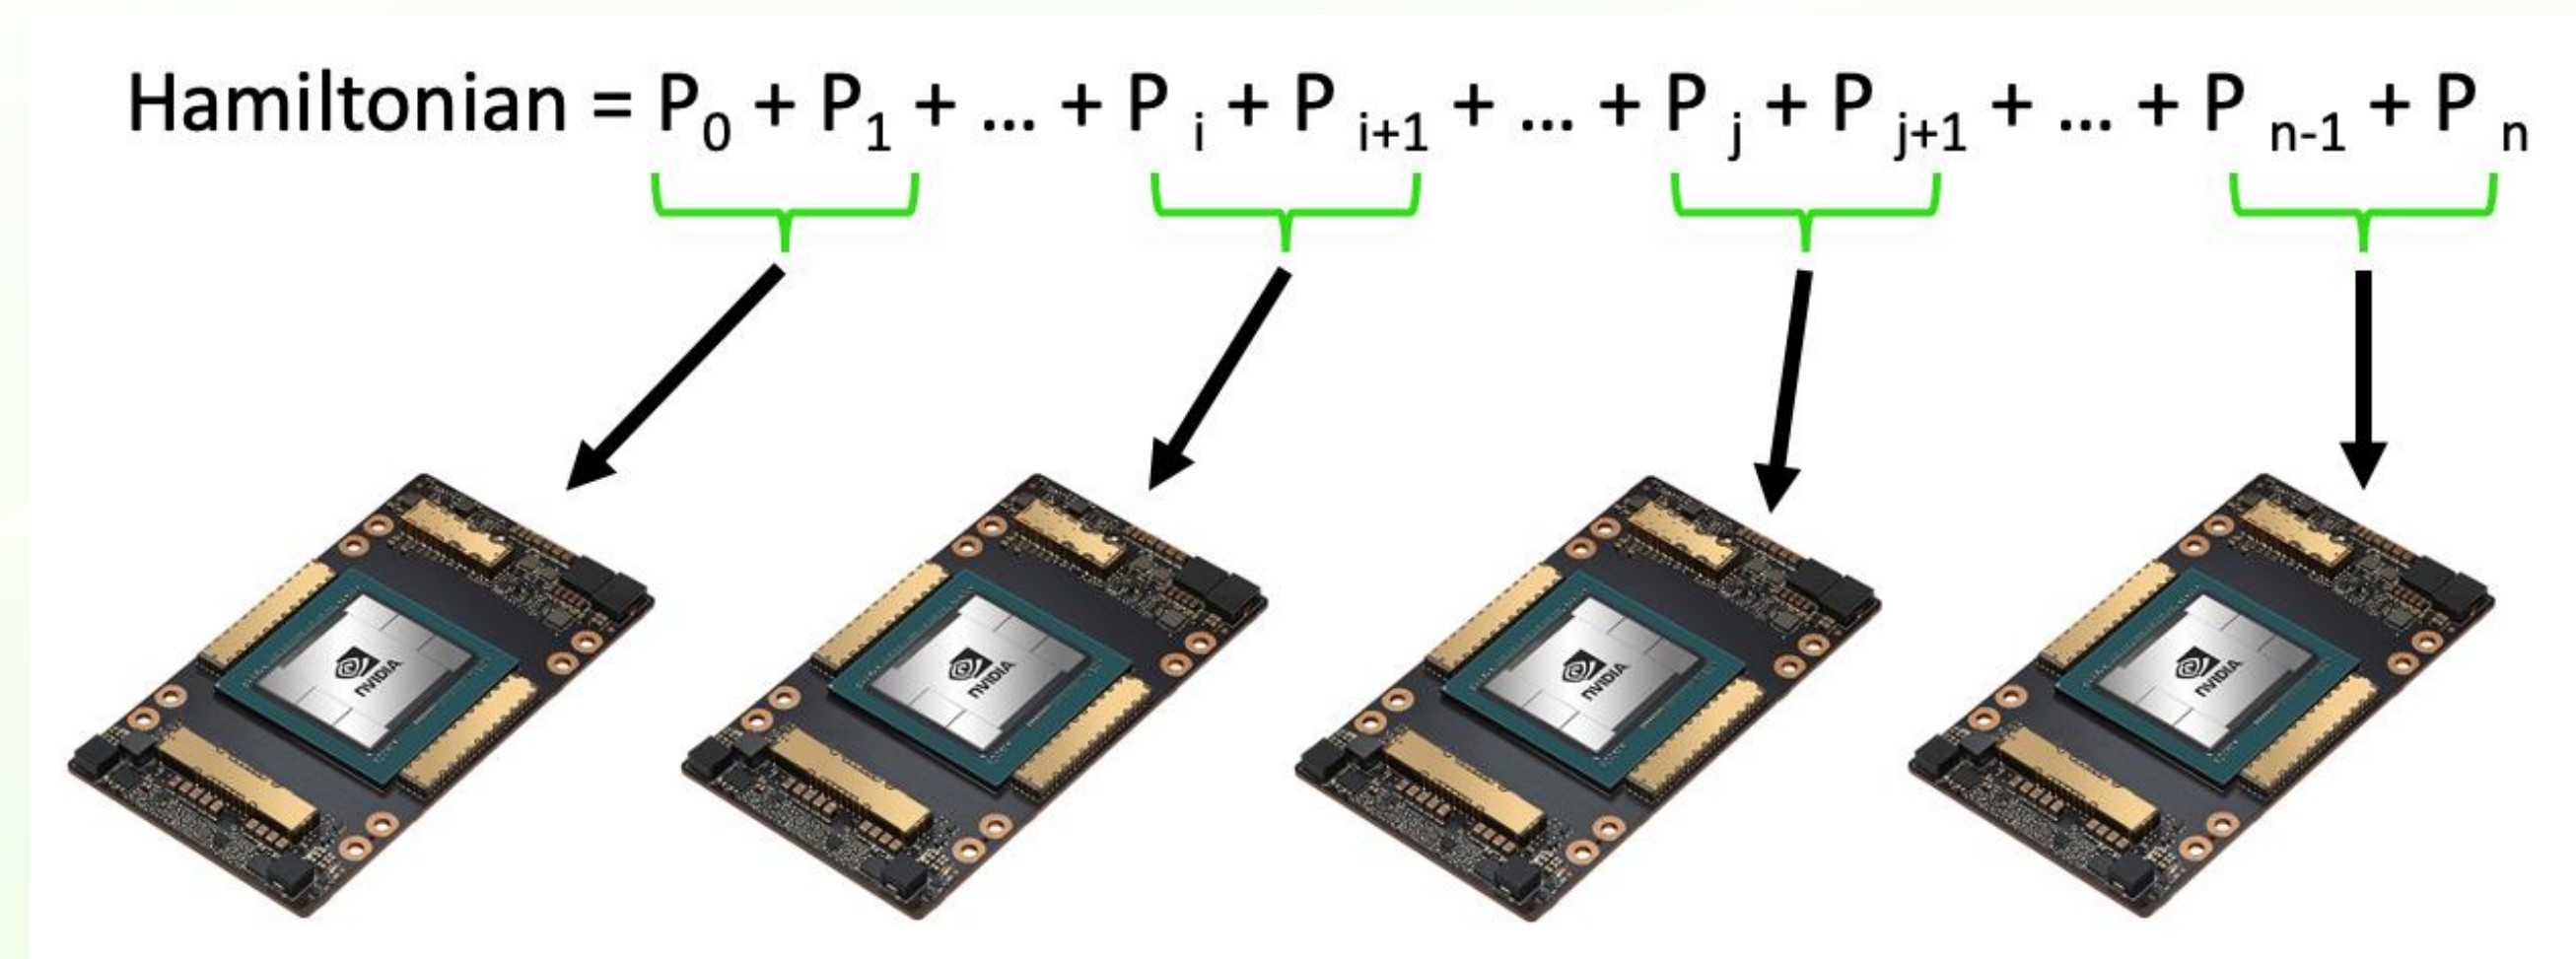
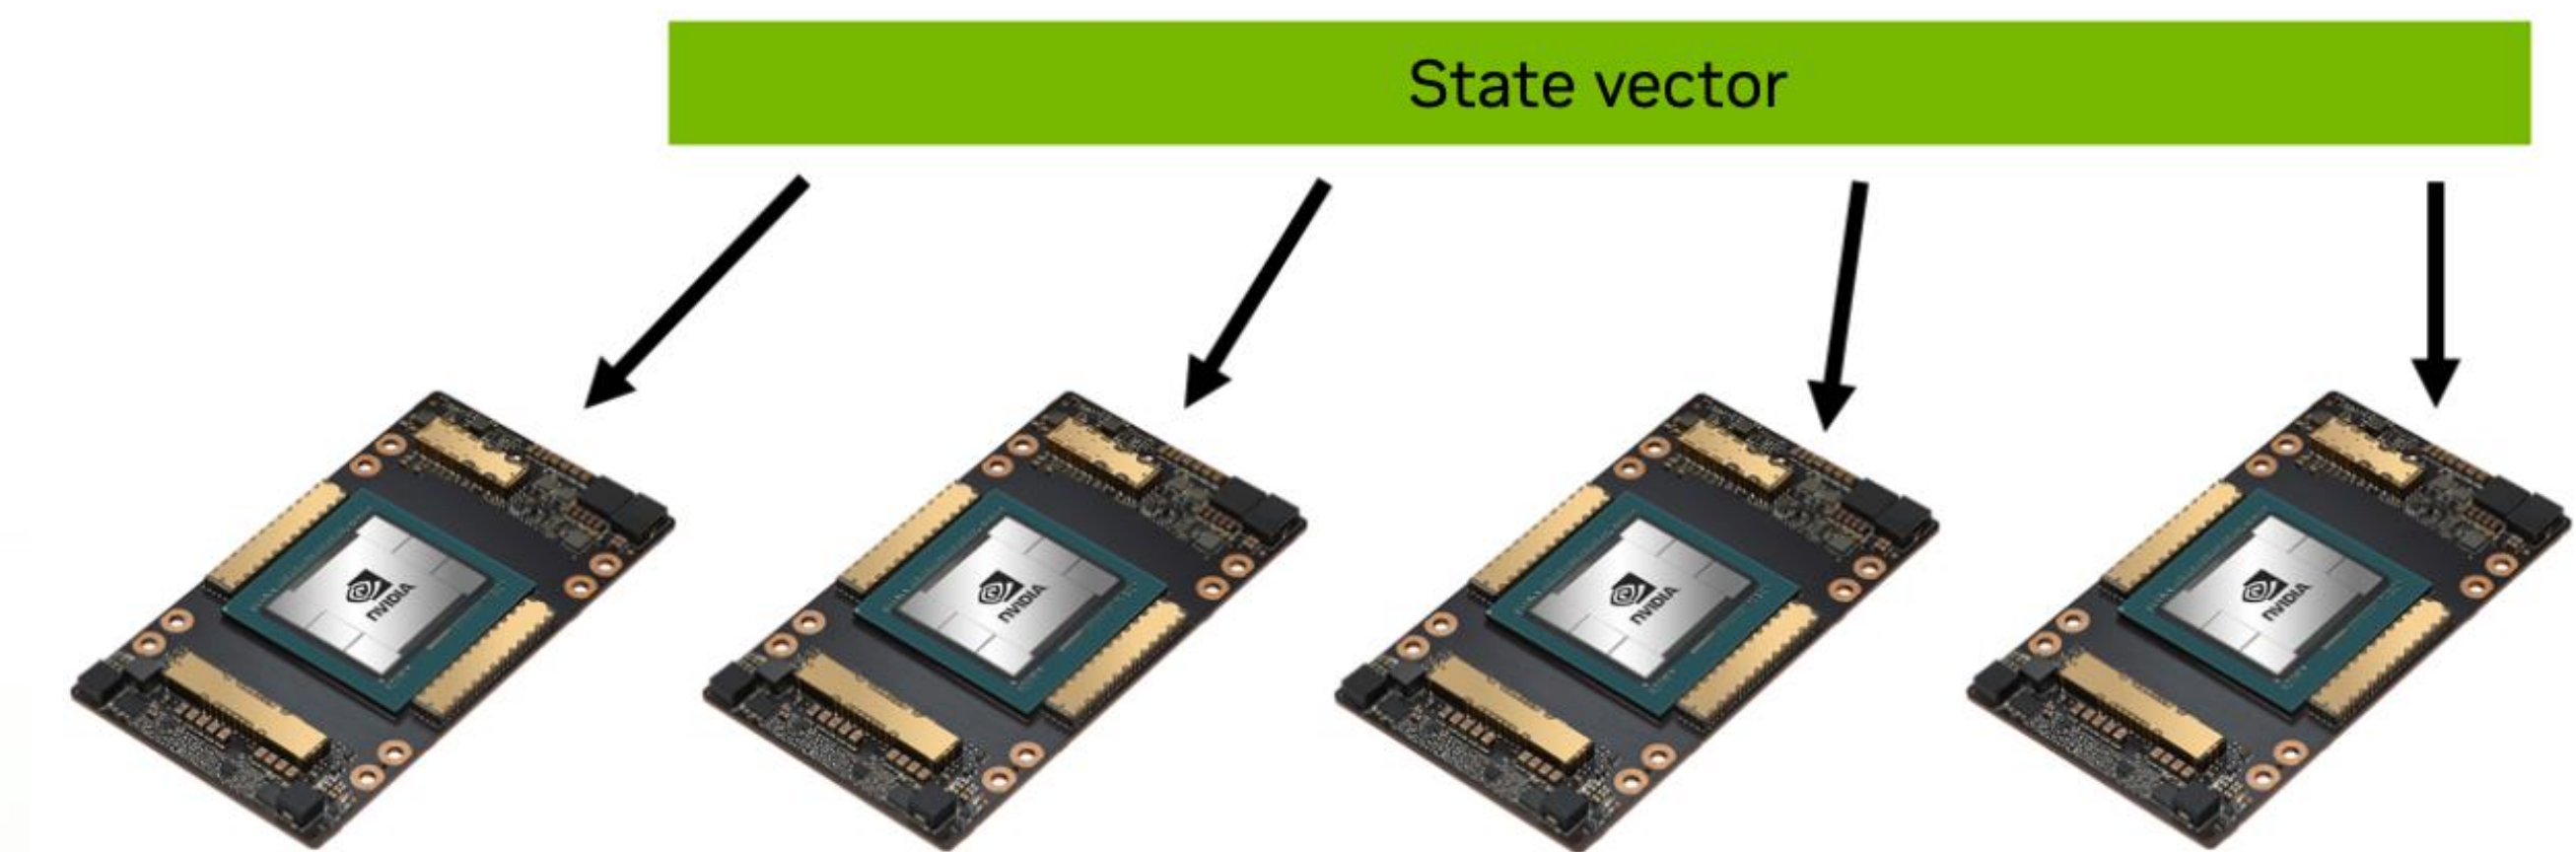
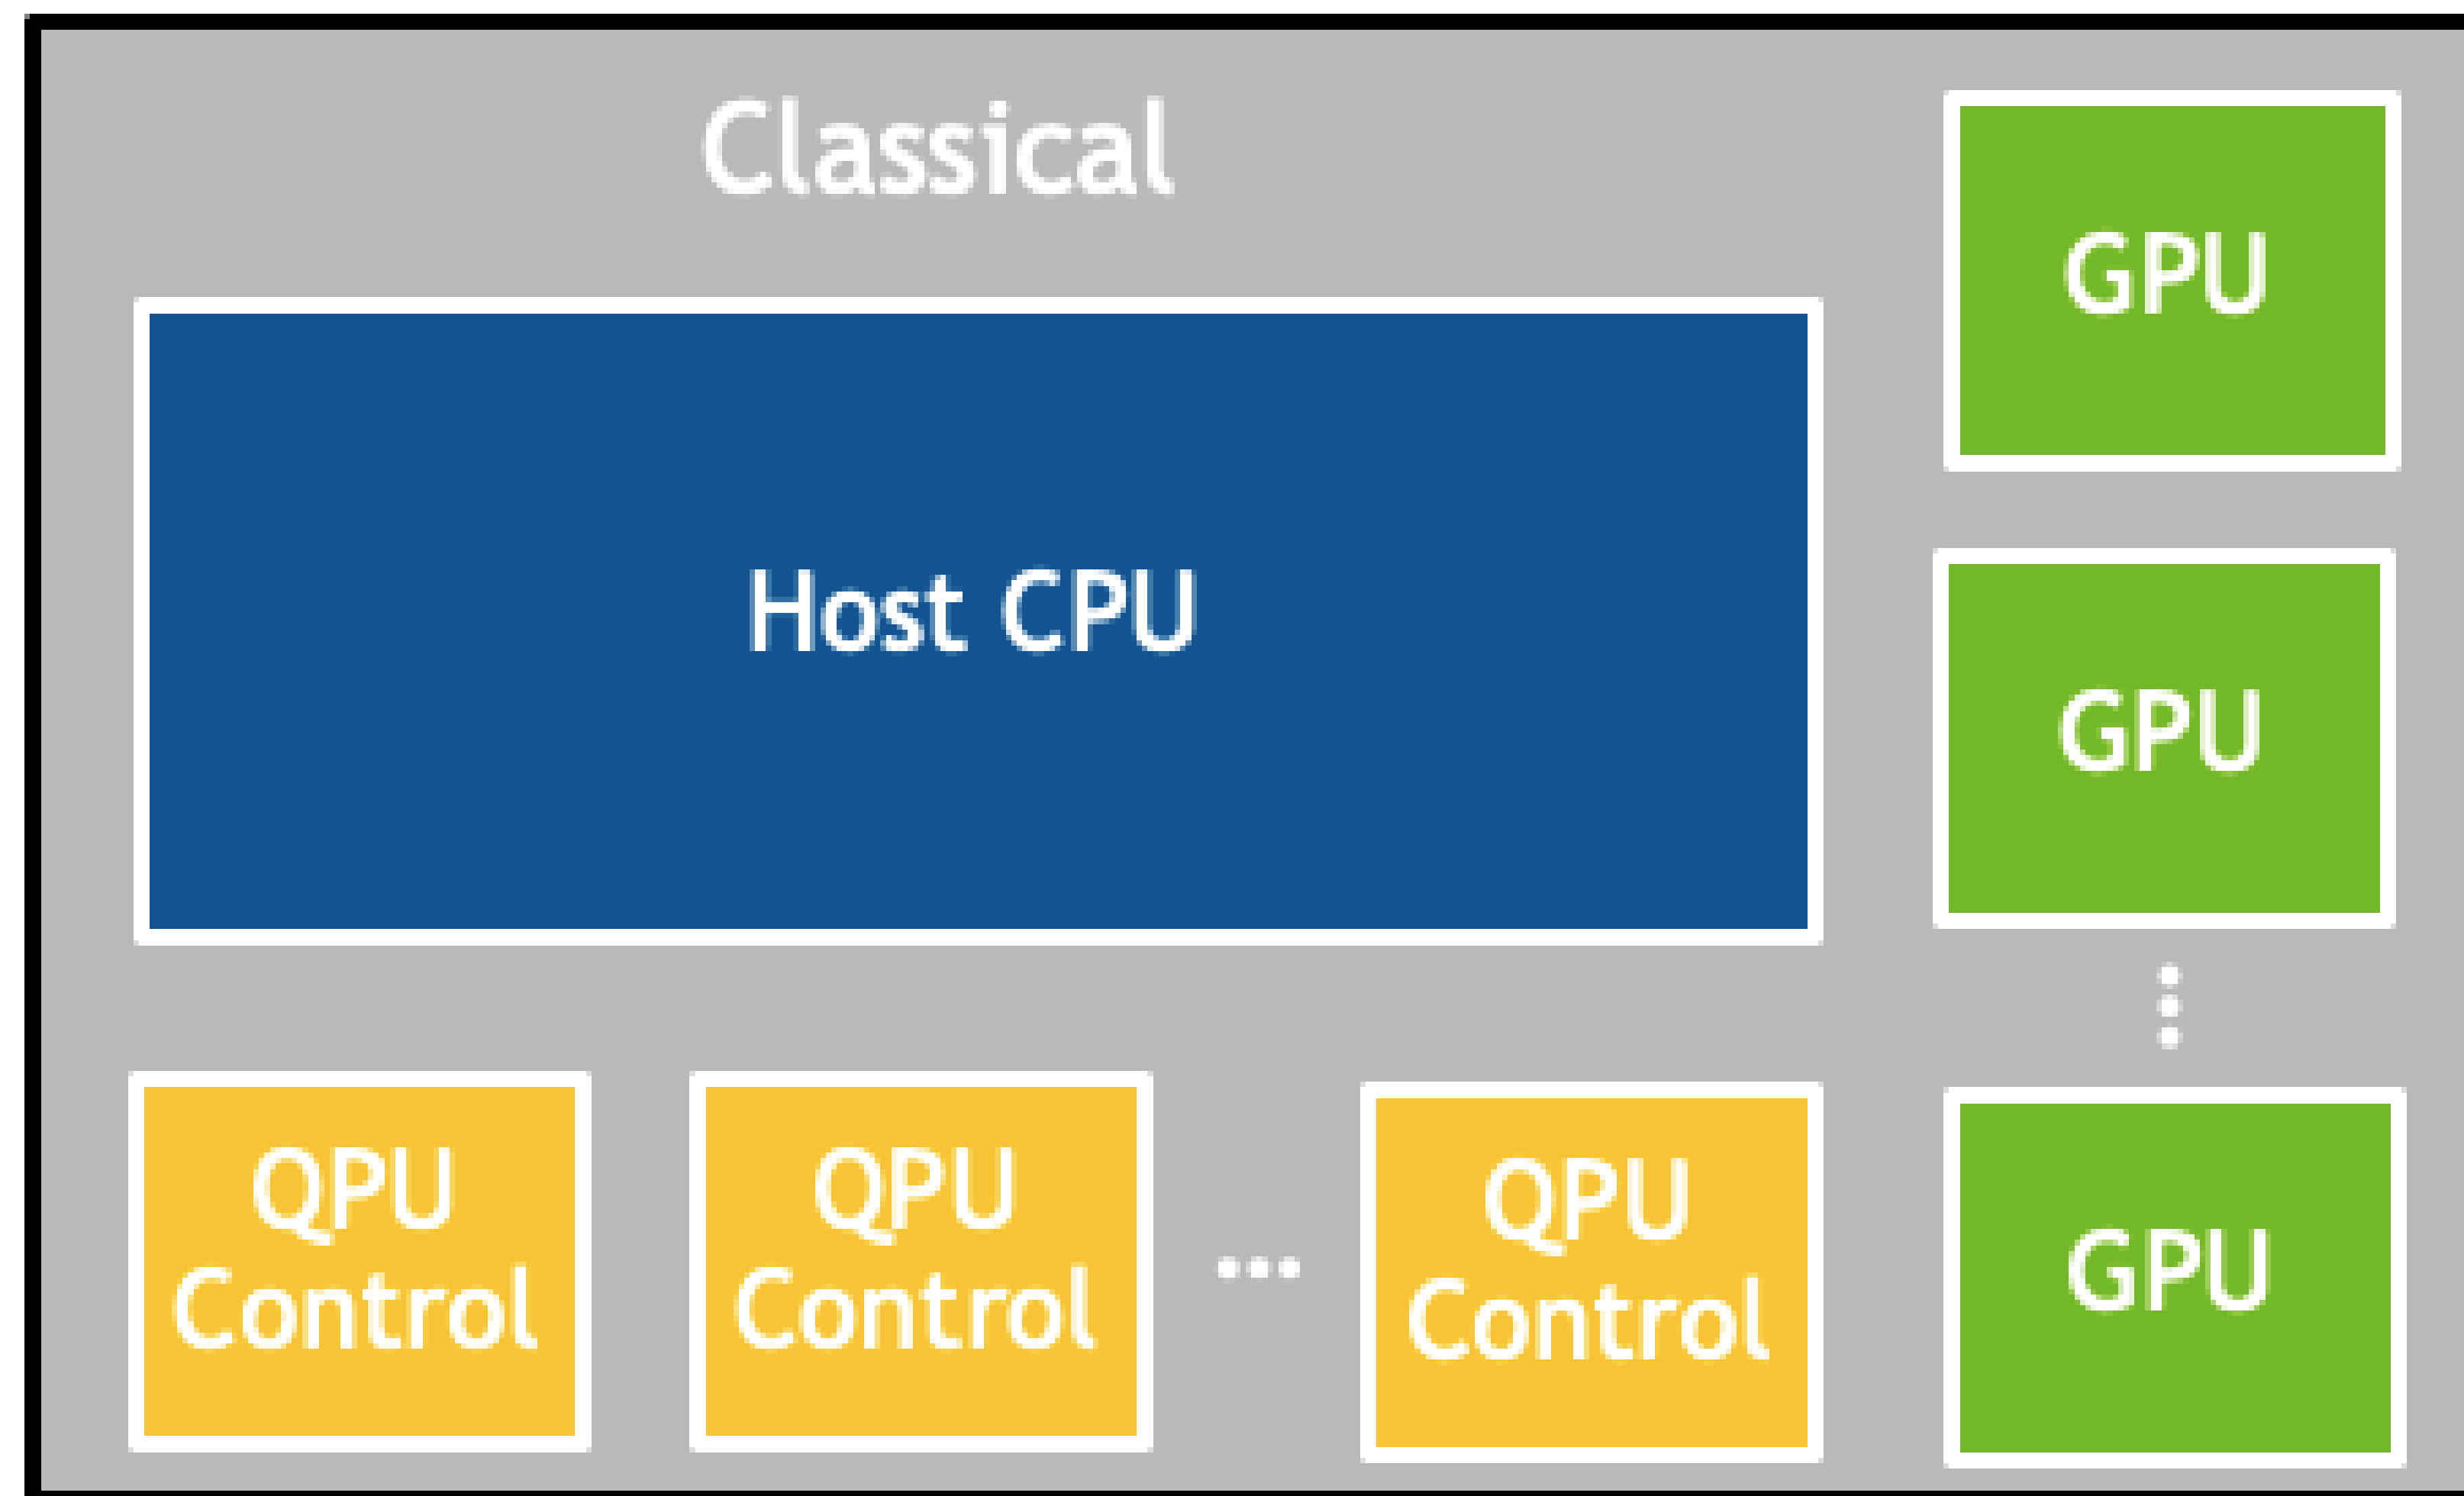
- How-to Guide to CUDA-Q

- Distributed Quantum Computing

- Conclusion

GPU-Accelerated Quantum Computing

Some High-Level Strategies for Parallelization





Agenda

- What is Quantum Accelerated Supercomputing

- Useful Quantum Simulation

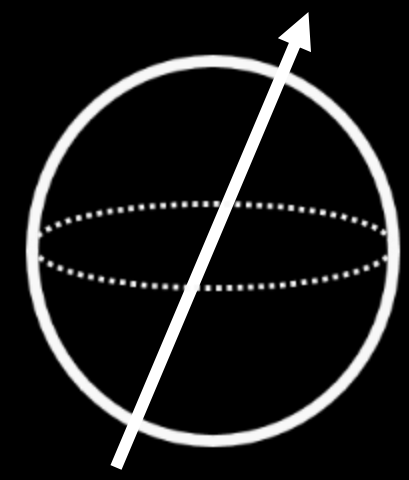
- How-to Guide to CUDA-Q

- Distributed Quantum Computing

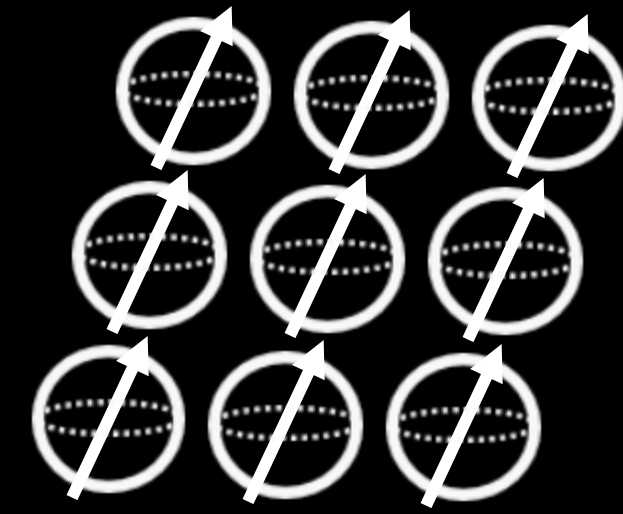
- Conclusion

Quantum: Not Just for Physicists

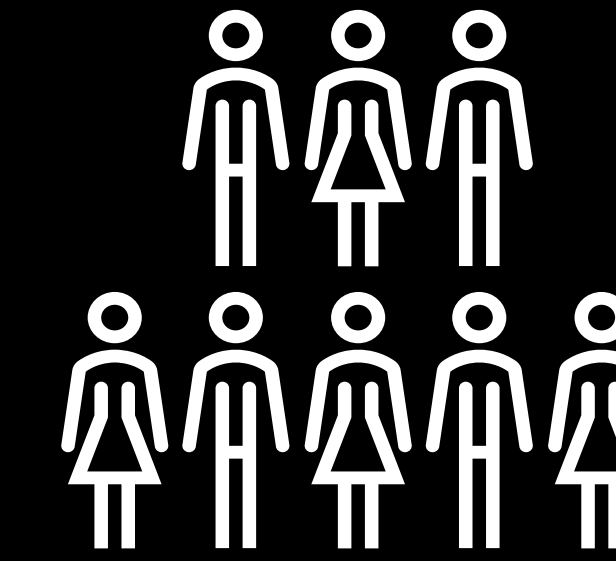
Overcoming these challenges requires broad spectrum of expertise



Qubit Fidelity
99.99% 2-Qubit Gate Fidelity

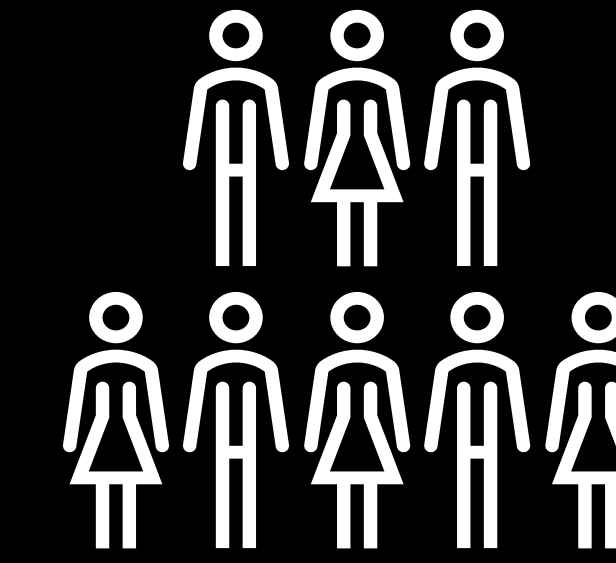


Qubit Scale
100k-1M+ Qubits for FTQC



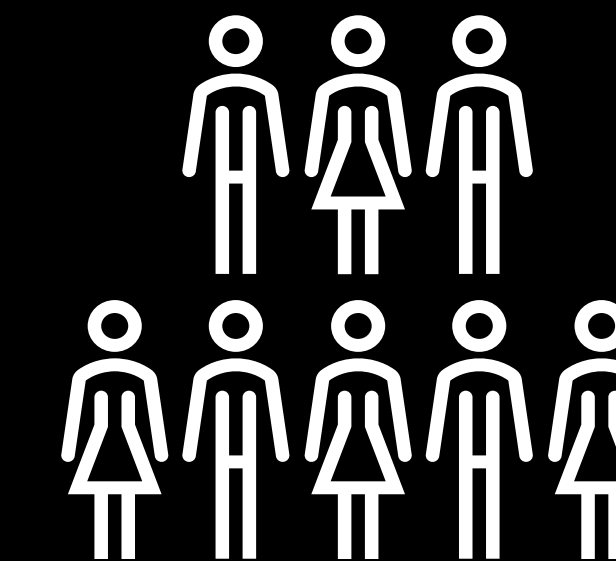
Physicists

Engineers



Computer Scientists

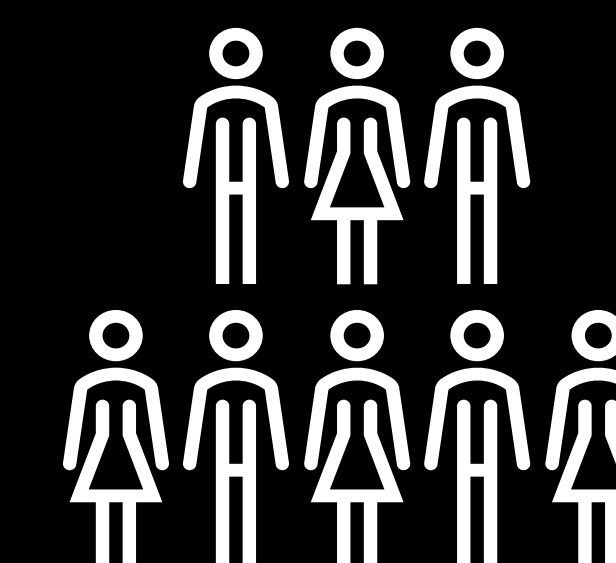
Developers



Mathematicians

Chemists

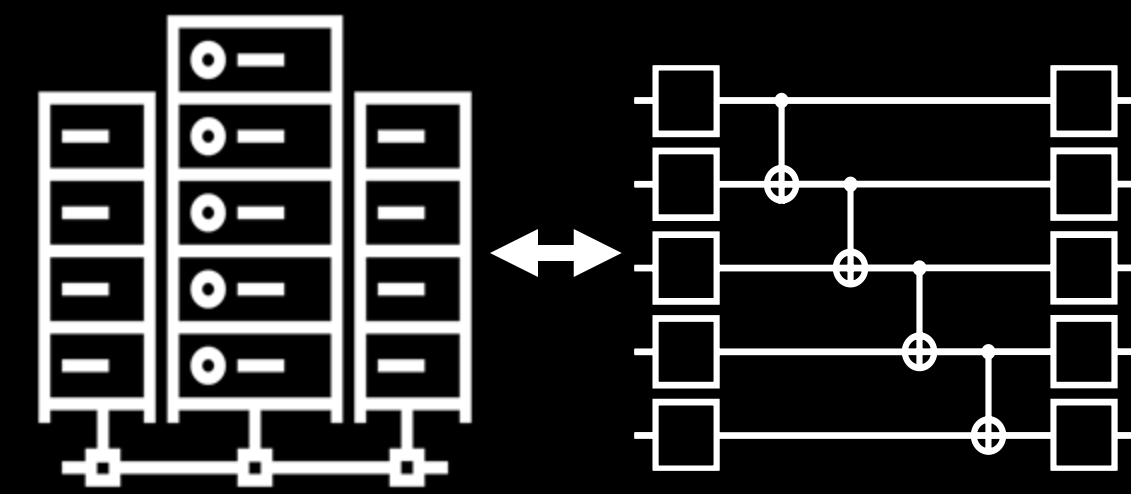
Biologists



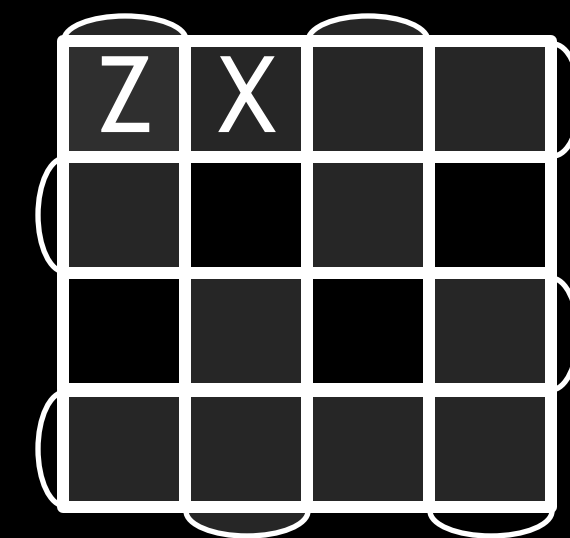
Subject Matter Experts

Students

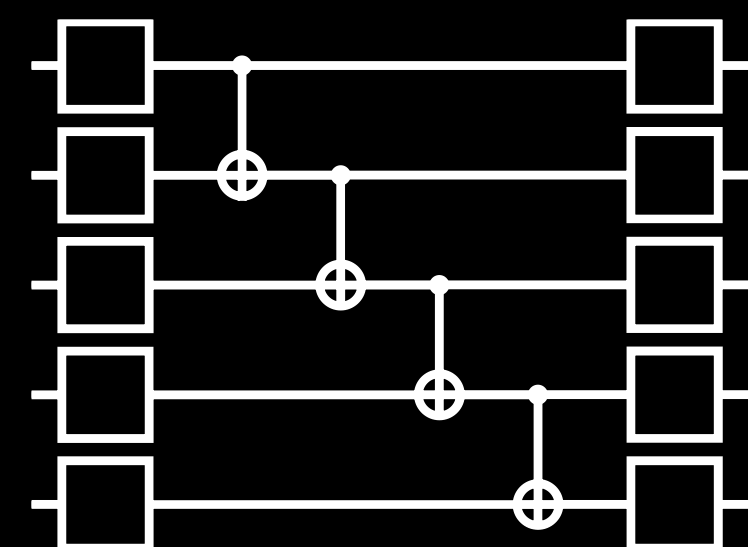
...



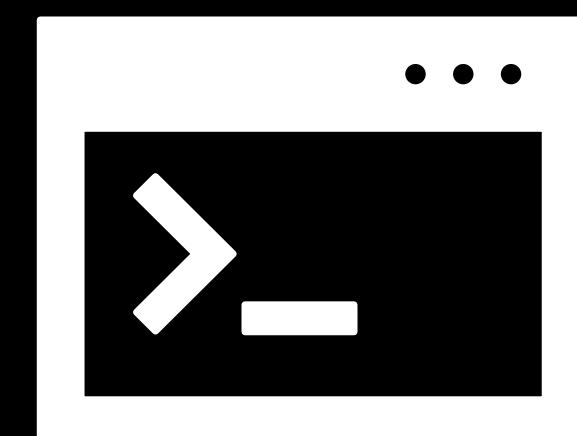
HPC Integration
Sub-Microsecond HPC-QC Latency



Error Correction
Methods that Scale to Large Quantum Systems



Algorithms
Algorithms with Exponential Speed-up



Developer Tools
Integrate with Scientific Computing
Familiar to non-Quantum Physicists



